

Fast Transparent Virtual Machine Migration in Distributed Edge Clouds

Lucas Chaufournier

Prateek Sharma

Franck Le

Erich Nahum

Prashant Shenoy

Don Towsley



UMASS
AMHERST



Cloud Computing

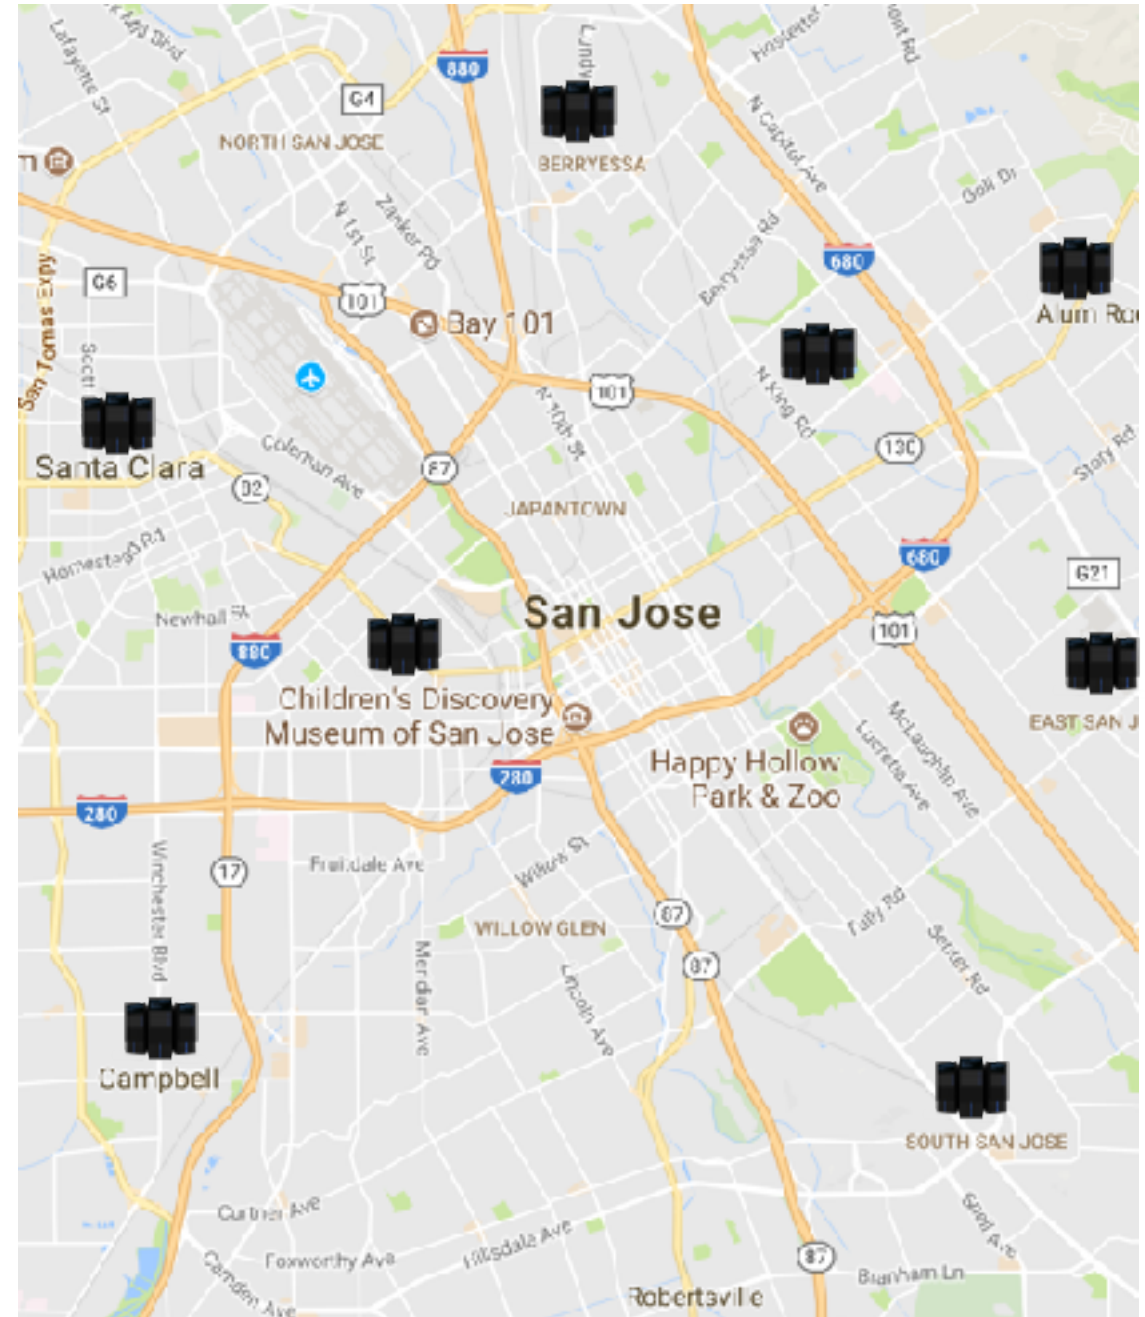
- Modern applications are deployed in centralized data centers.
- Few locations each with lots of servers.
- Benefits:
 - Easy to deploy and scale applications
- Limitations:
 - High Latencies
 - Low WAN Bandwidth

AWS Regions



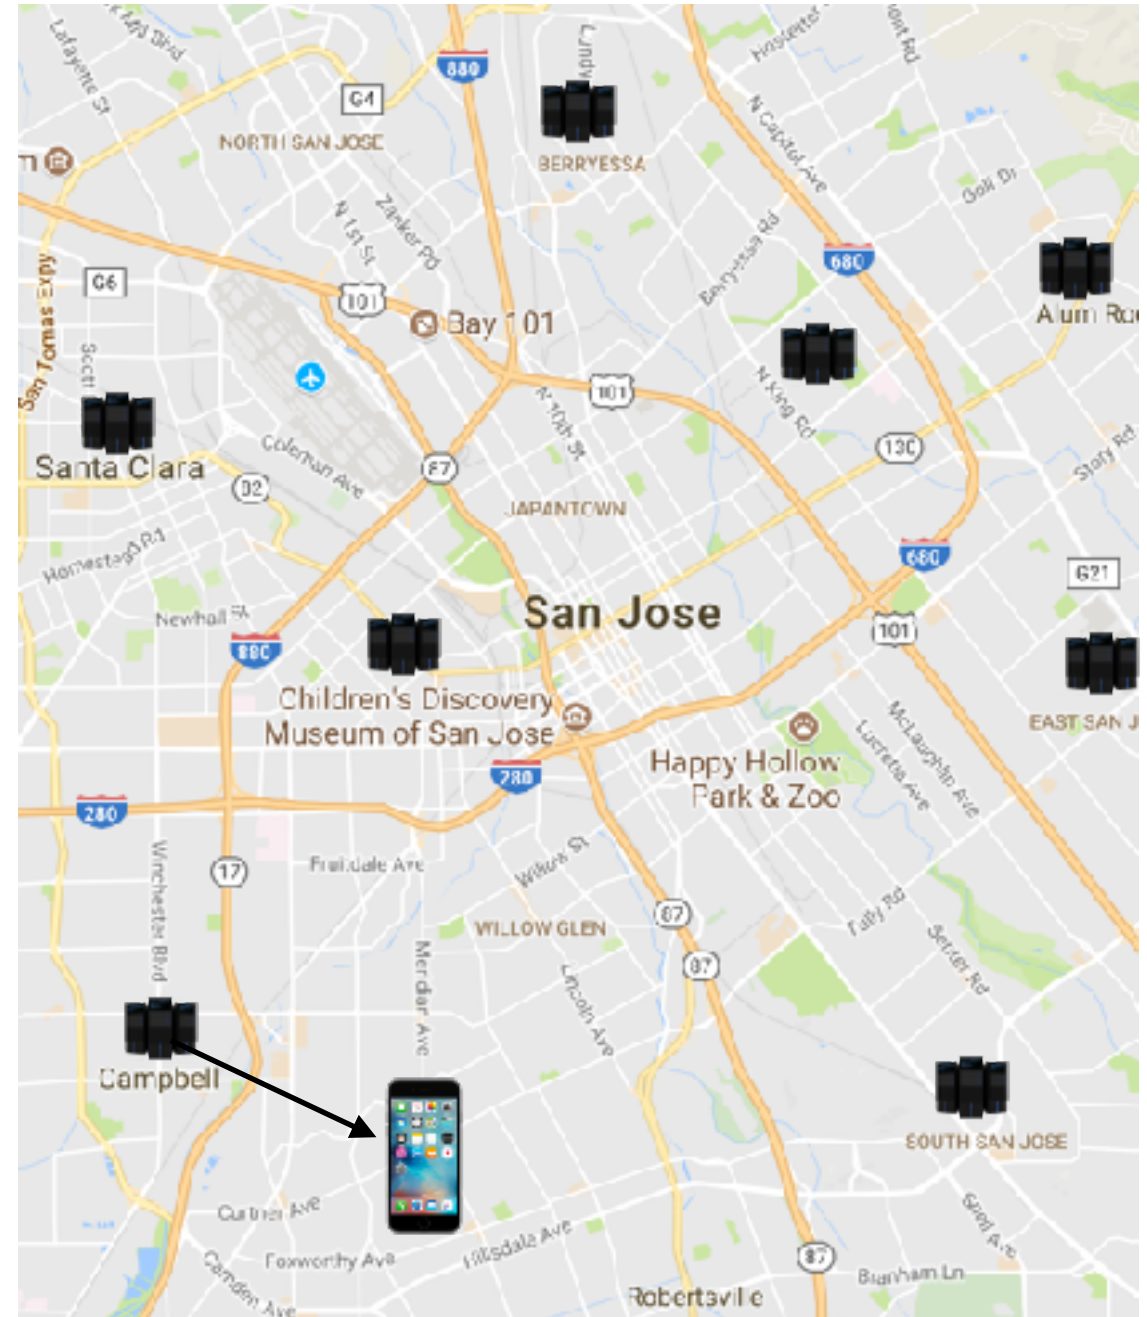
Edge Clouds

- Paradigm where servers are located at the edge rather than remote data centers.
- Lots of locations each with few servers.
- Allows you to host applications closer to users with lower latencies.



Edge Clouds

- Users increasingly mobile
- User mobility implies Latency increase
- Potential Solution: Migrating workloads to the closest site.



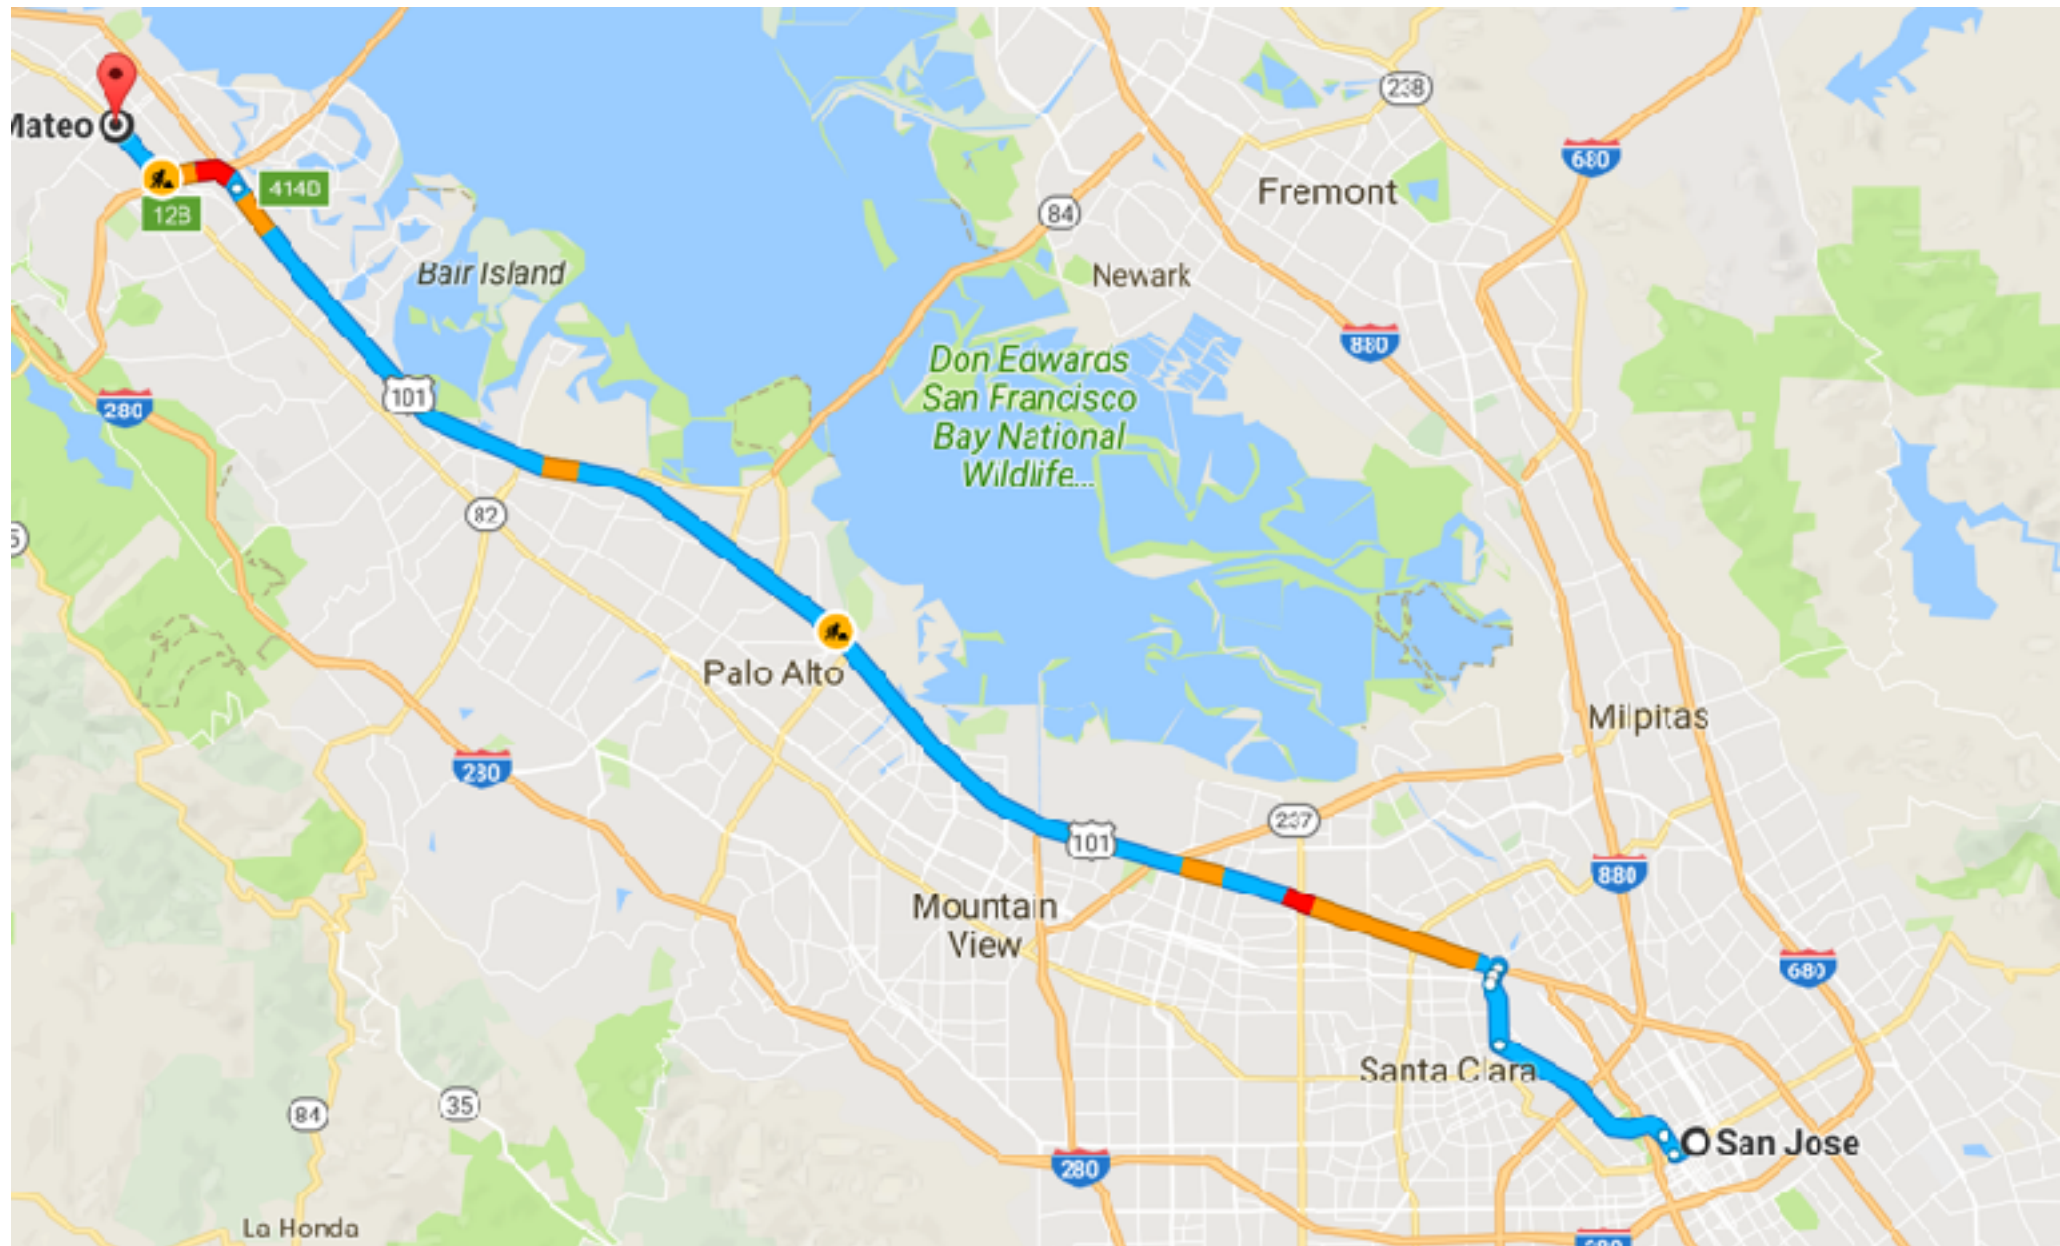
Meet Alice



Meet Alice

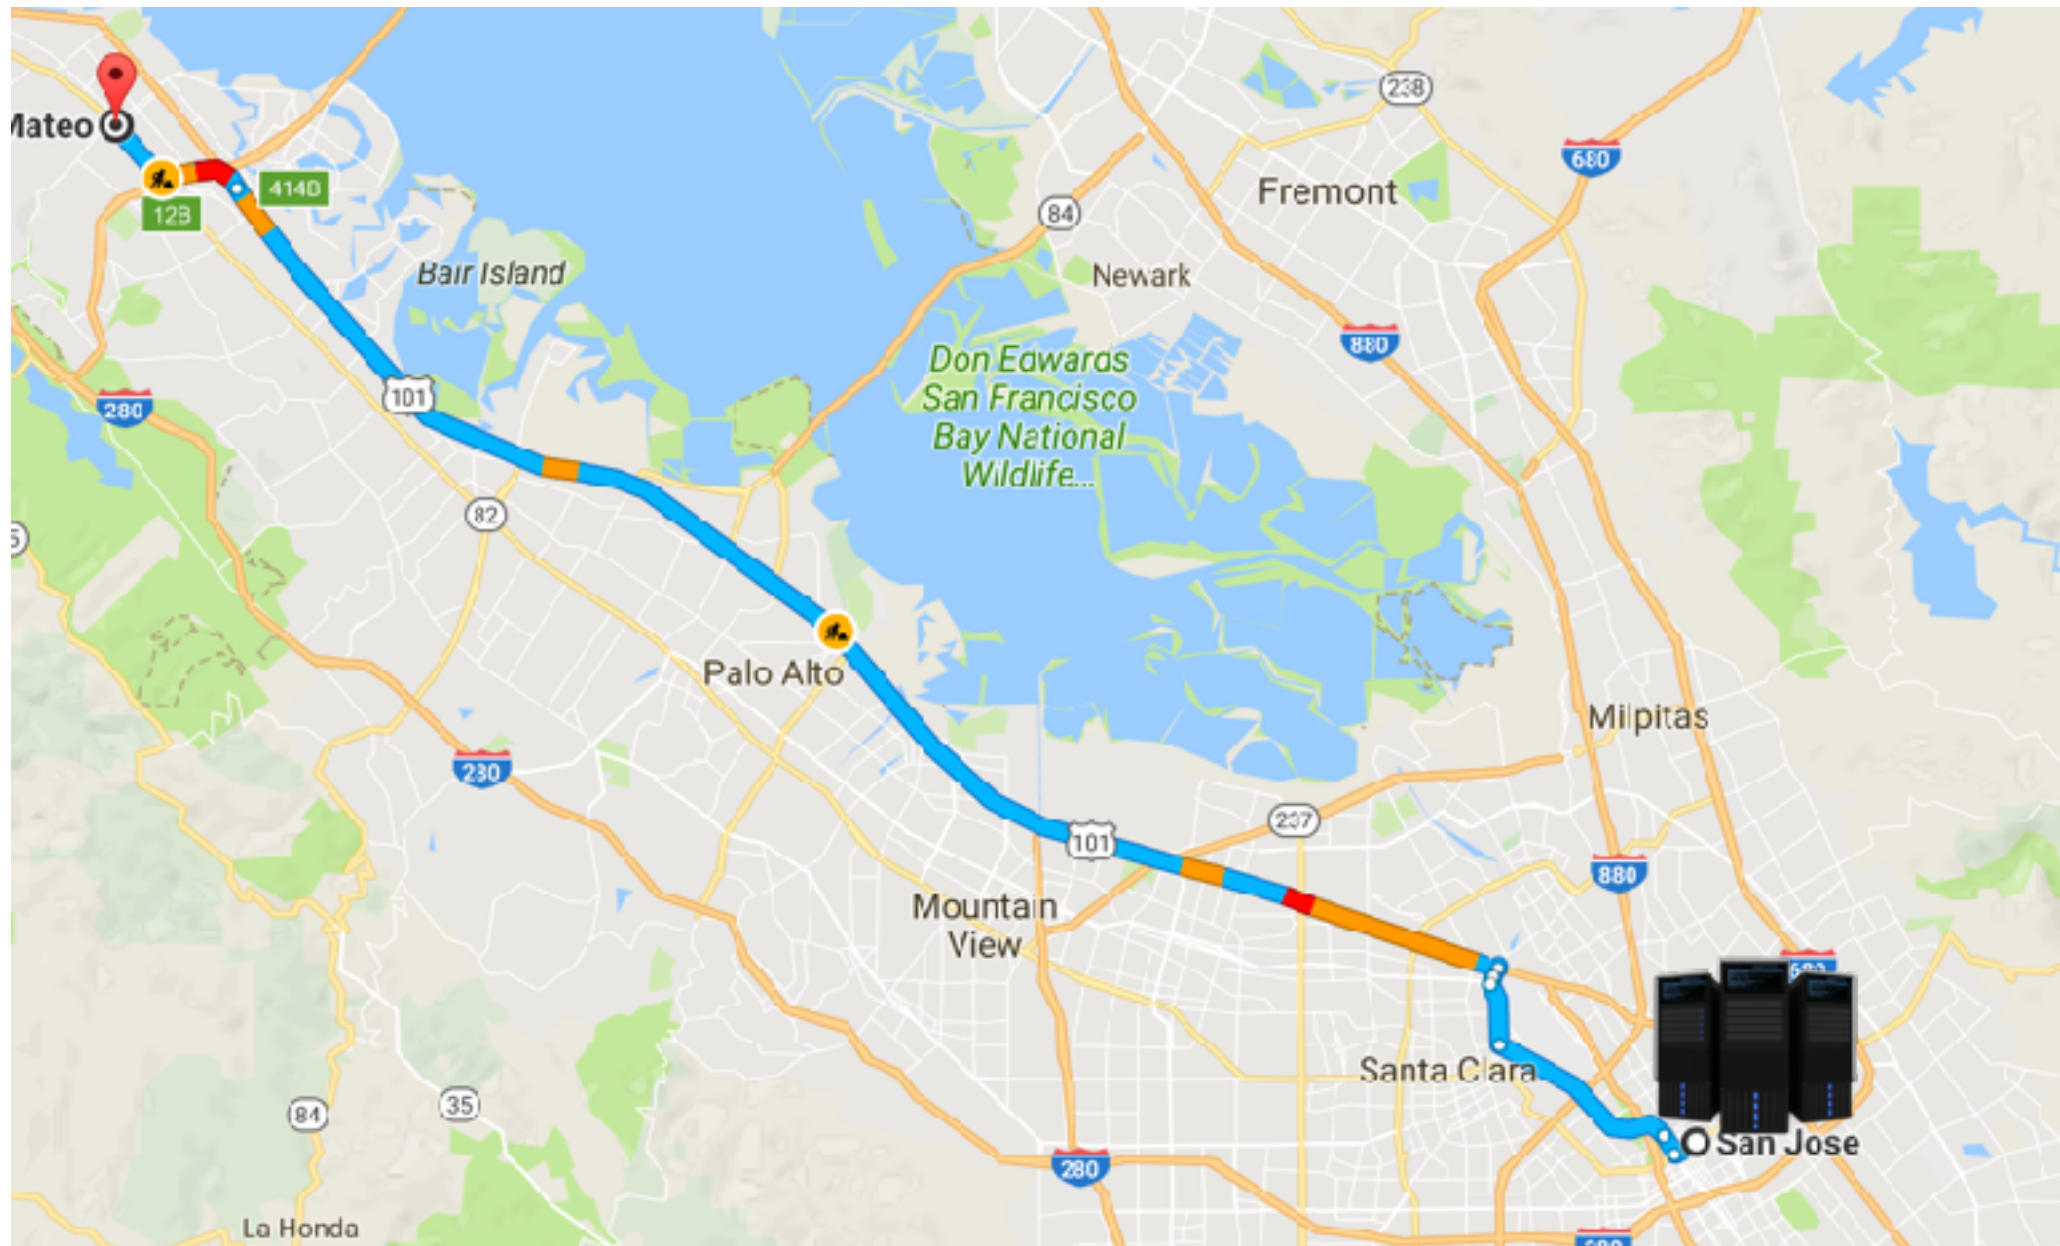


- A software engineer
- Develops Applications for Augmented Reality(AR)
- Often works from home.
- Tests her AR apps during her commute.



Alice works in San Jose but lives in San Mateo.

An hour long commute by train...



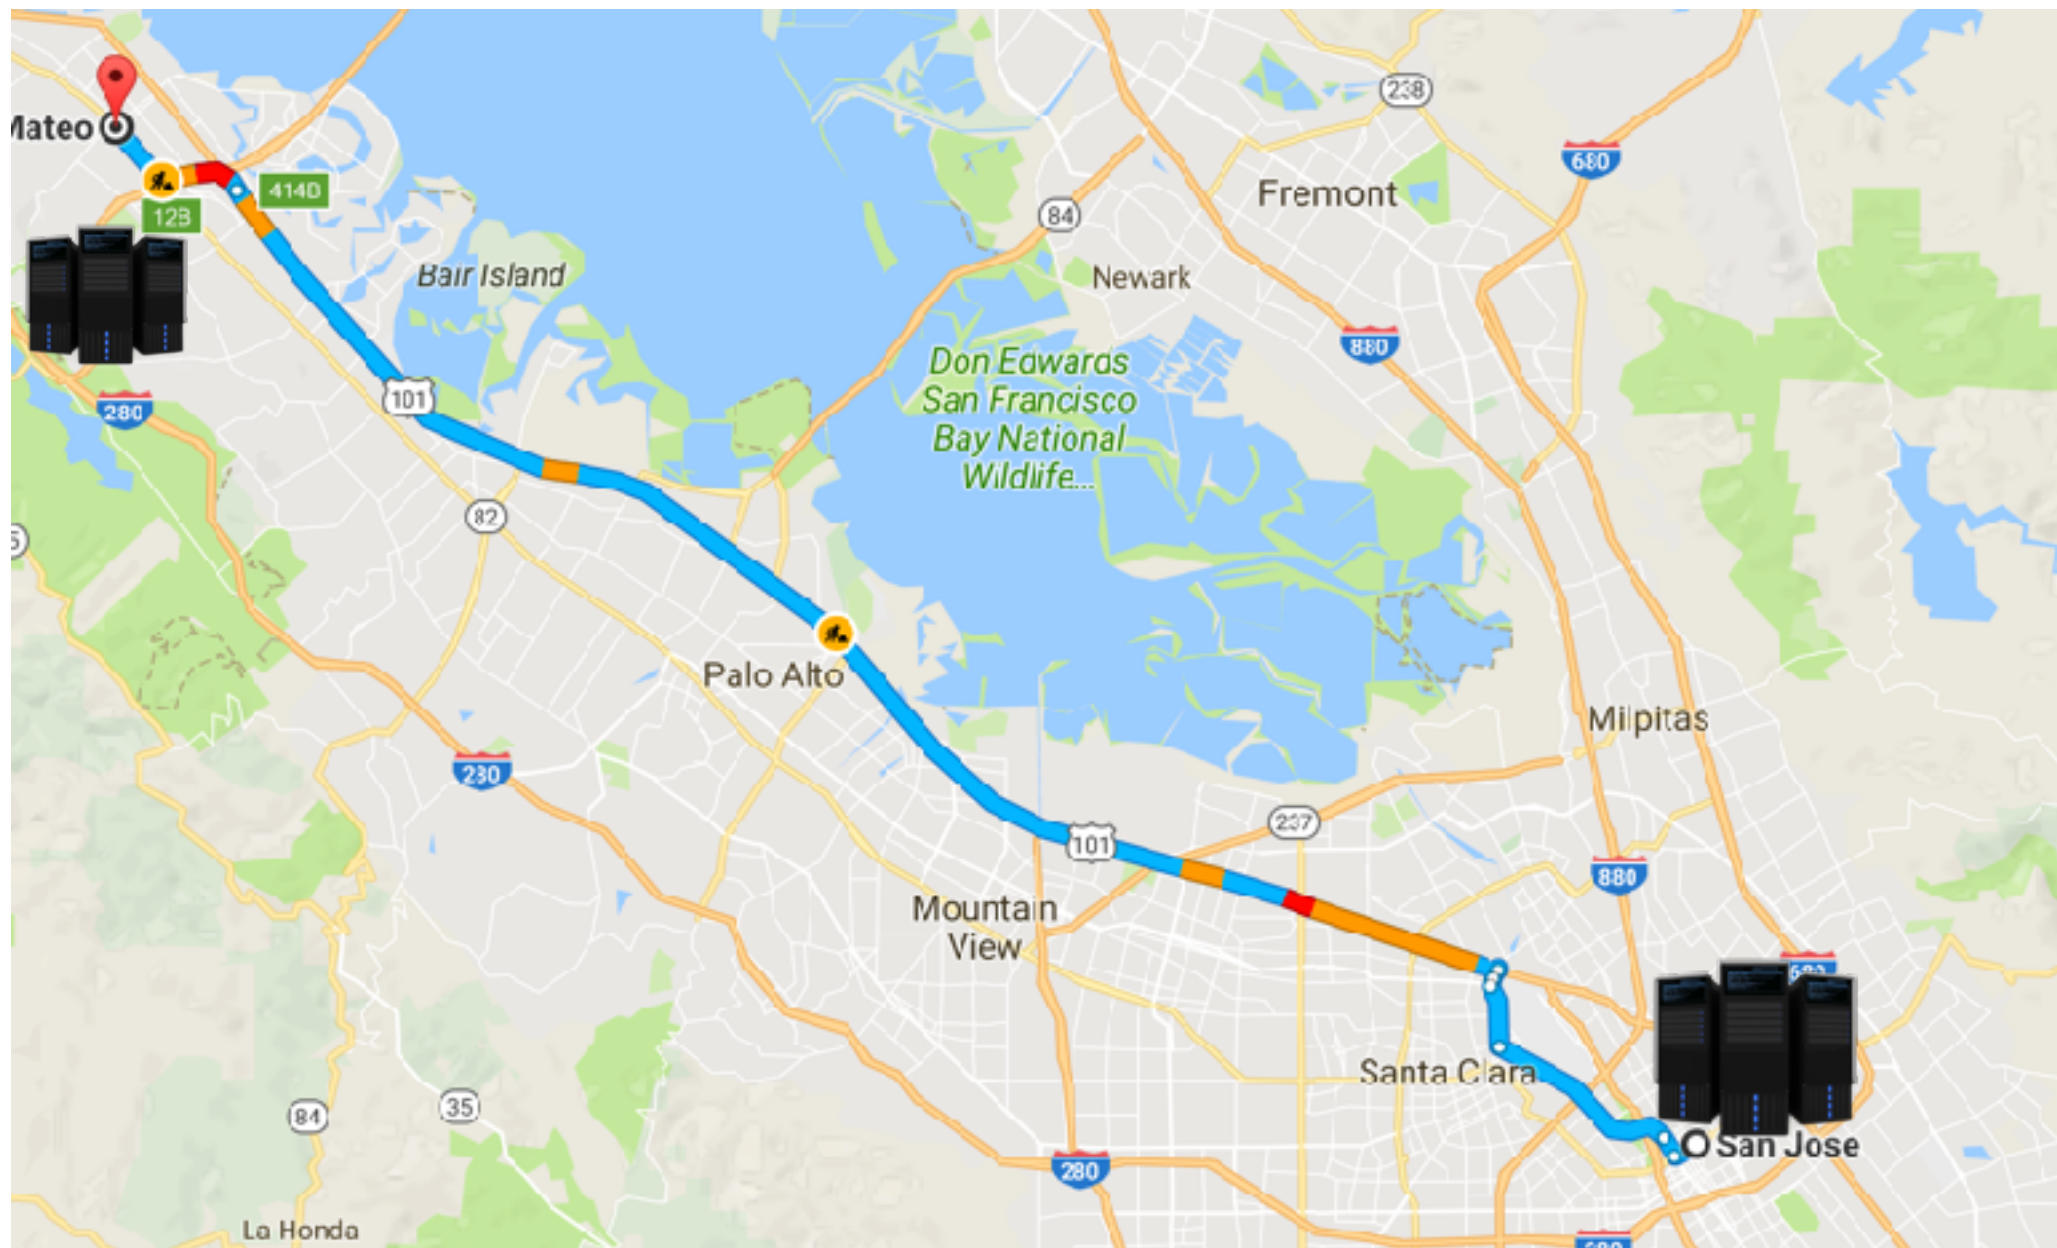
**Her company servers are located at a
San Jose edge cloud.**

This makes testing her AR apps and working from home a problem...

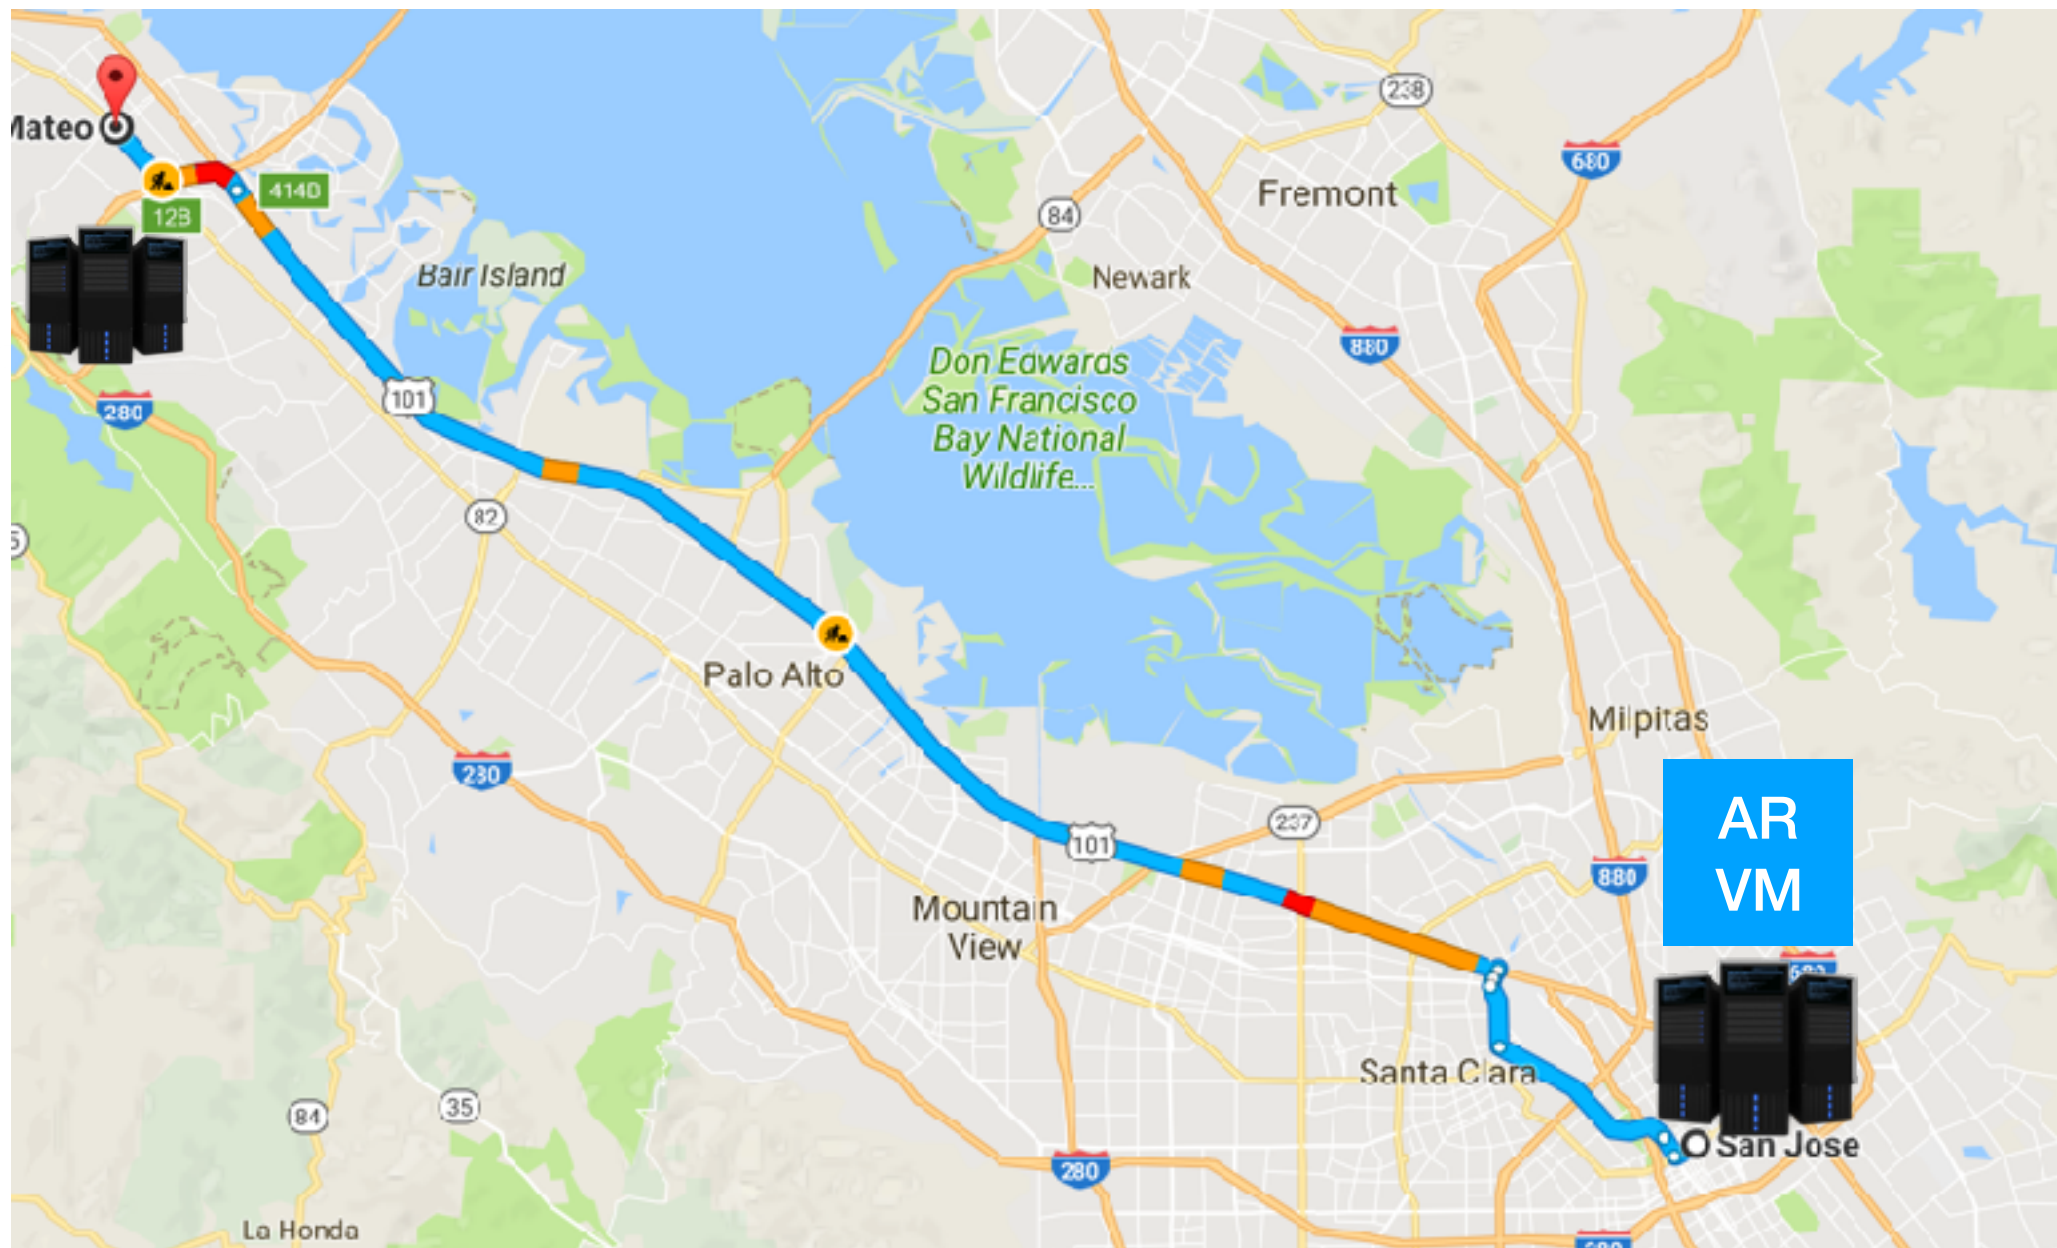


As Alice gets closer to home, Latencies start to increase.

Higher latencies make it impossible to test her apps.
Working from home means she can't test her apps.

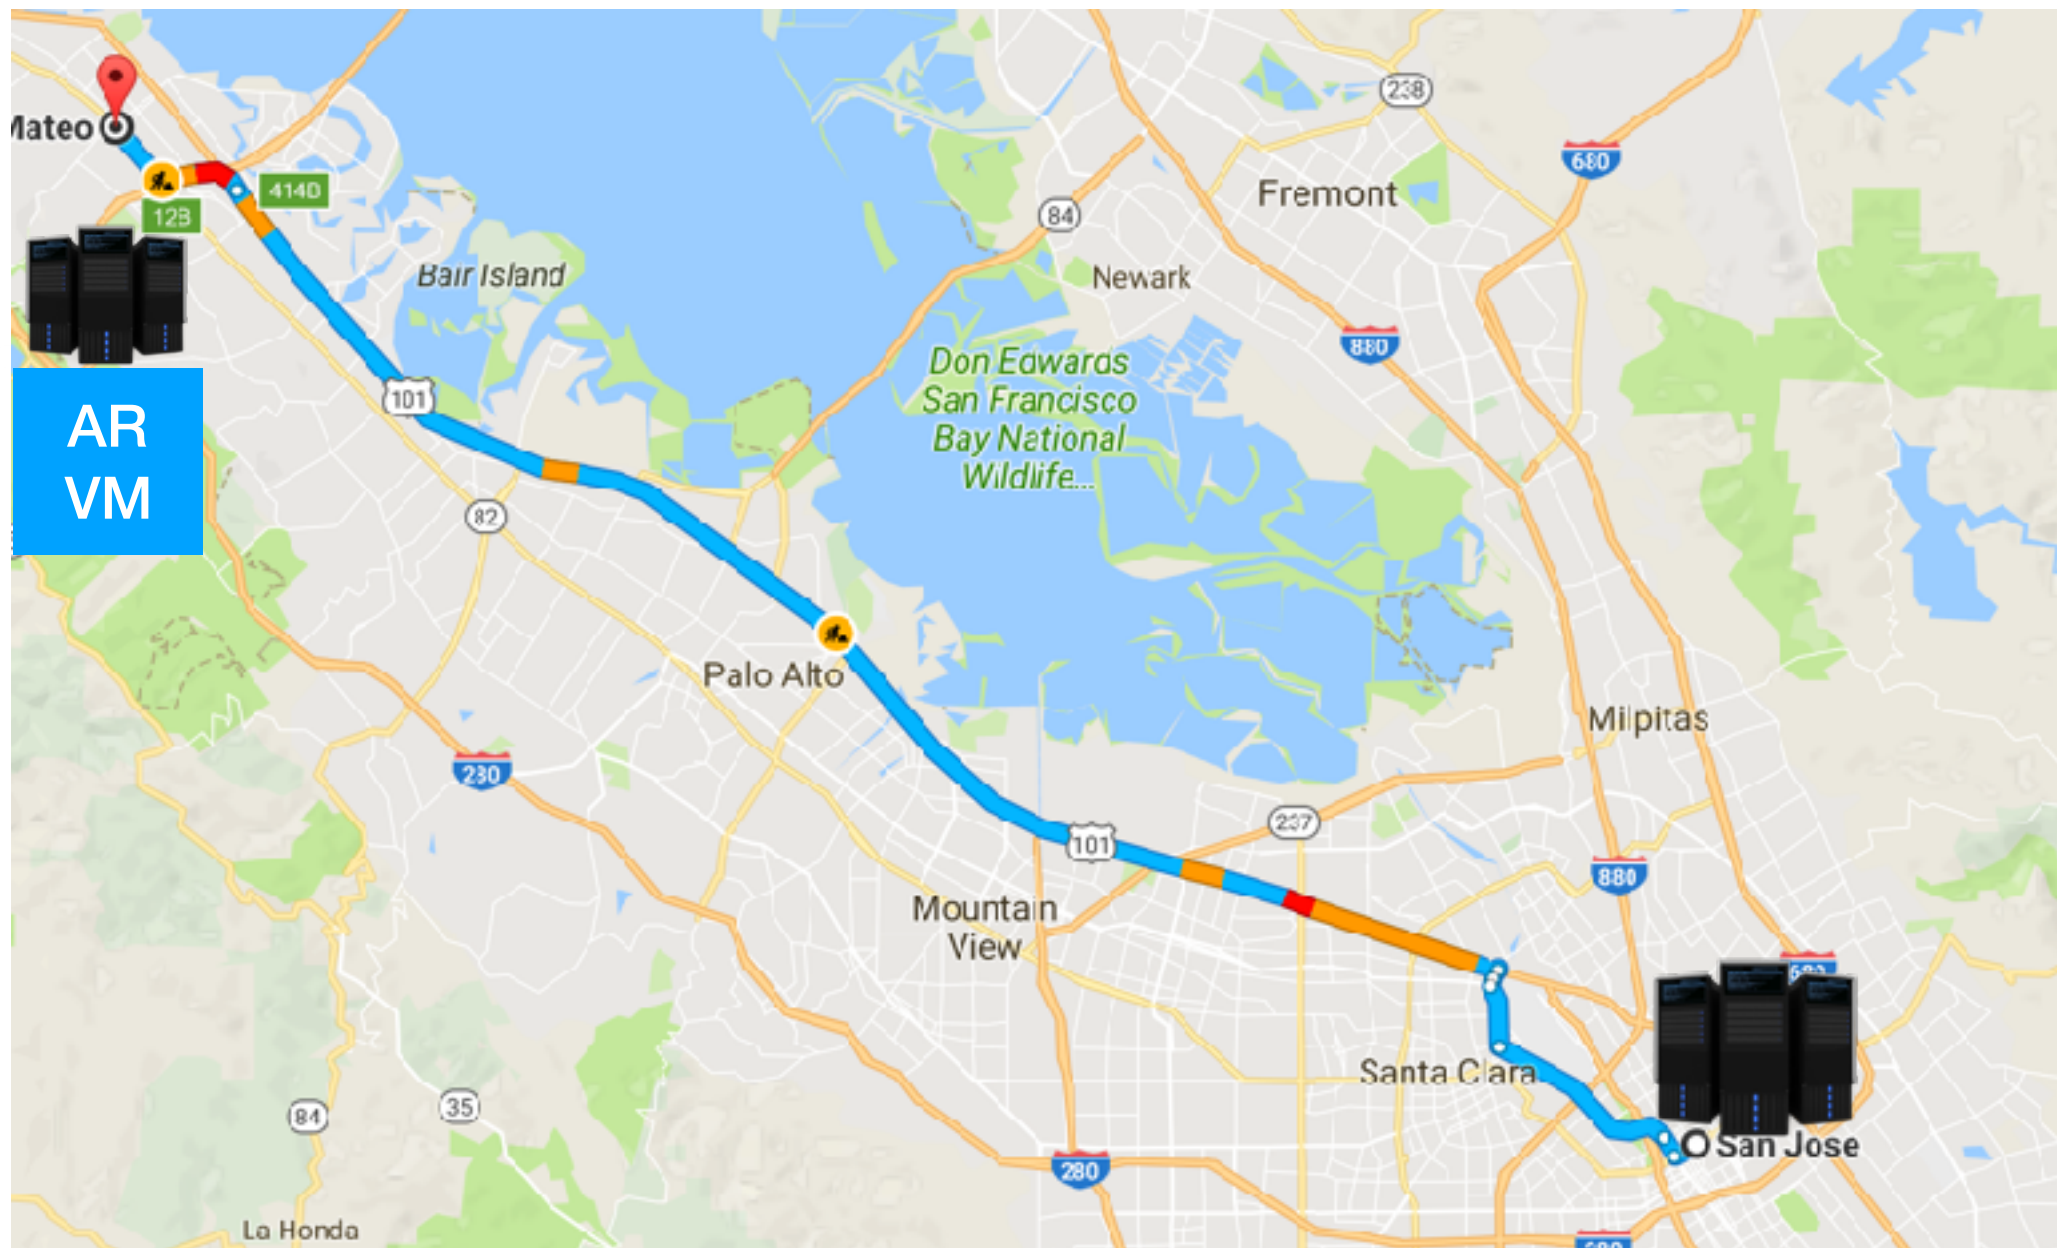


Luckily, there's an edge cloud location closer to her home.



Luckily, there's an edge cloud location closer to her home.

Cloud can migrate her AR application to a closer location.



Luckily, there's an edge cloud location closer to her home.

Cloud can migrate her AR application to a closer location.

Problem Statement


- **Potential Solution:** Live Migrate the VM from one edge cloud location to another.
- Two potential problems in VM Migrations in Edge Clouds
 - Low WAN Bandwidth slows down transfers.
 - IP address changes between locations hurts application latencies.

Technique for Transparent Live VM Migrations in Edge Clouds.



Outline

- Motivation and problem statement
- Migration & MPTCP Background
- Implementation
- Results
- Conclusion

Anatomy of VM Migration




San Jose




San Mateo

Migration: Three States



**Memory
State**

Migration: Three States




**Memory
State**



**Disk
State**

Migration: Three States



**Memory
State**



**Disk
State**

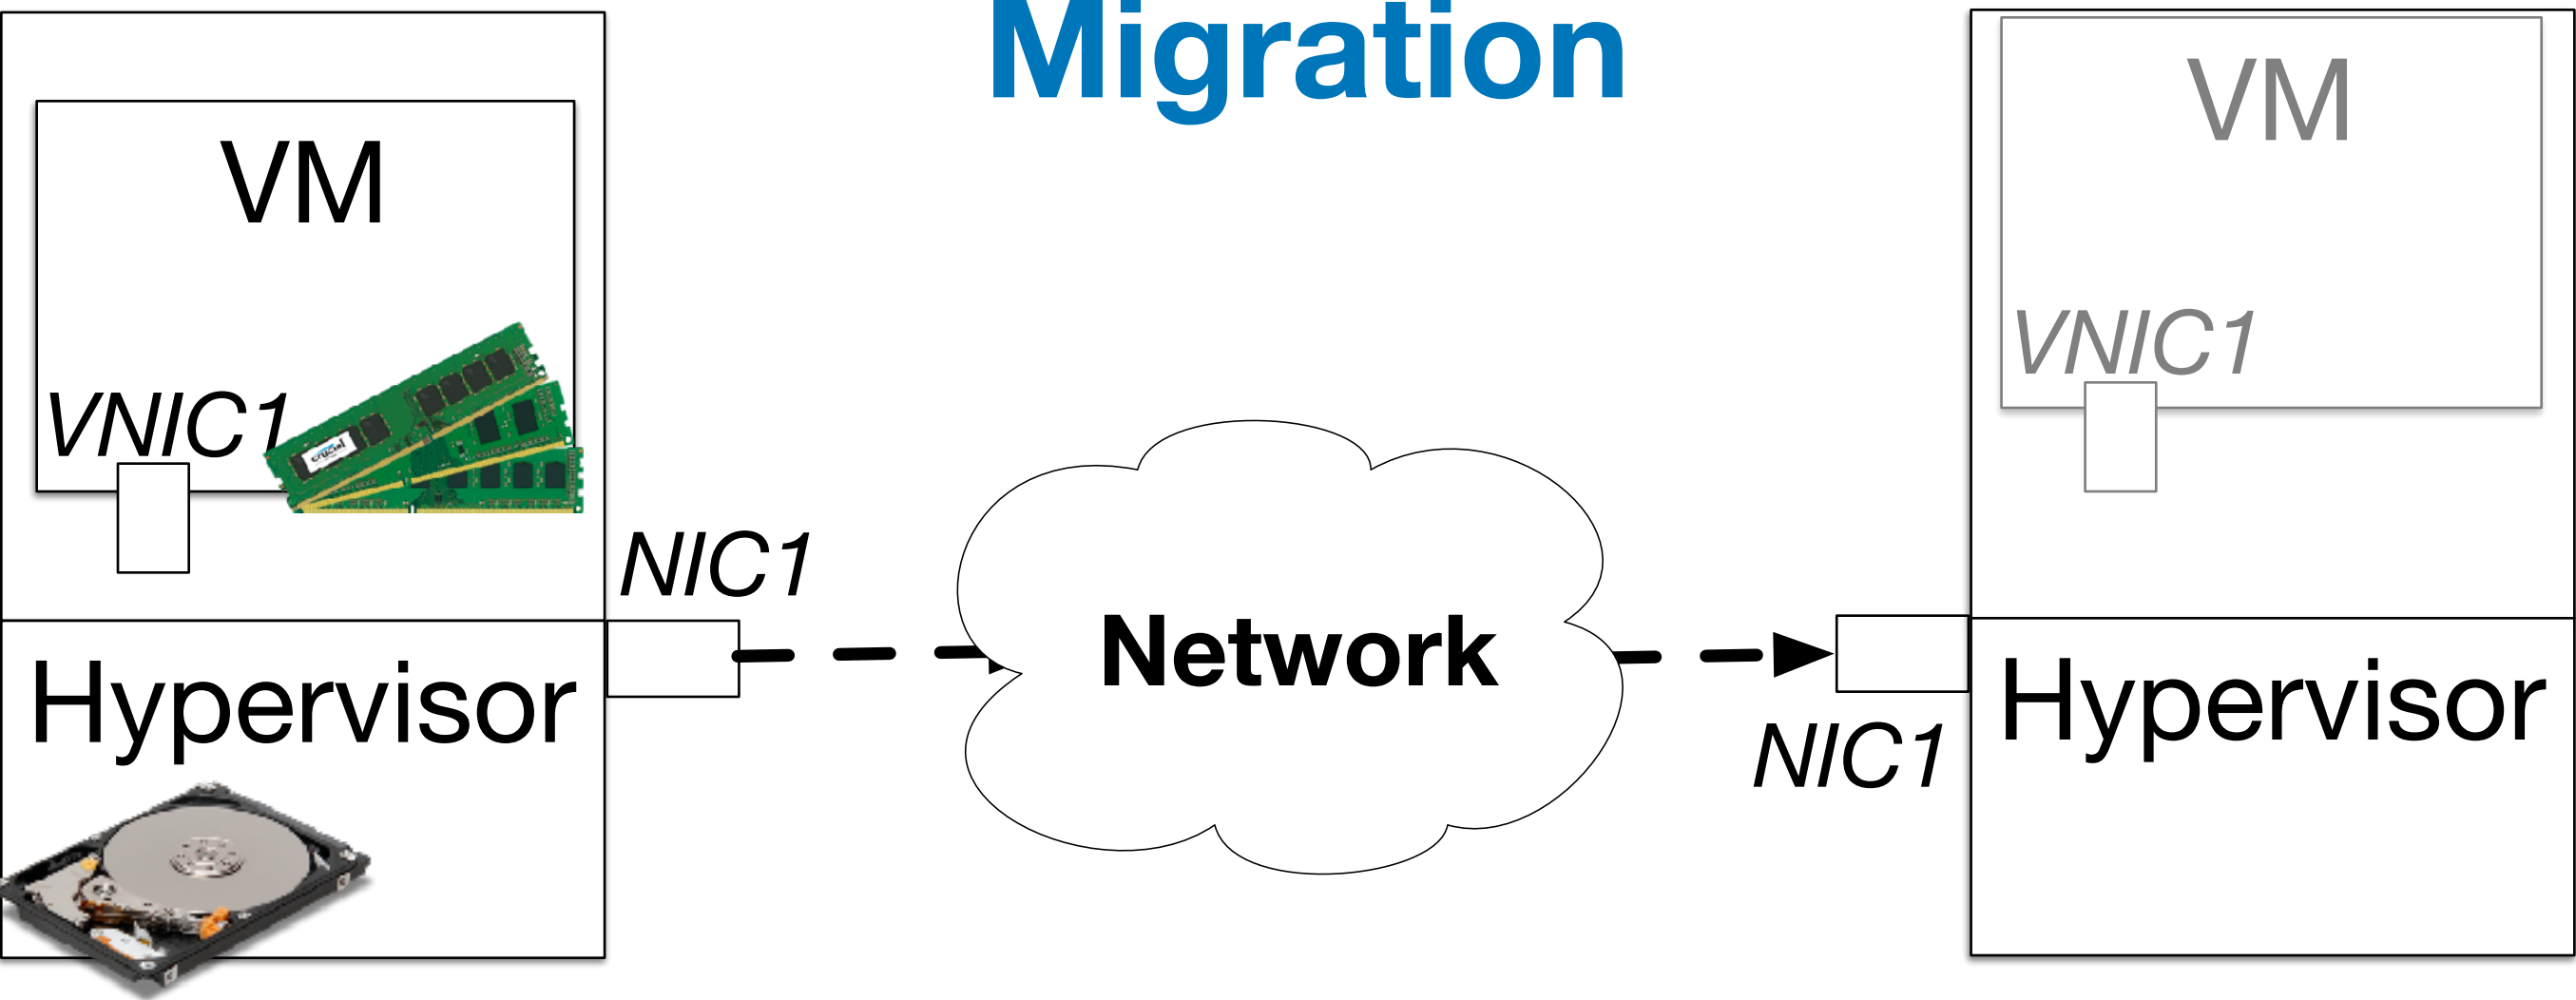


192.168.1.1

**Network
State**

Anatomy of VM Migration

Begin Migration

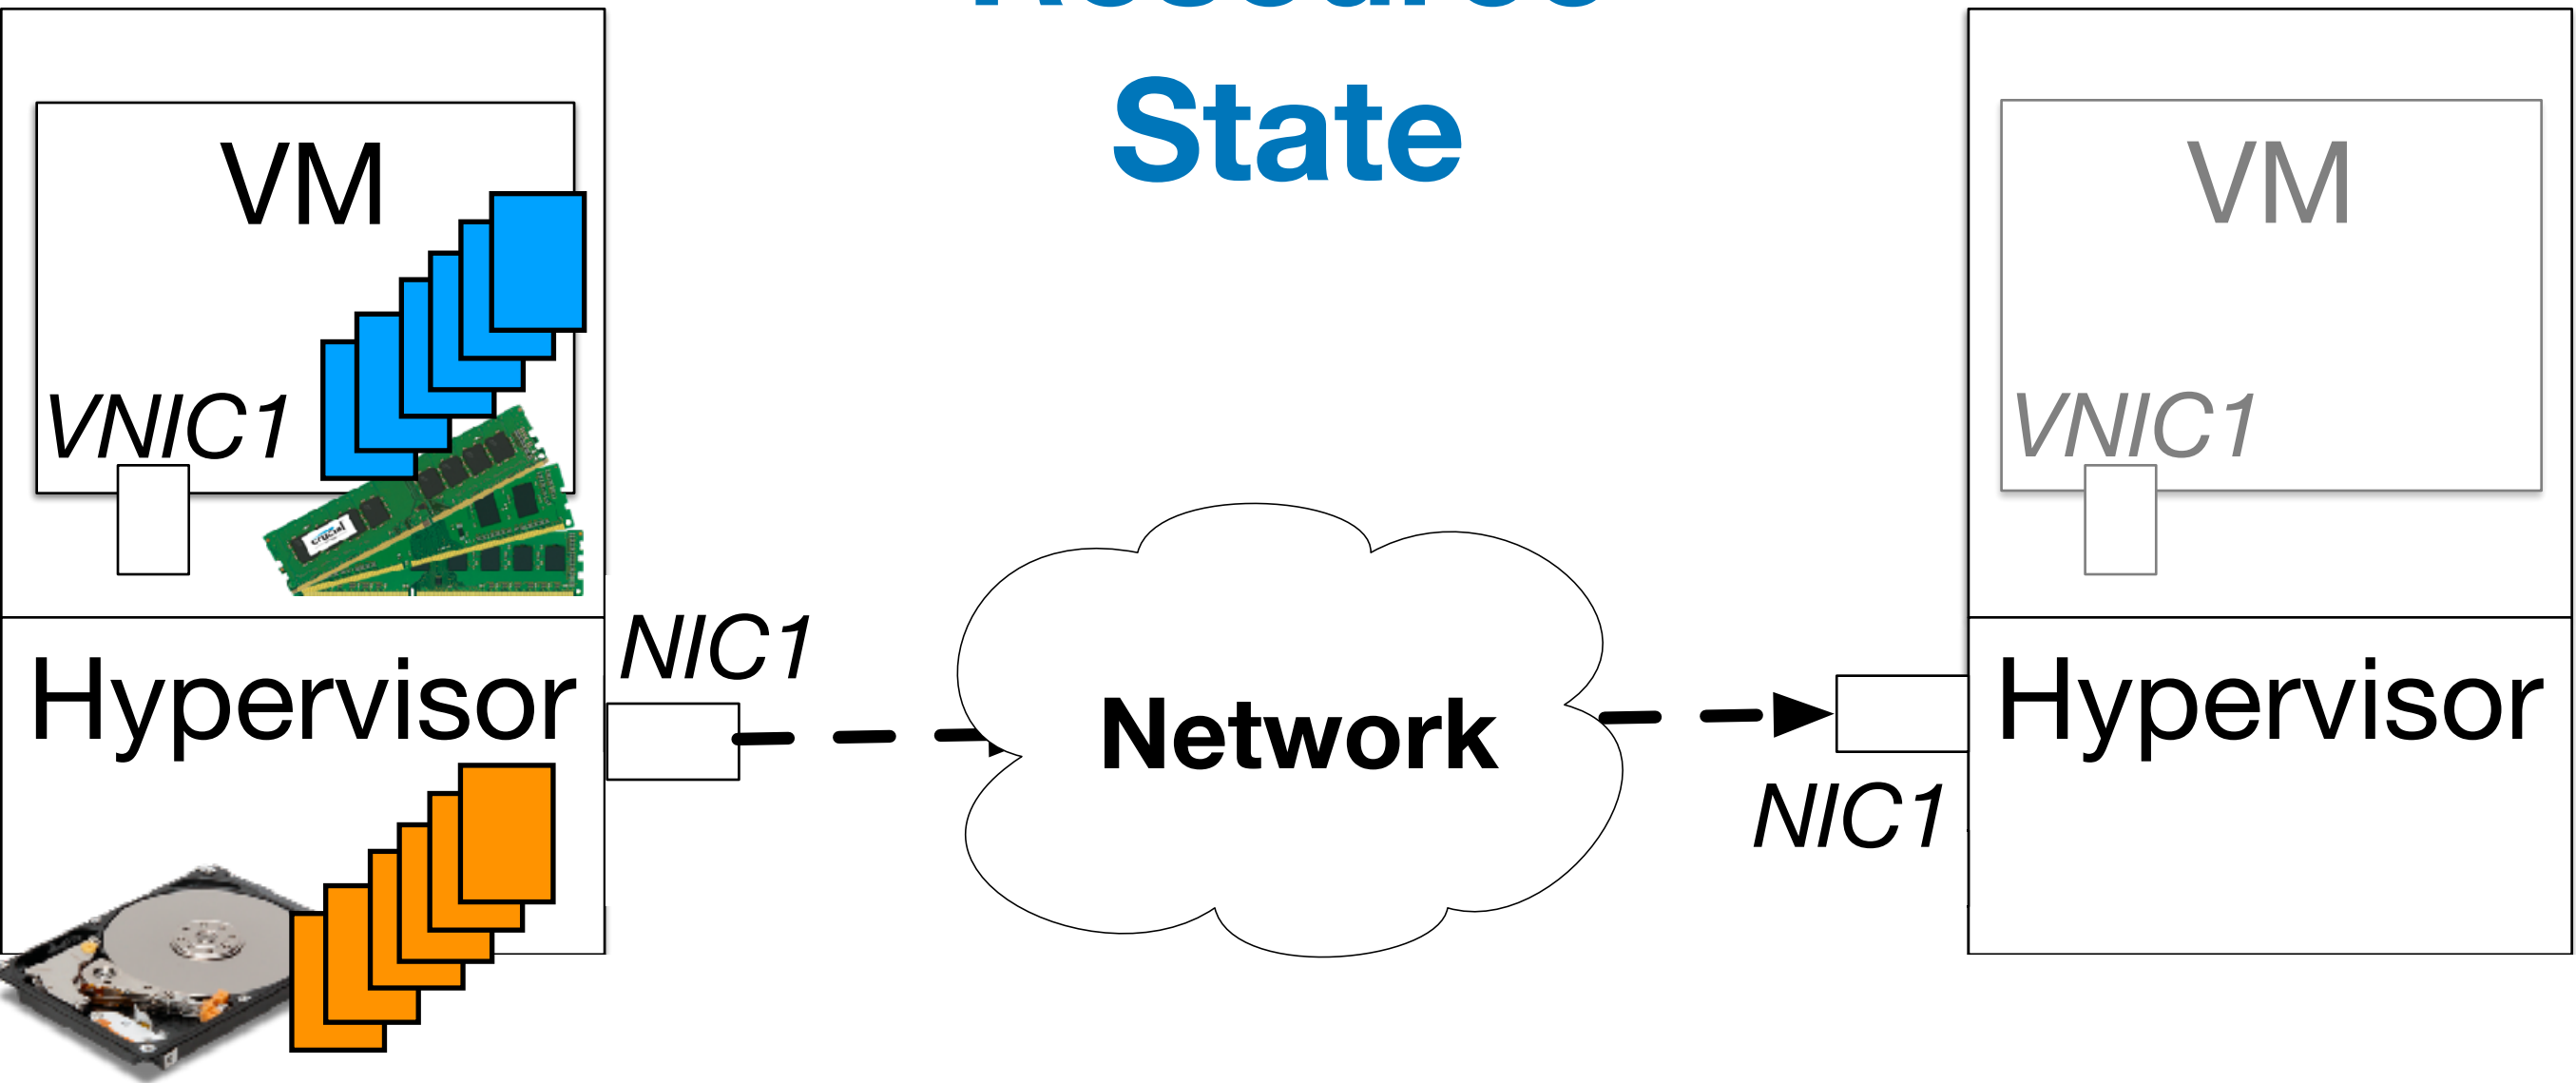


San Jose

San Mateo

Anatomy of VM Migration

Resource State

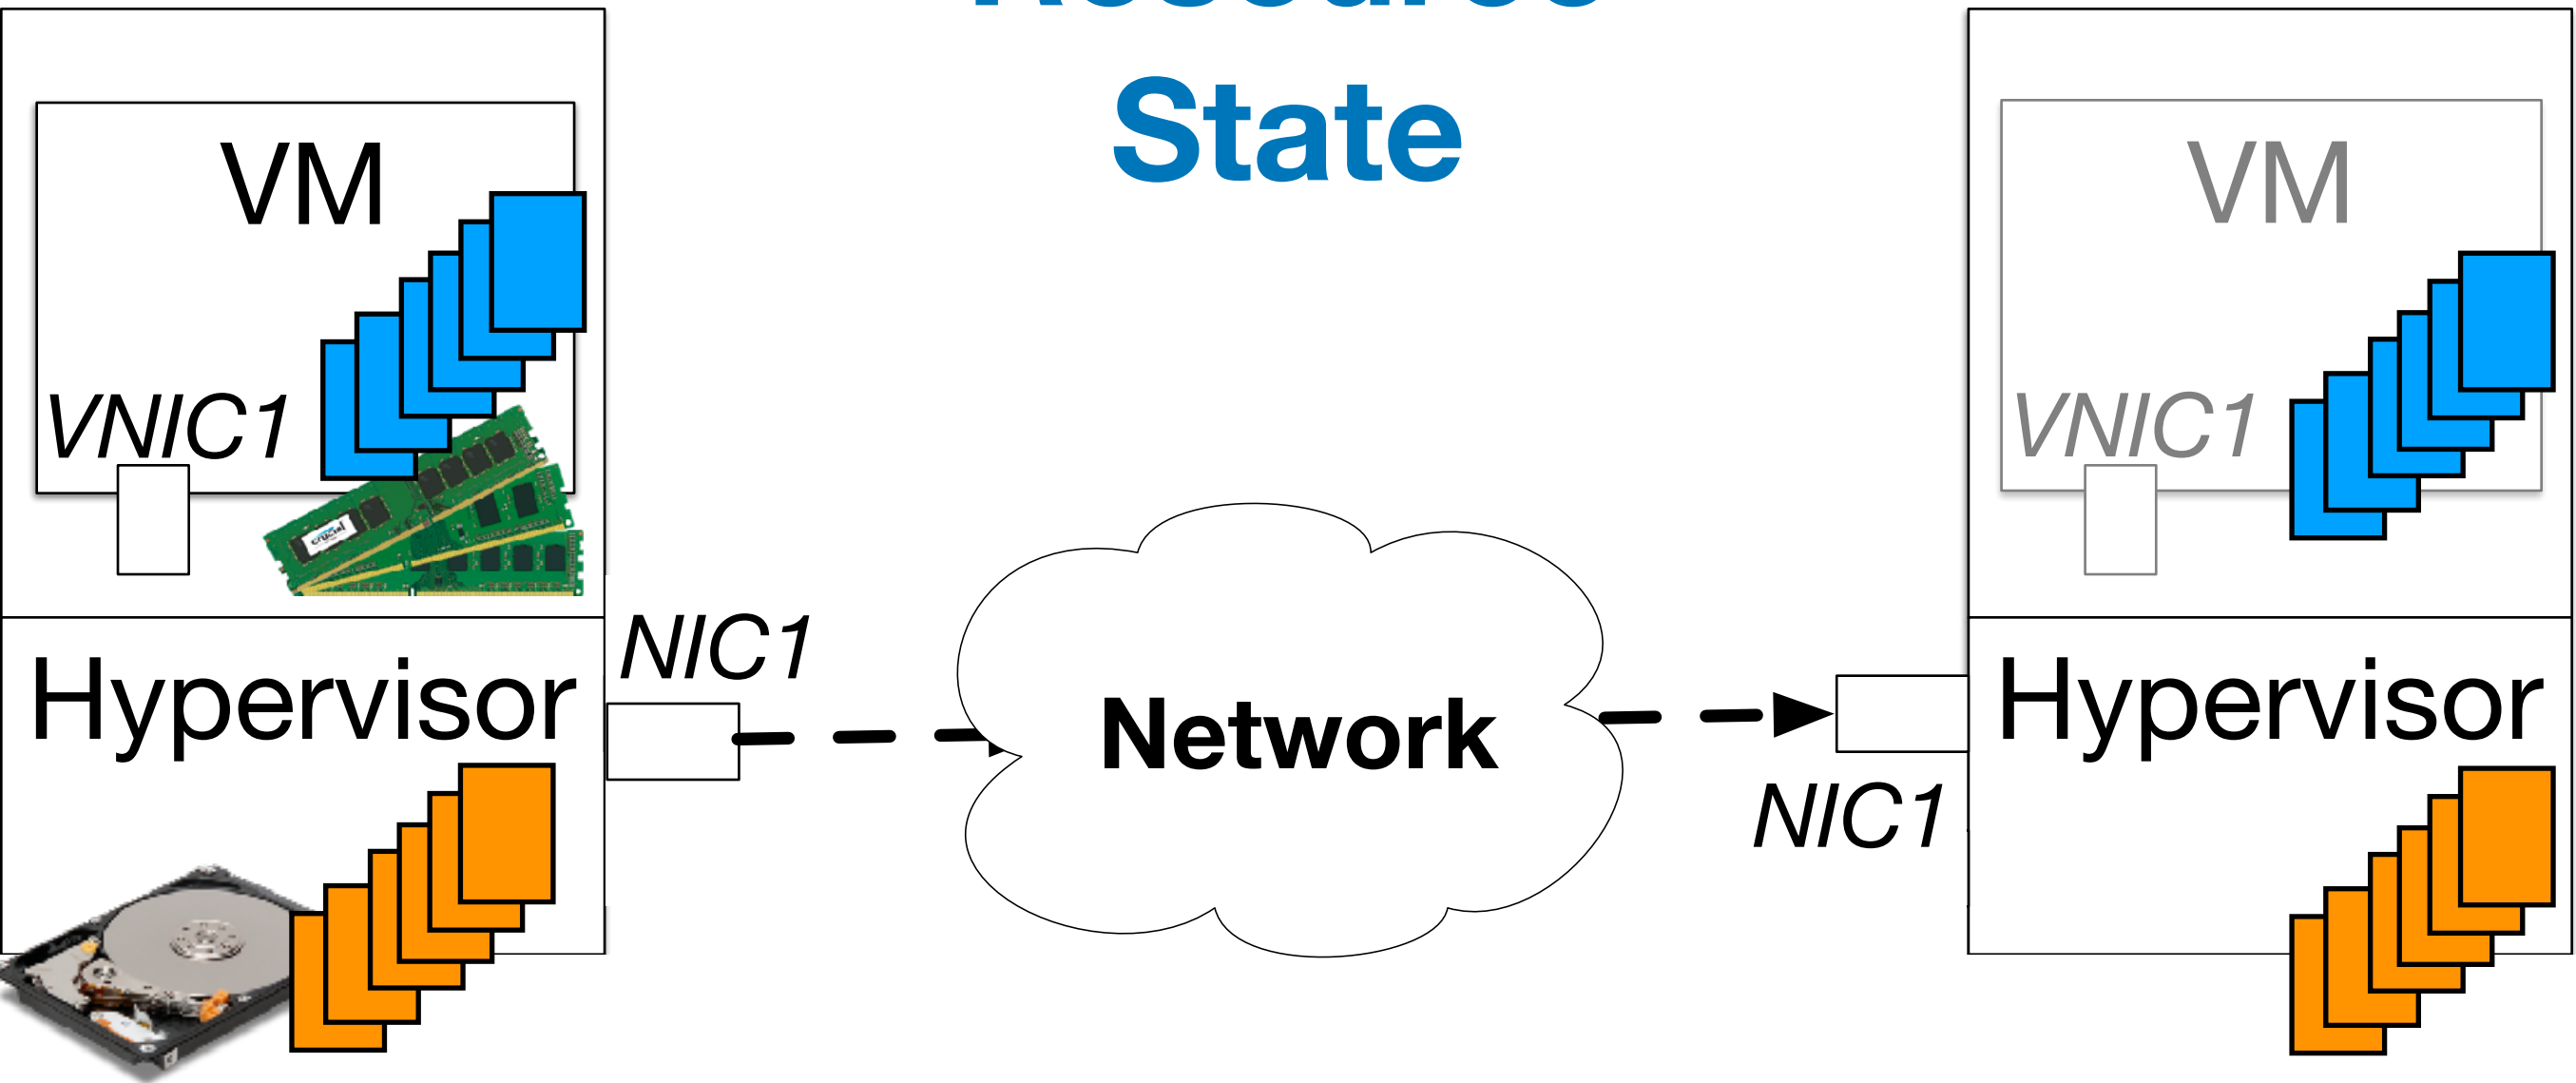


San Jose

San Mateo

Anatomy of VM Migration

Resource State

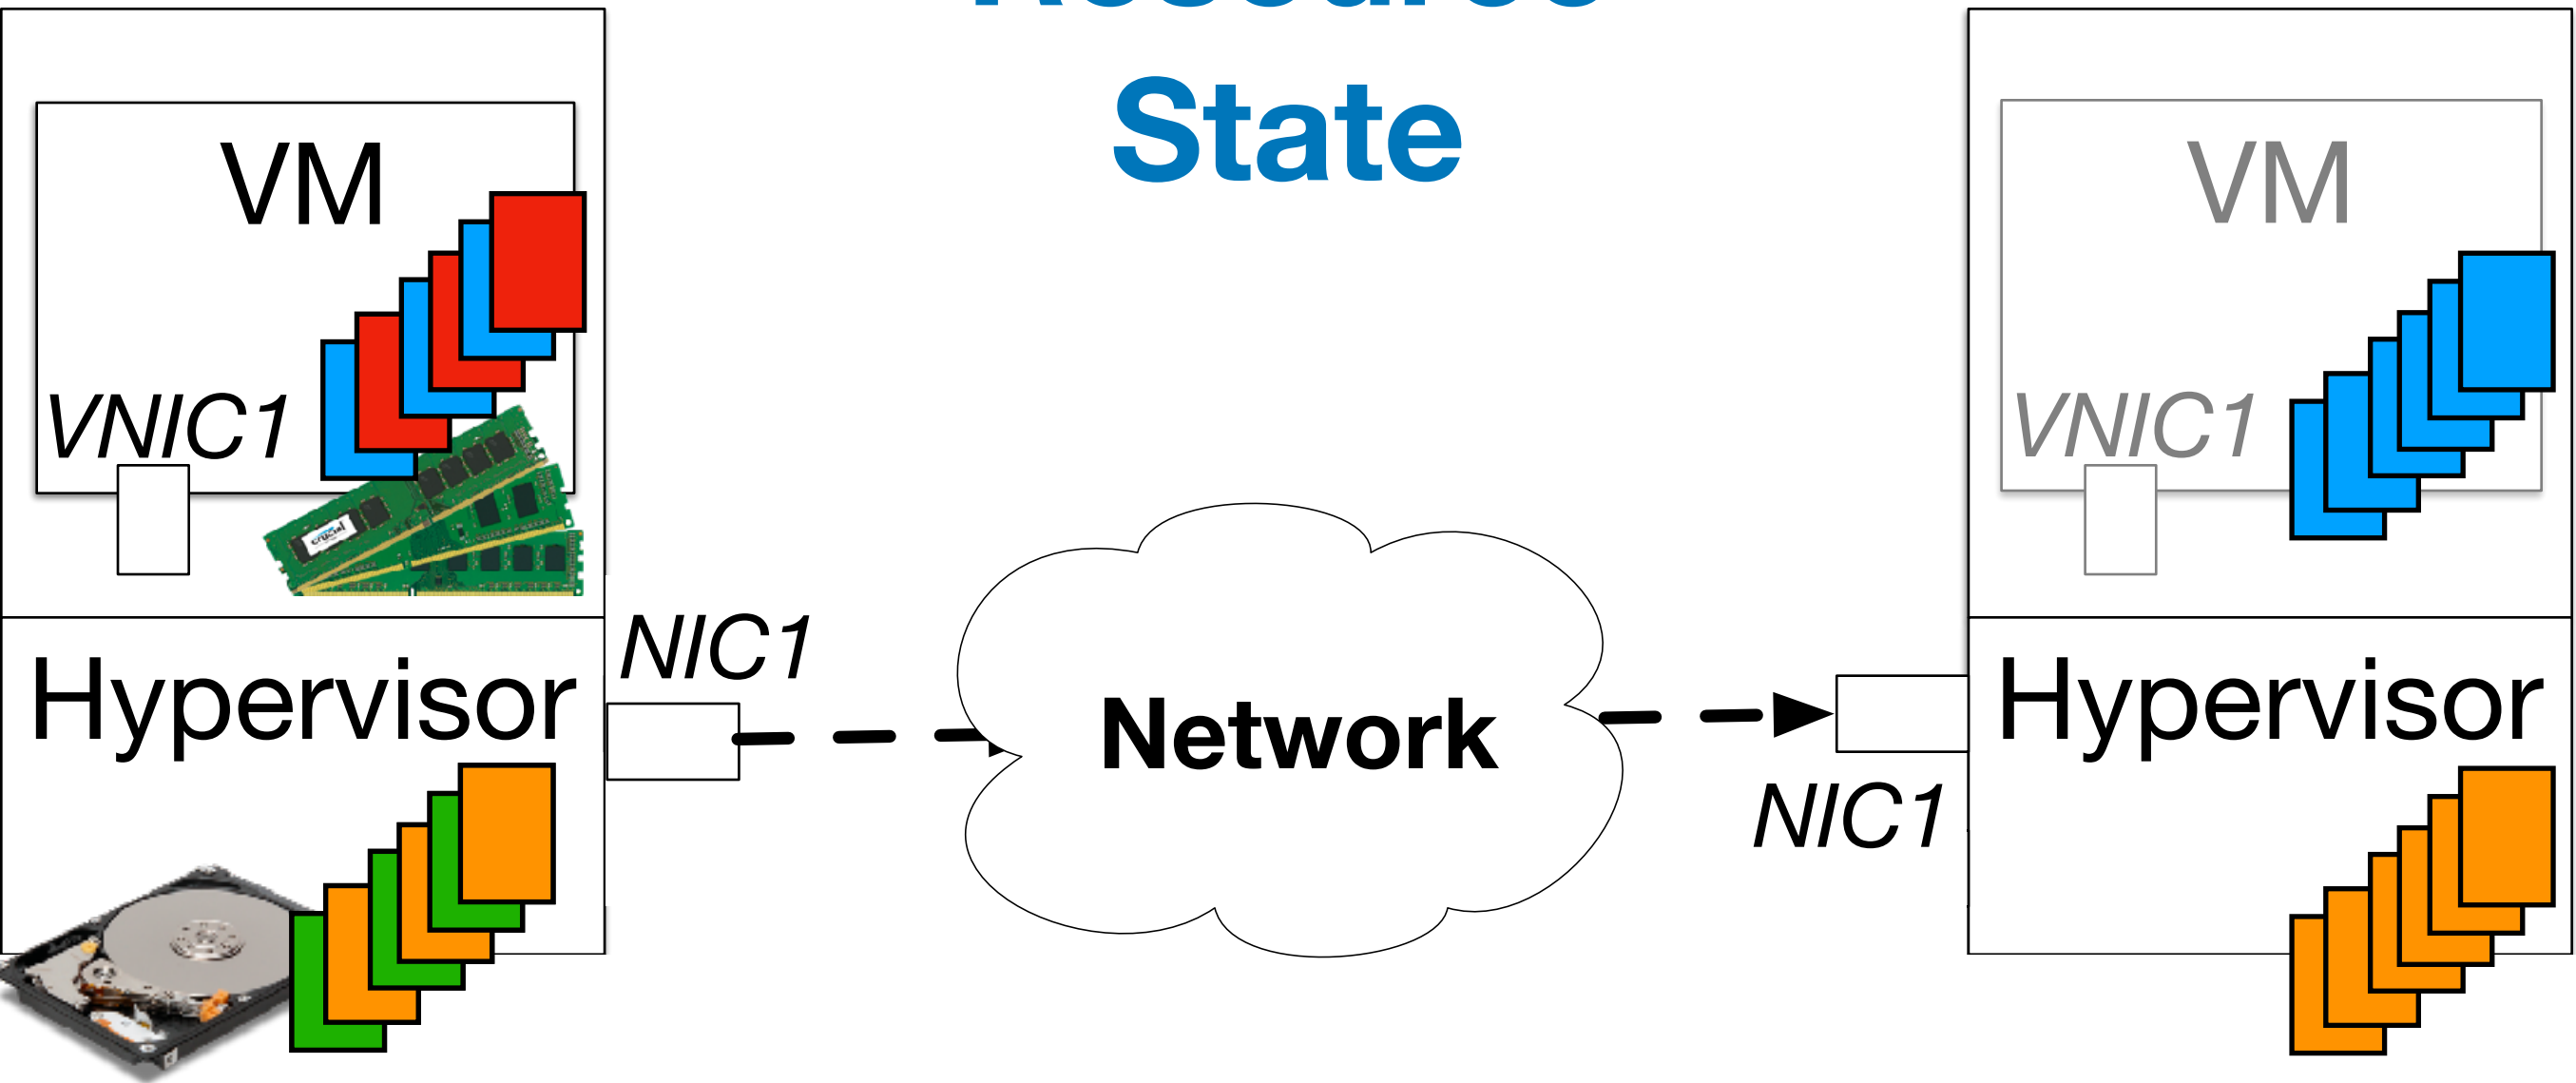


San Jose

San Mateo

Anatomy of VM Migration

Resource State

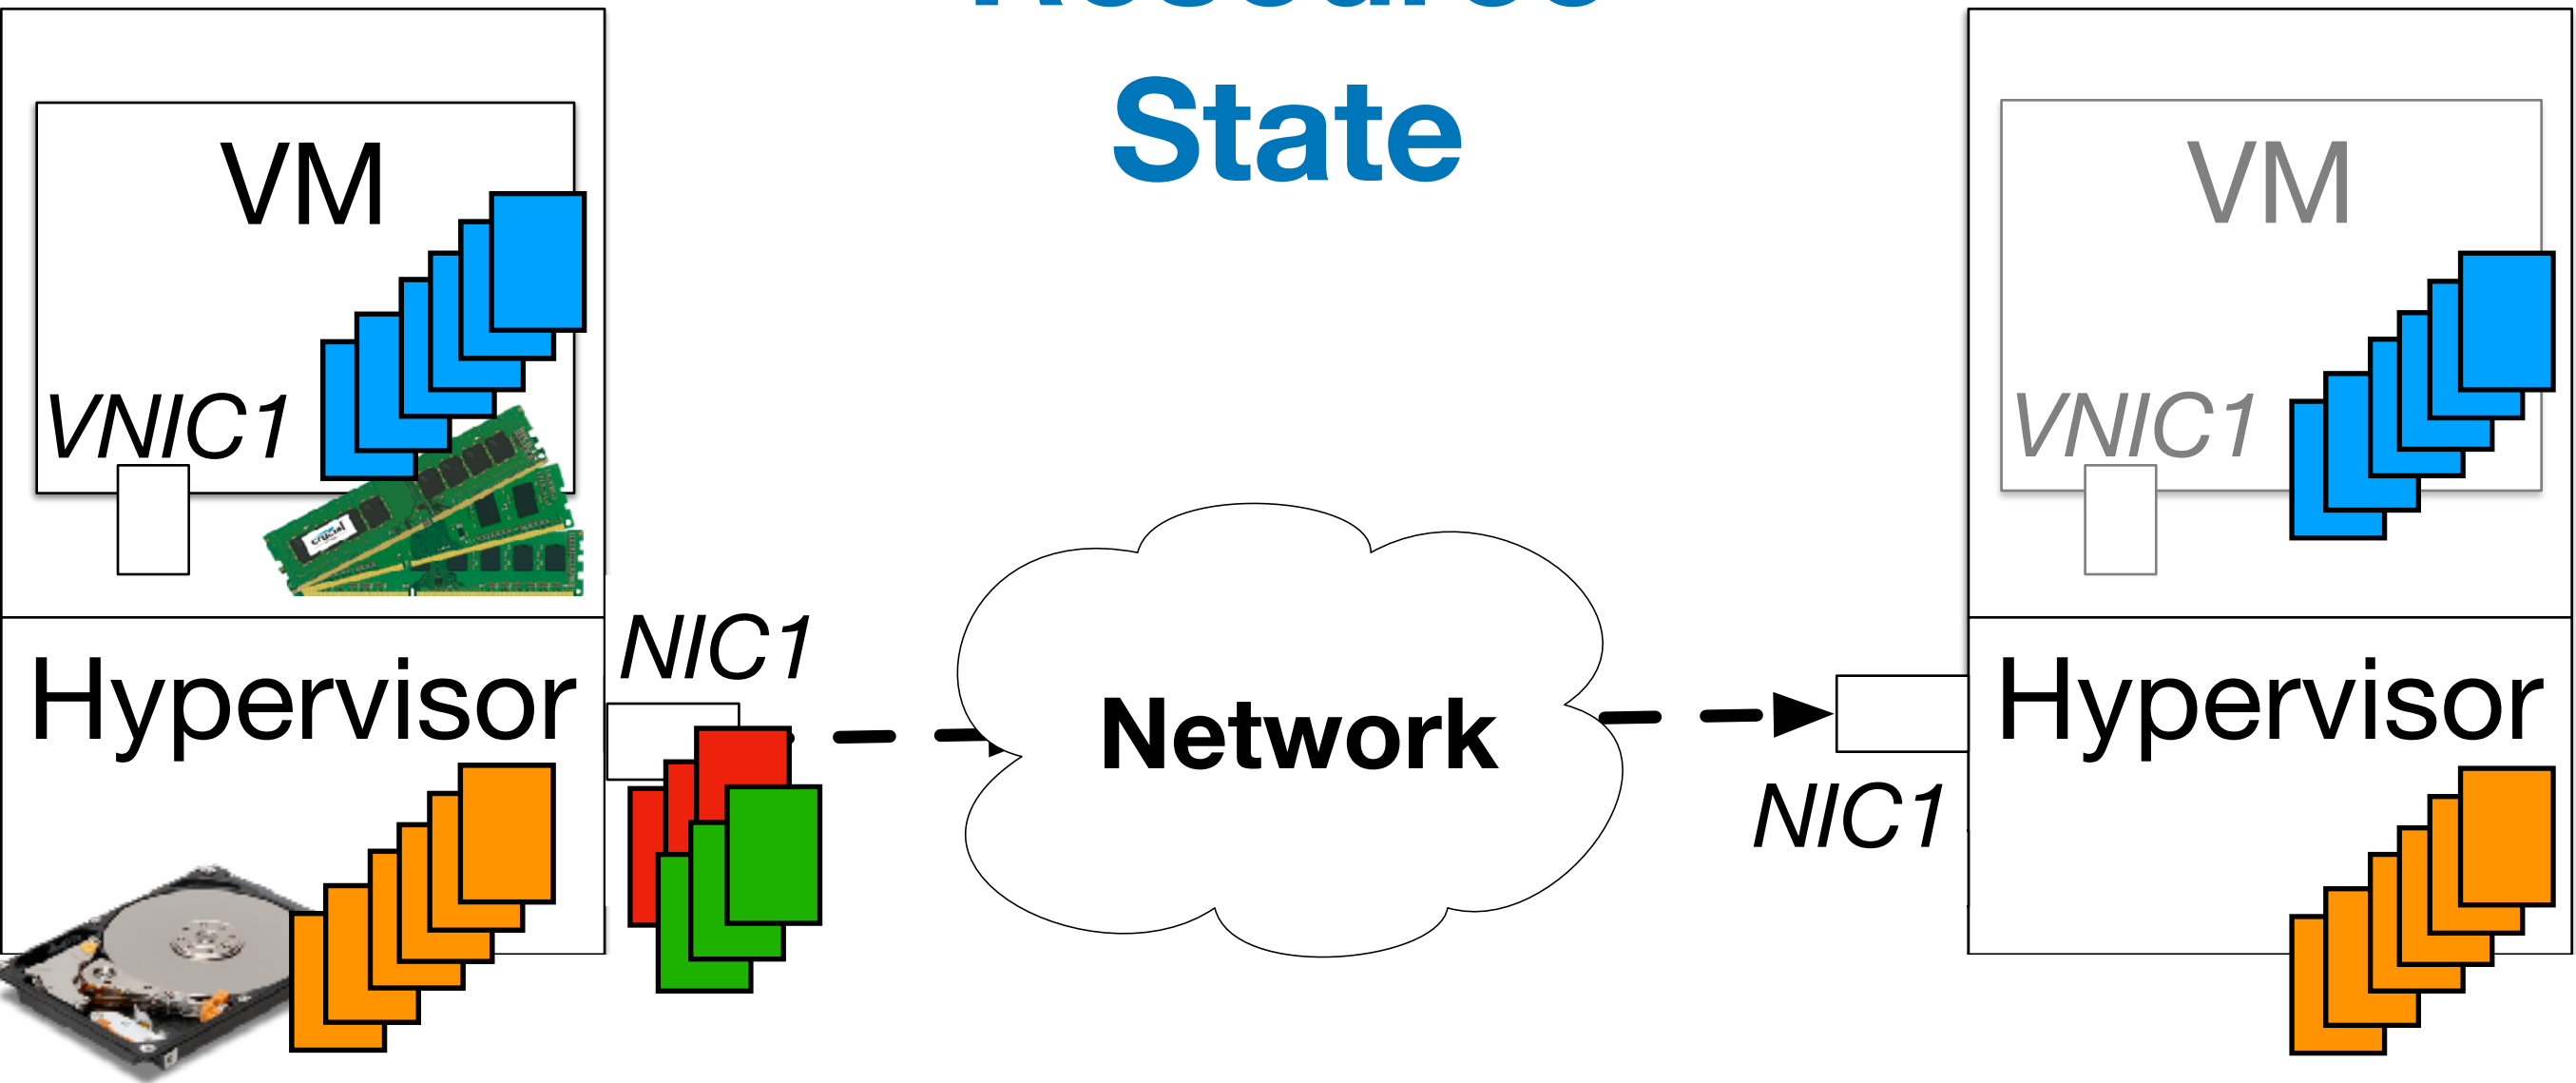


San Jose

San Mateo

Anatomy of VM Migration

Resource State

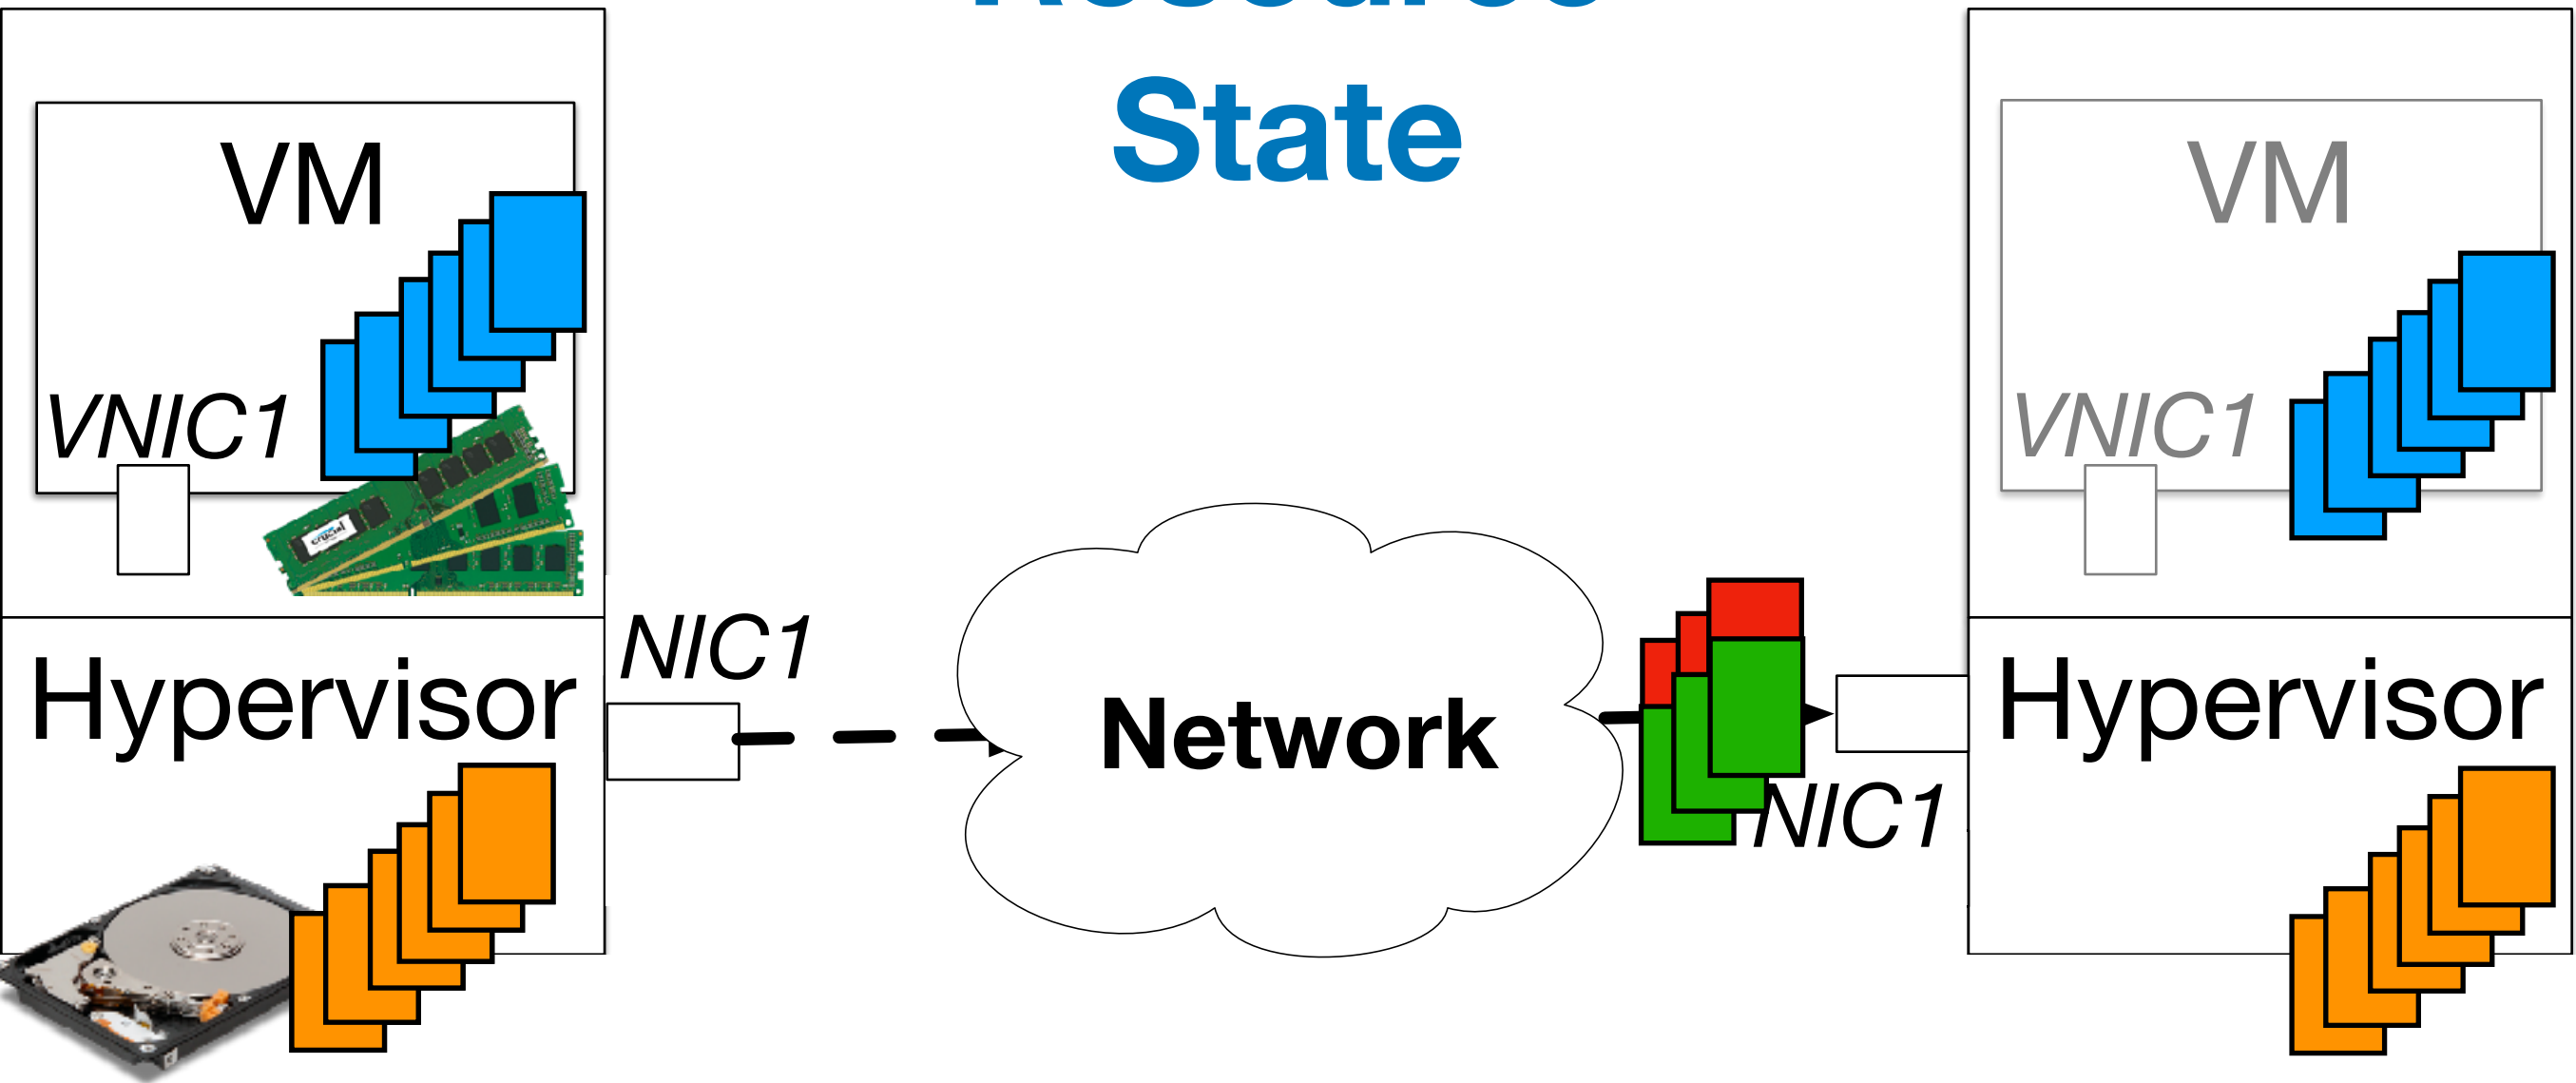


San Jose

San Mateo

Anatomy of VM Migration

Resource State

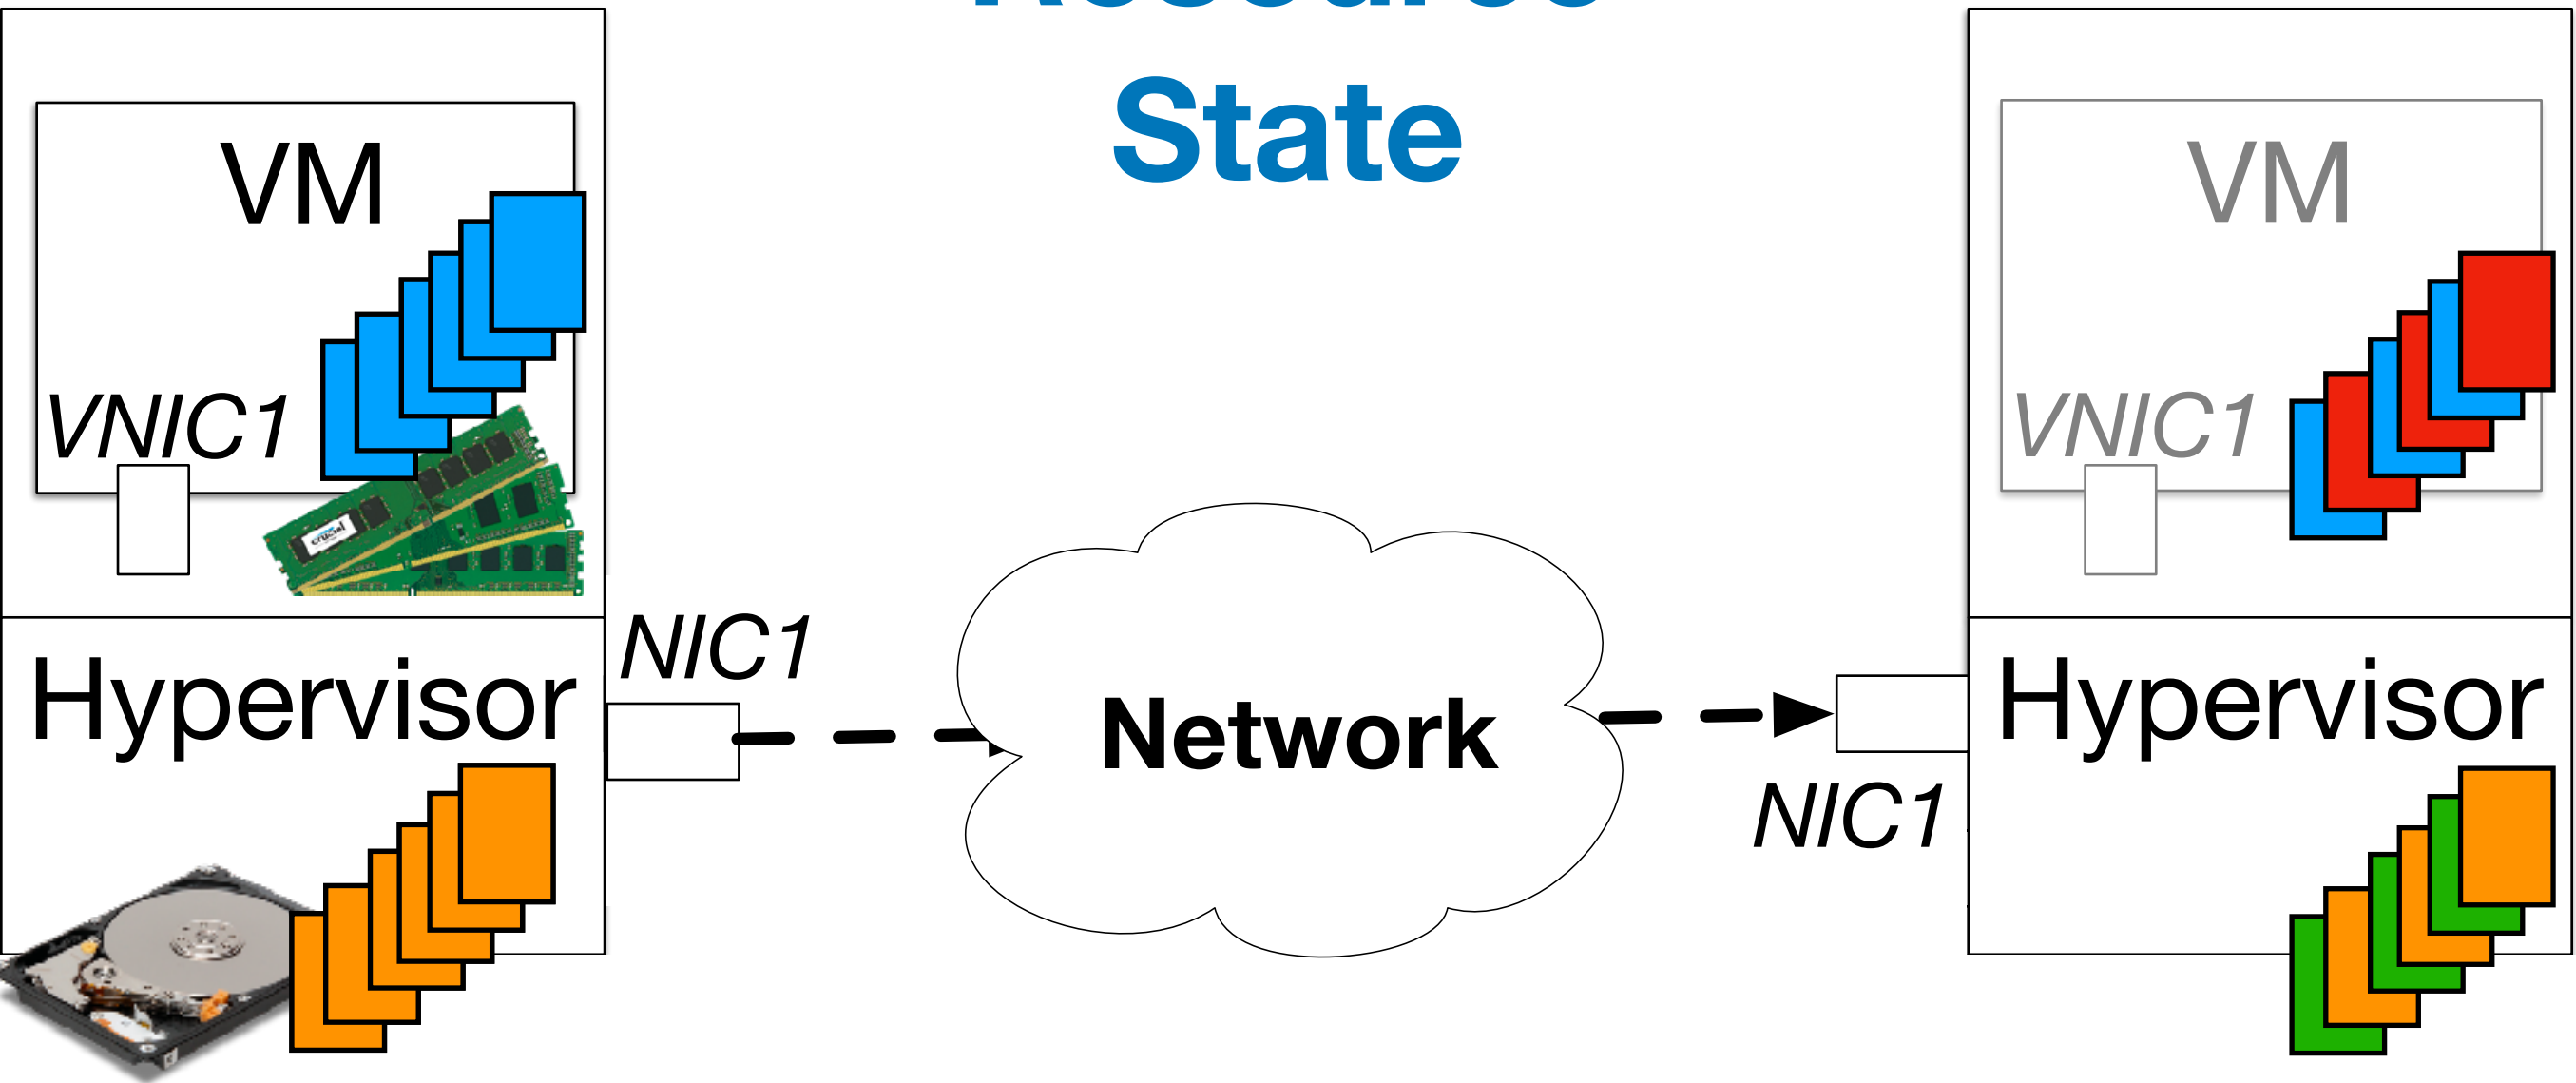


San Jose

San Mateo

Anatomy of VM Migration

Resource State

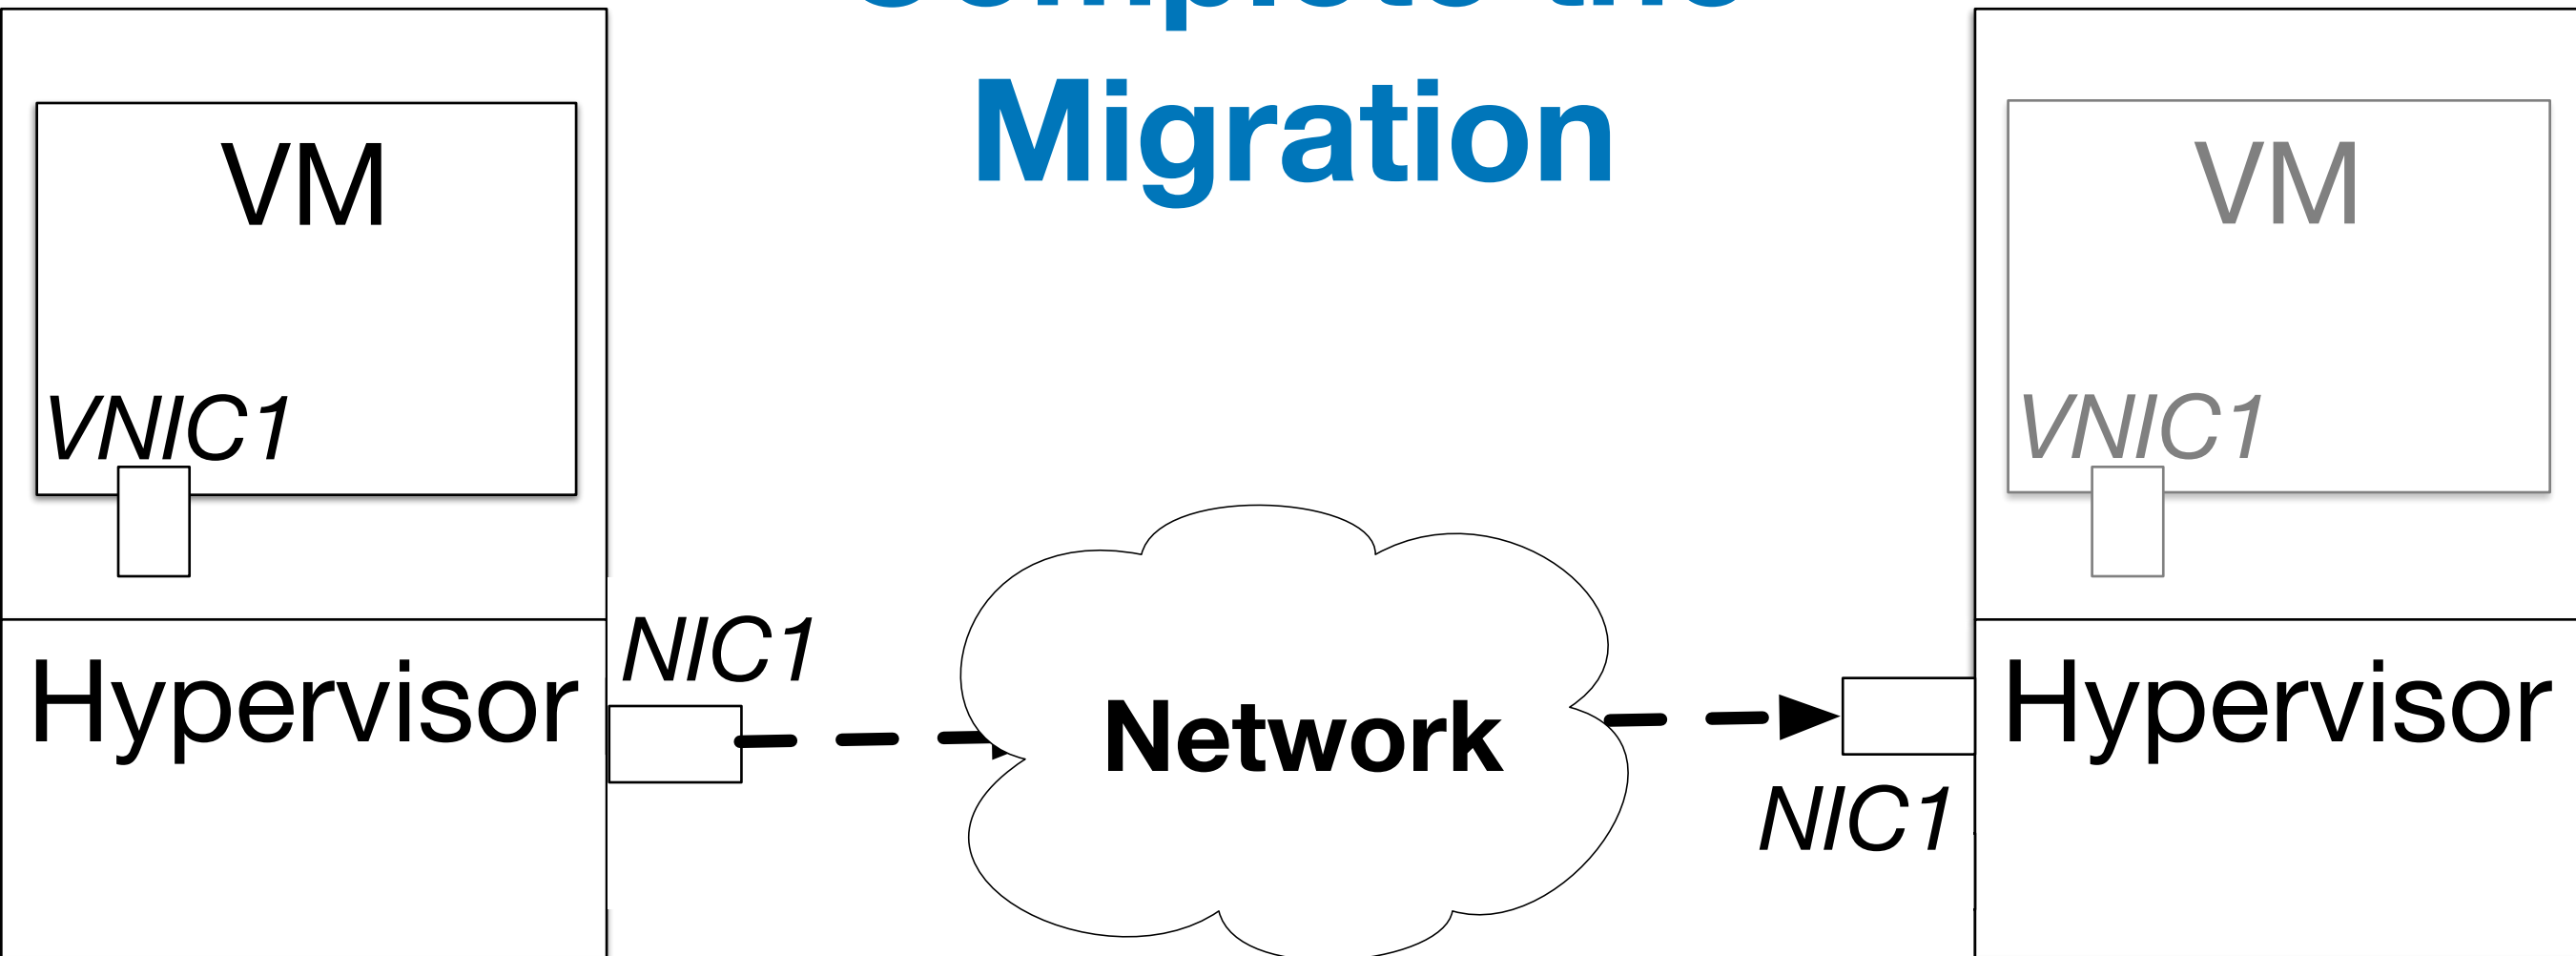


San Jose

San Mateo

Anatomy of VM Migration

Complete the Migration

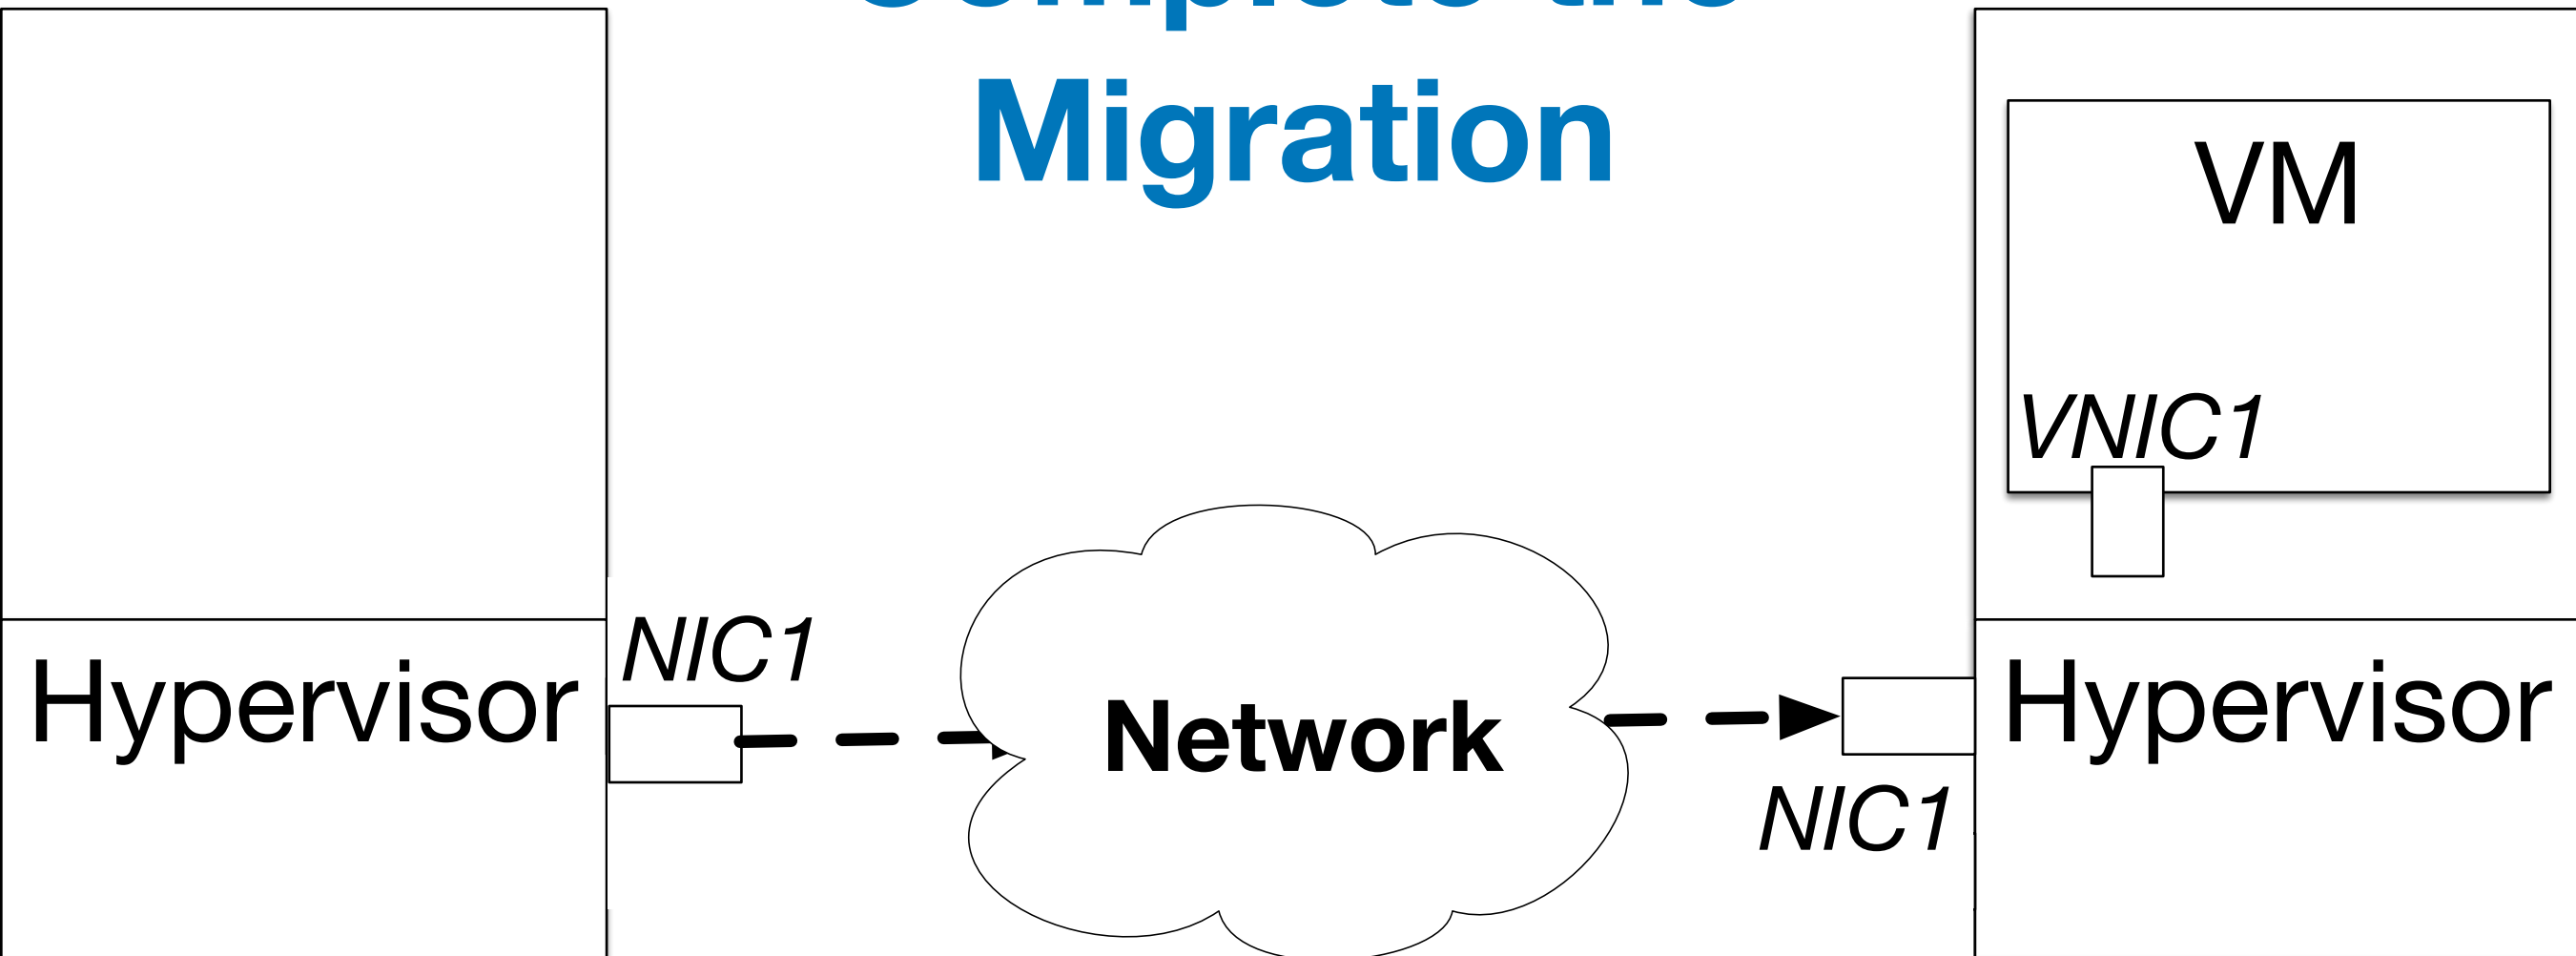


San Jose

San Mateo

Anatomy of VM Migration

Complete the Migration




San Jose

San Mateo

Anatomy of VM Migration

Disk & Memory State



**Low Bandwidth Speeds
Mean Long Transfer
Times.**

San Jose

San Mateo

Migration over Slow Networks

- **Problem:** Low Bandwidth Speeds Mean Long Transfer Times.


Migration over Slow Networks

- **Problem:** Low Bandwidth Speeds Mean Long Transfer Times.
- **Observation:** Most commodity servers have multiple interfaces and multiple paths are common between edge cloud locations


Migration over Slow Networks

- **Problem:** Low Bandwidth Speeds Mean Long Transfer Times.
- **Observation:** Most commodity servers have multiple interfaces and multiple paths are common between edge cloud locations
- **Solution:** Take advantage of multiple network paths to aggregate bandwidth.

A Typical TCP Connection




A Typical TCP Connection




What if we could take advantage of our second interface?

An MPTCP Socket




Migrating over MPTCP

- Parallel Paths allow us to aggregate bandwidth
- A net effect of providing faster transfer times for VM state.



Migration

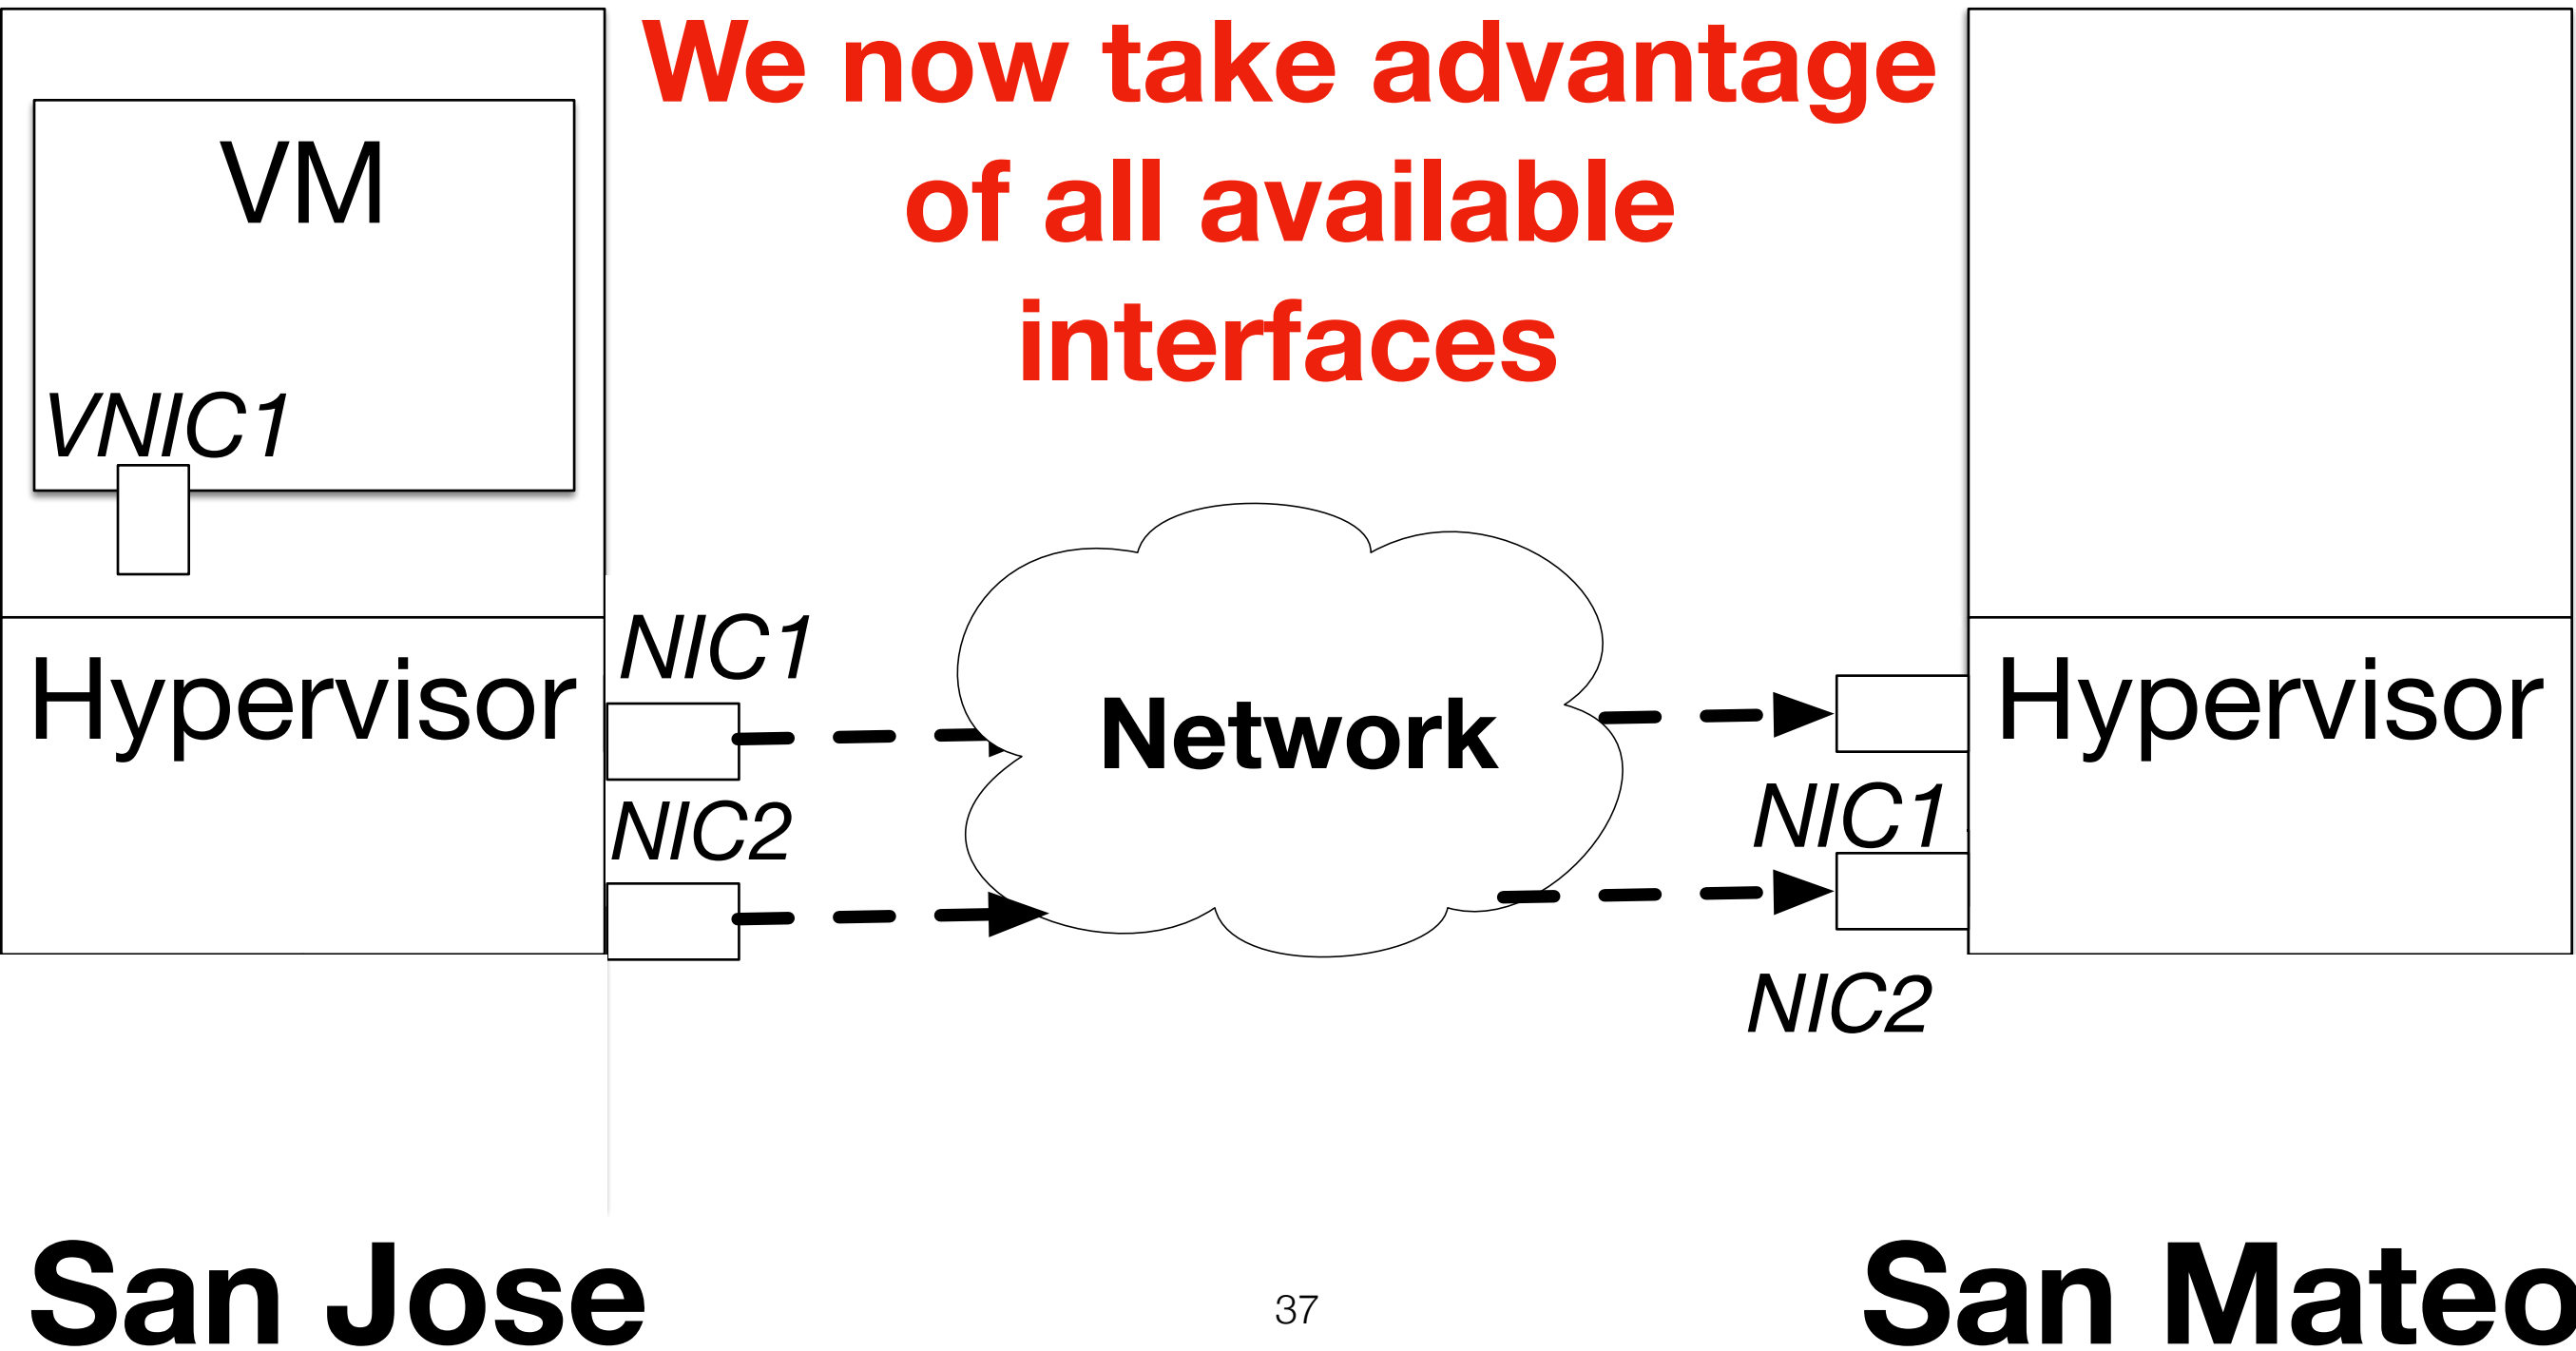


San Jose

San Mateo

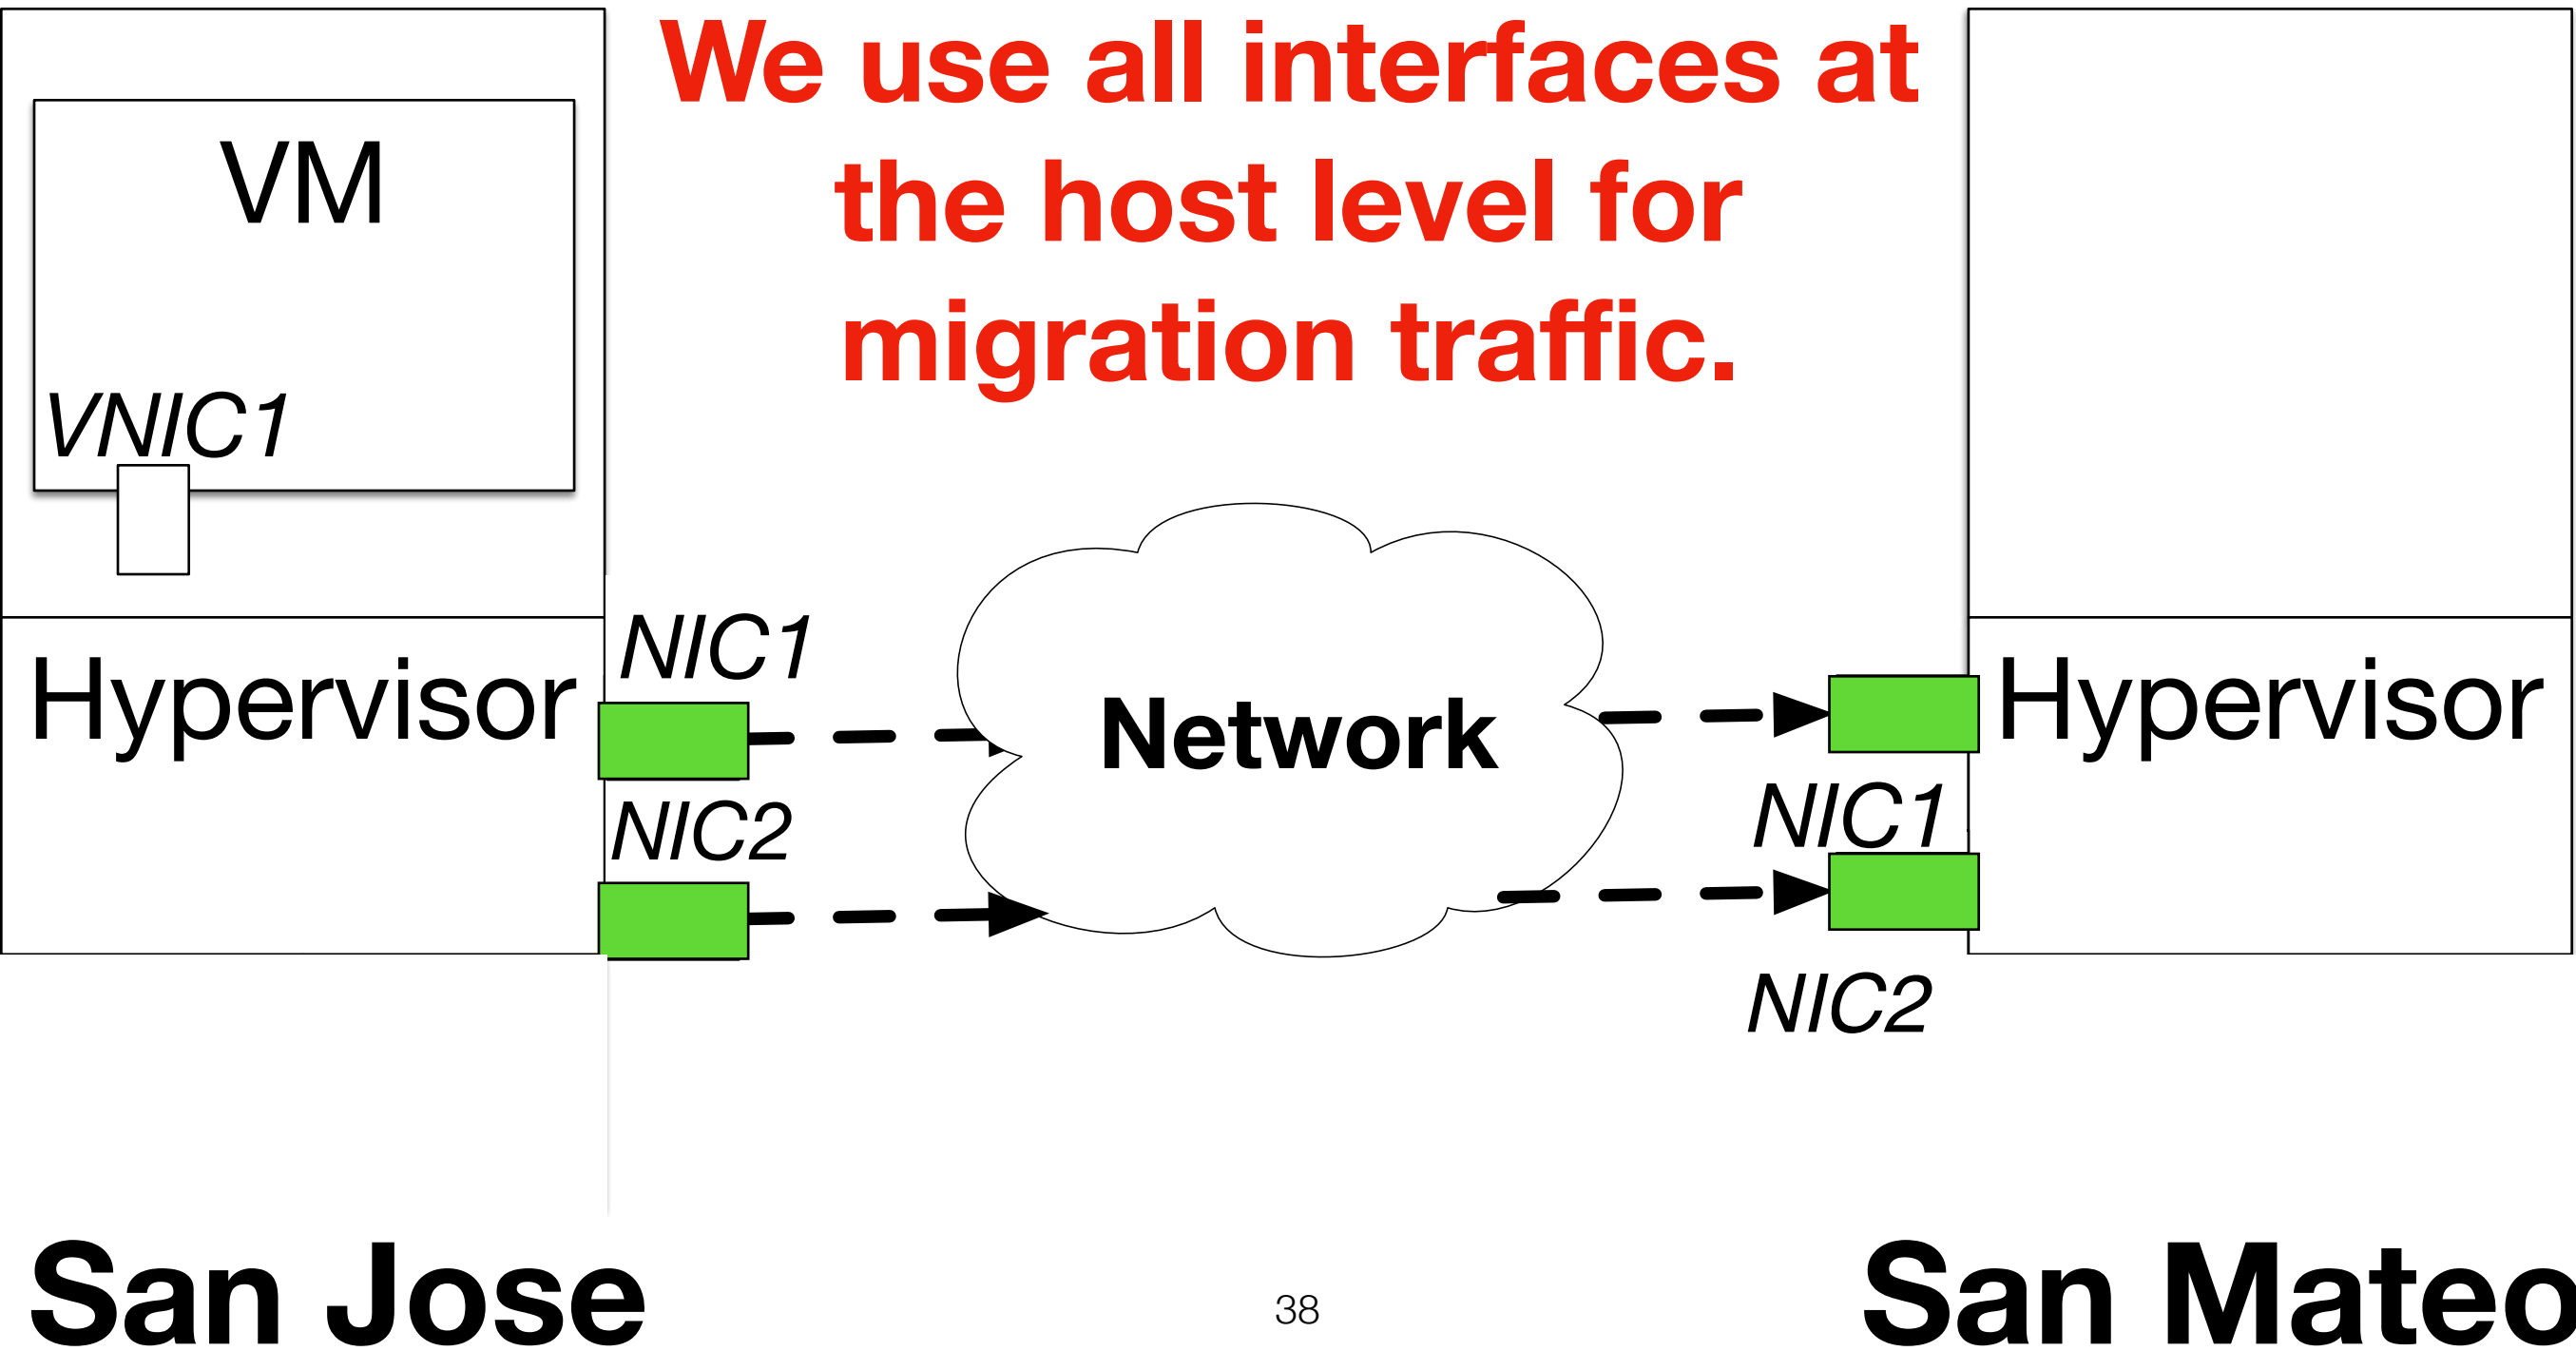
Migration

**We now take advantage
of all available
interfaces**



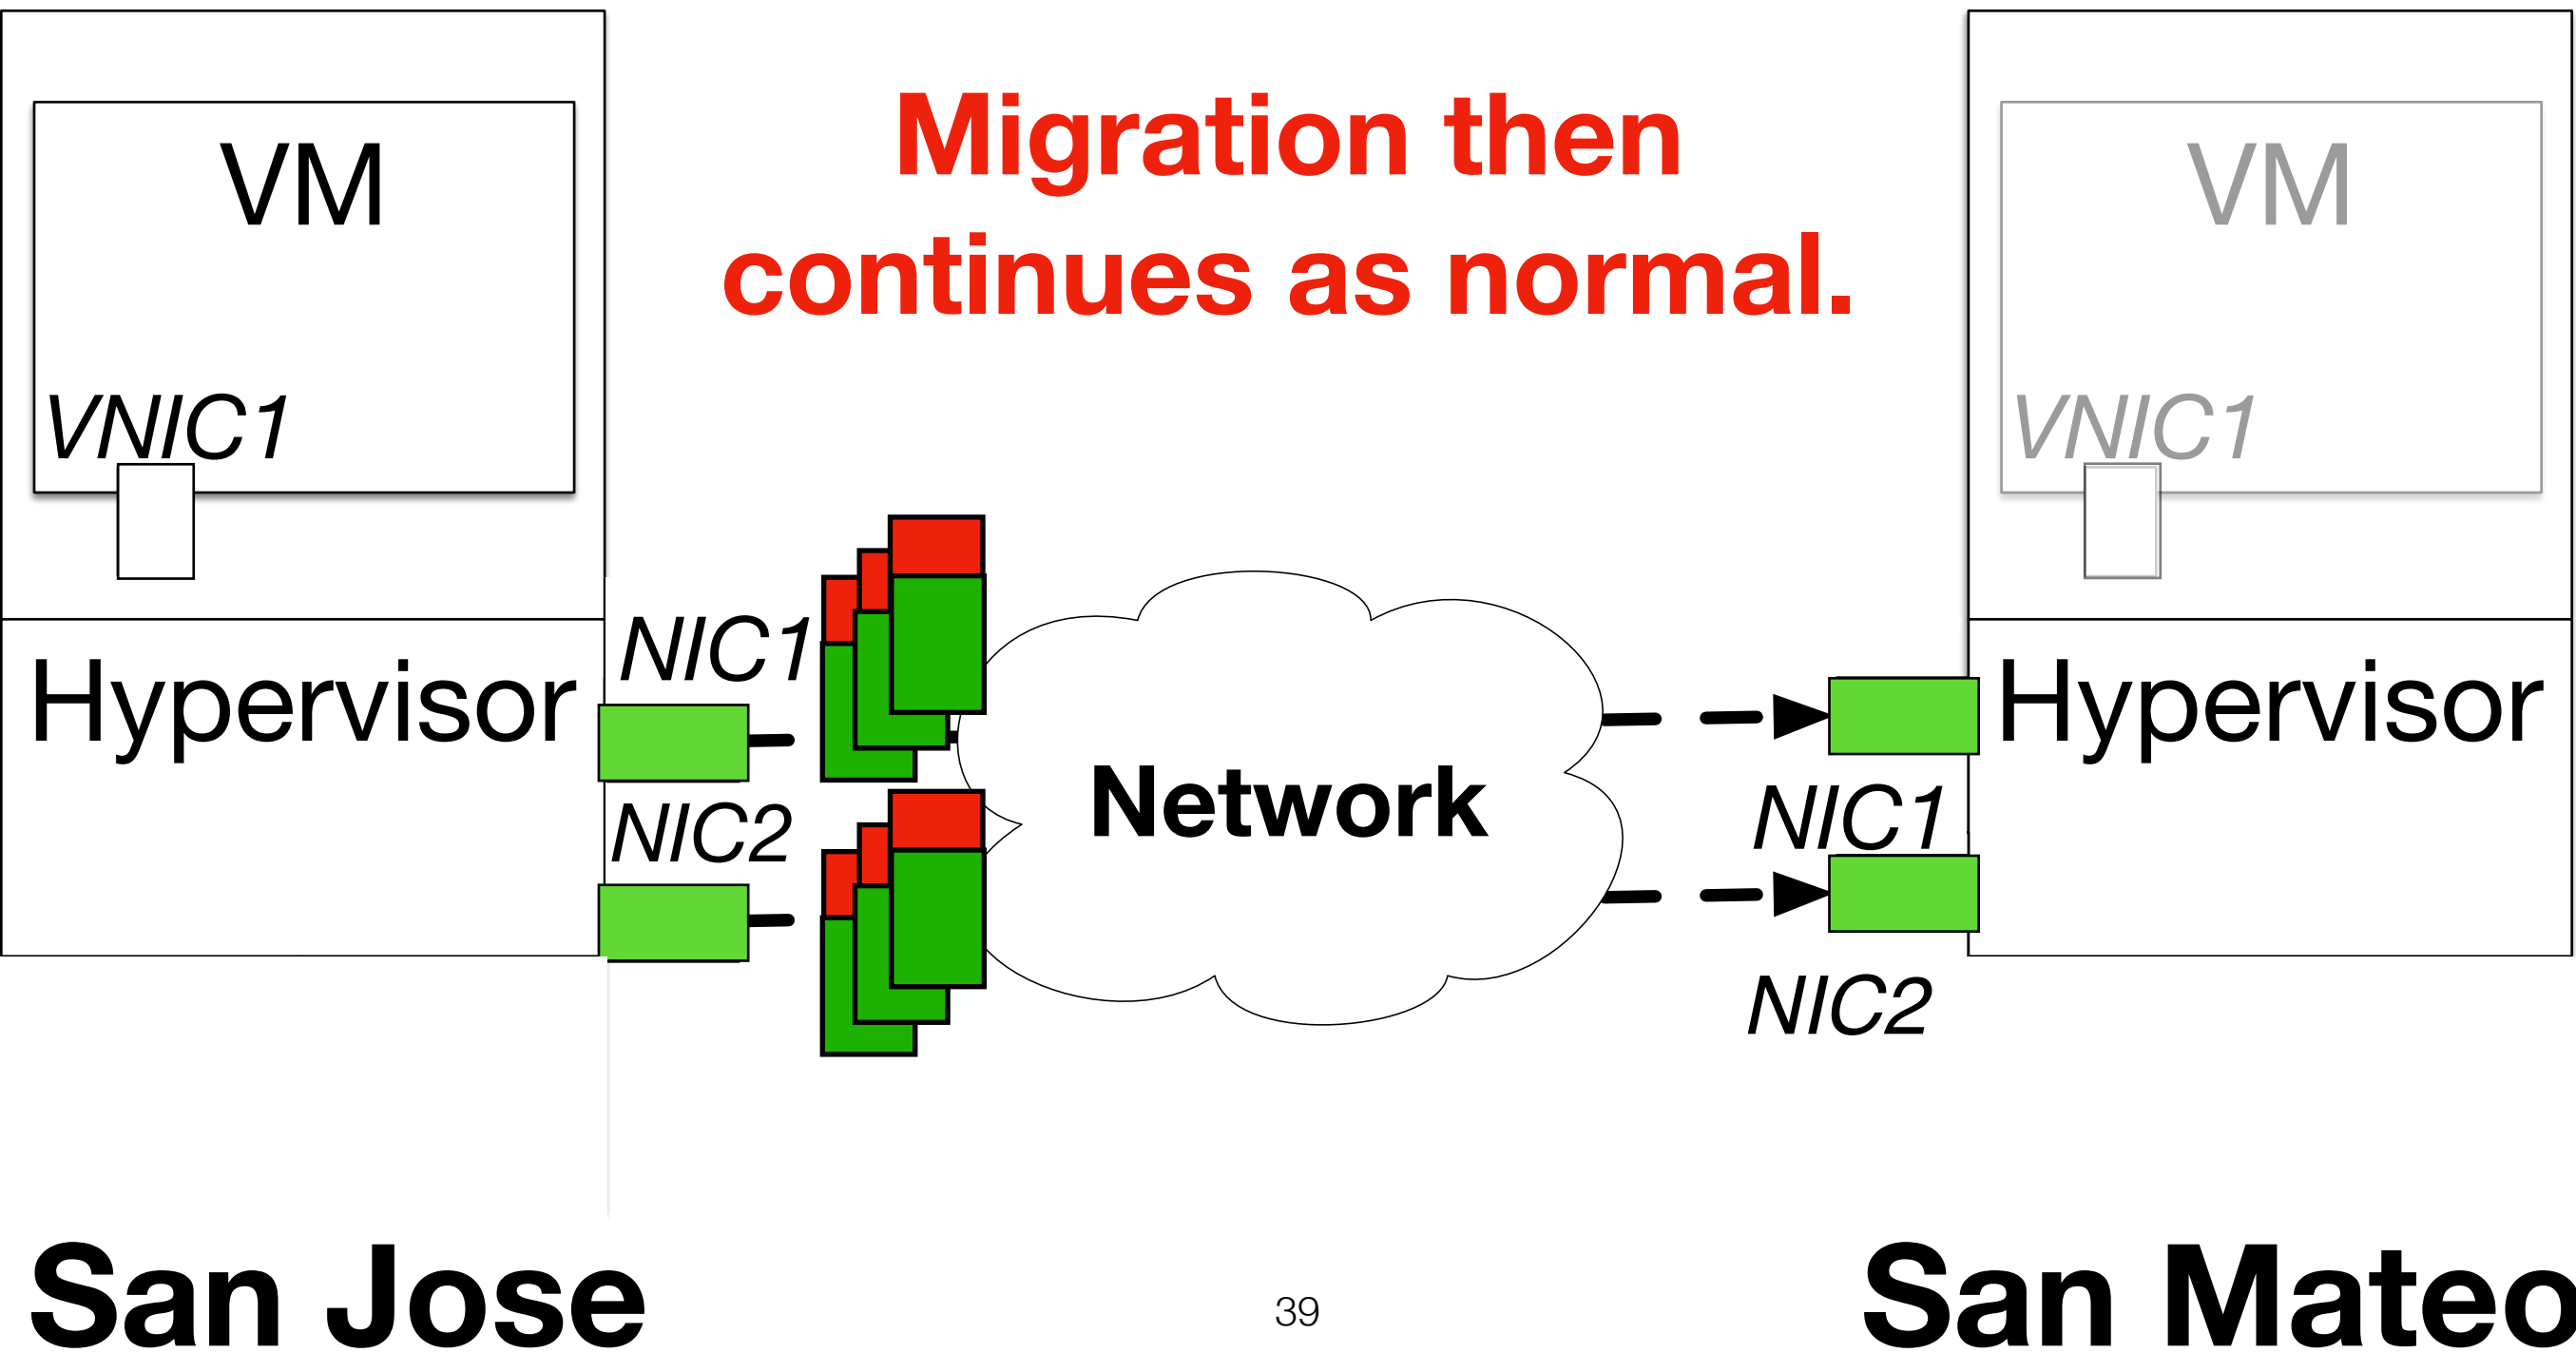
Migration

We use all interfaces at the host level for migration traffic.



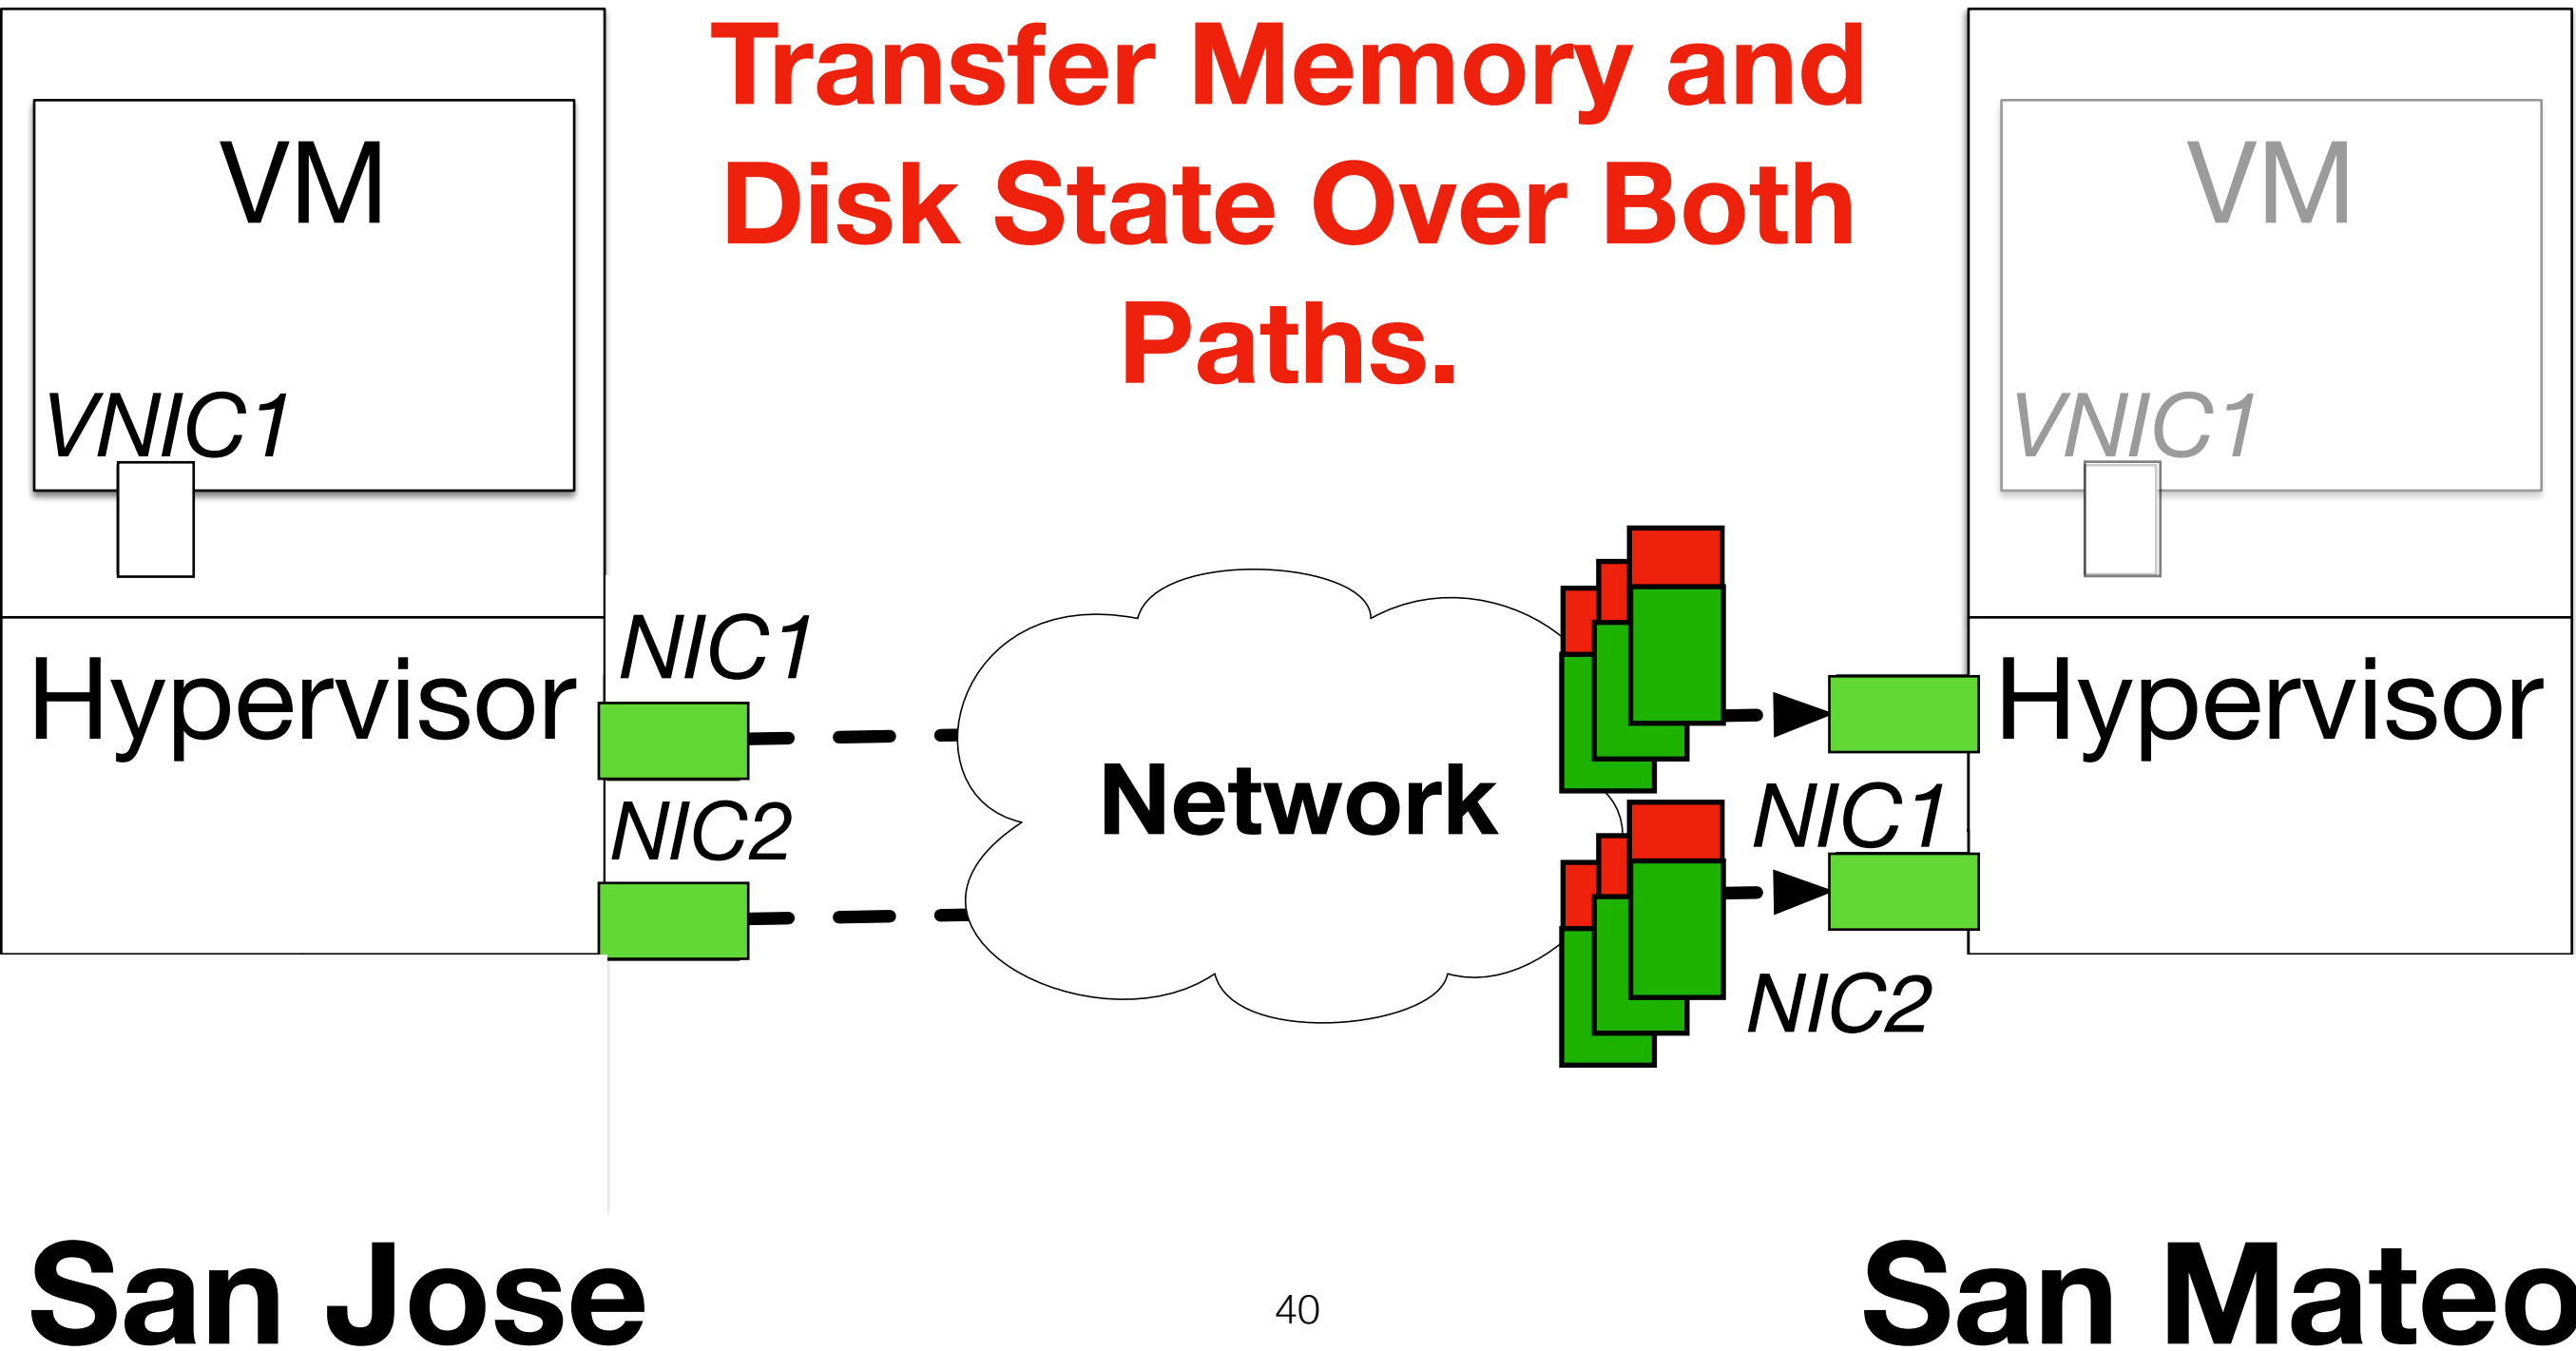
Migration

**Migration then
continues as normal.**



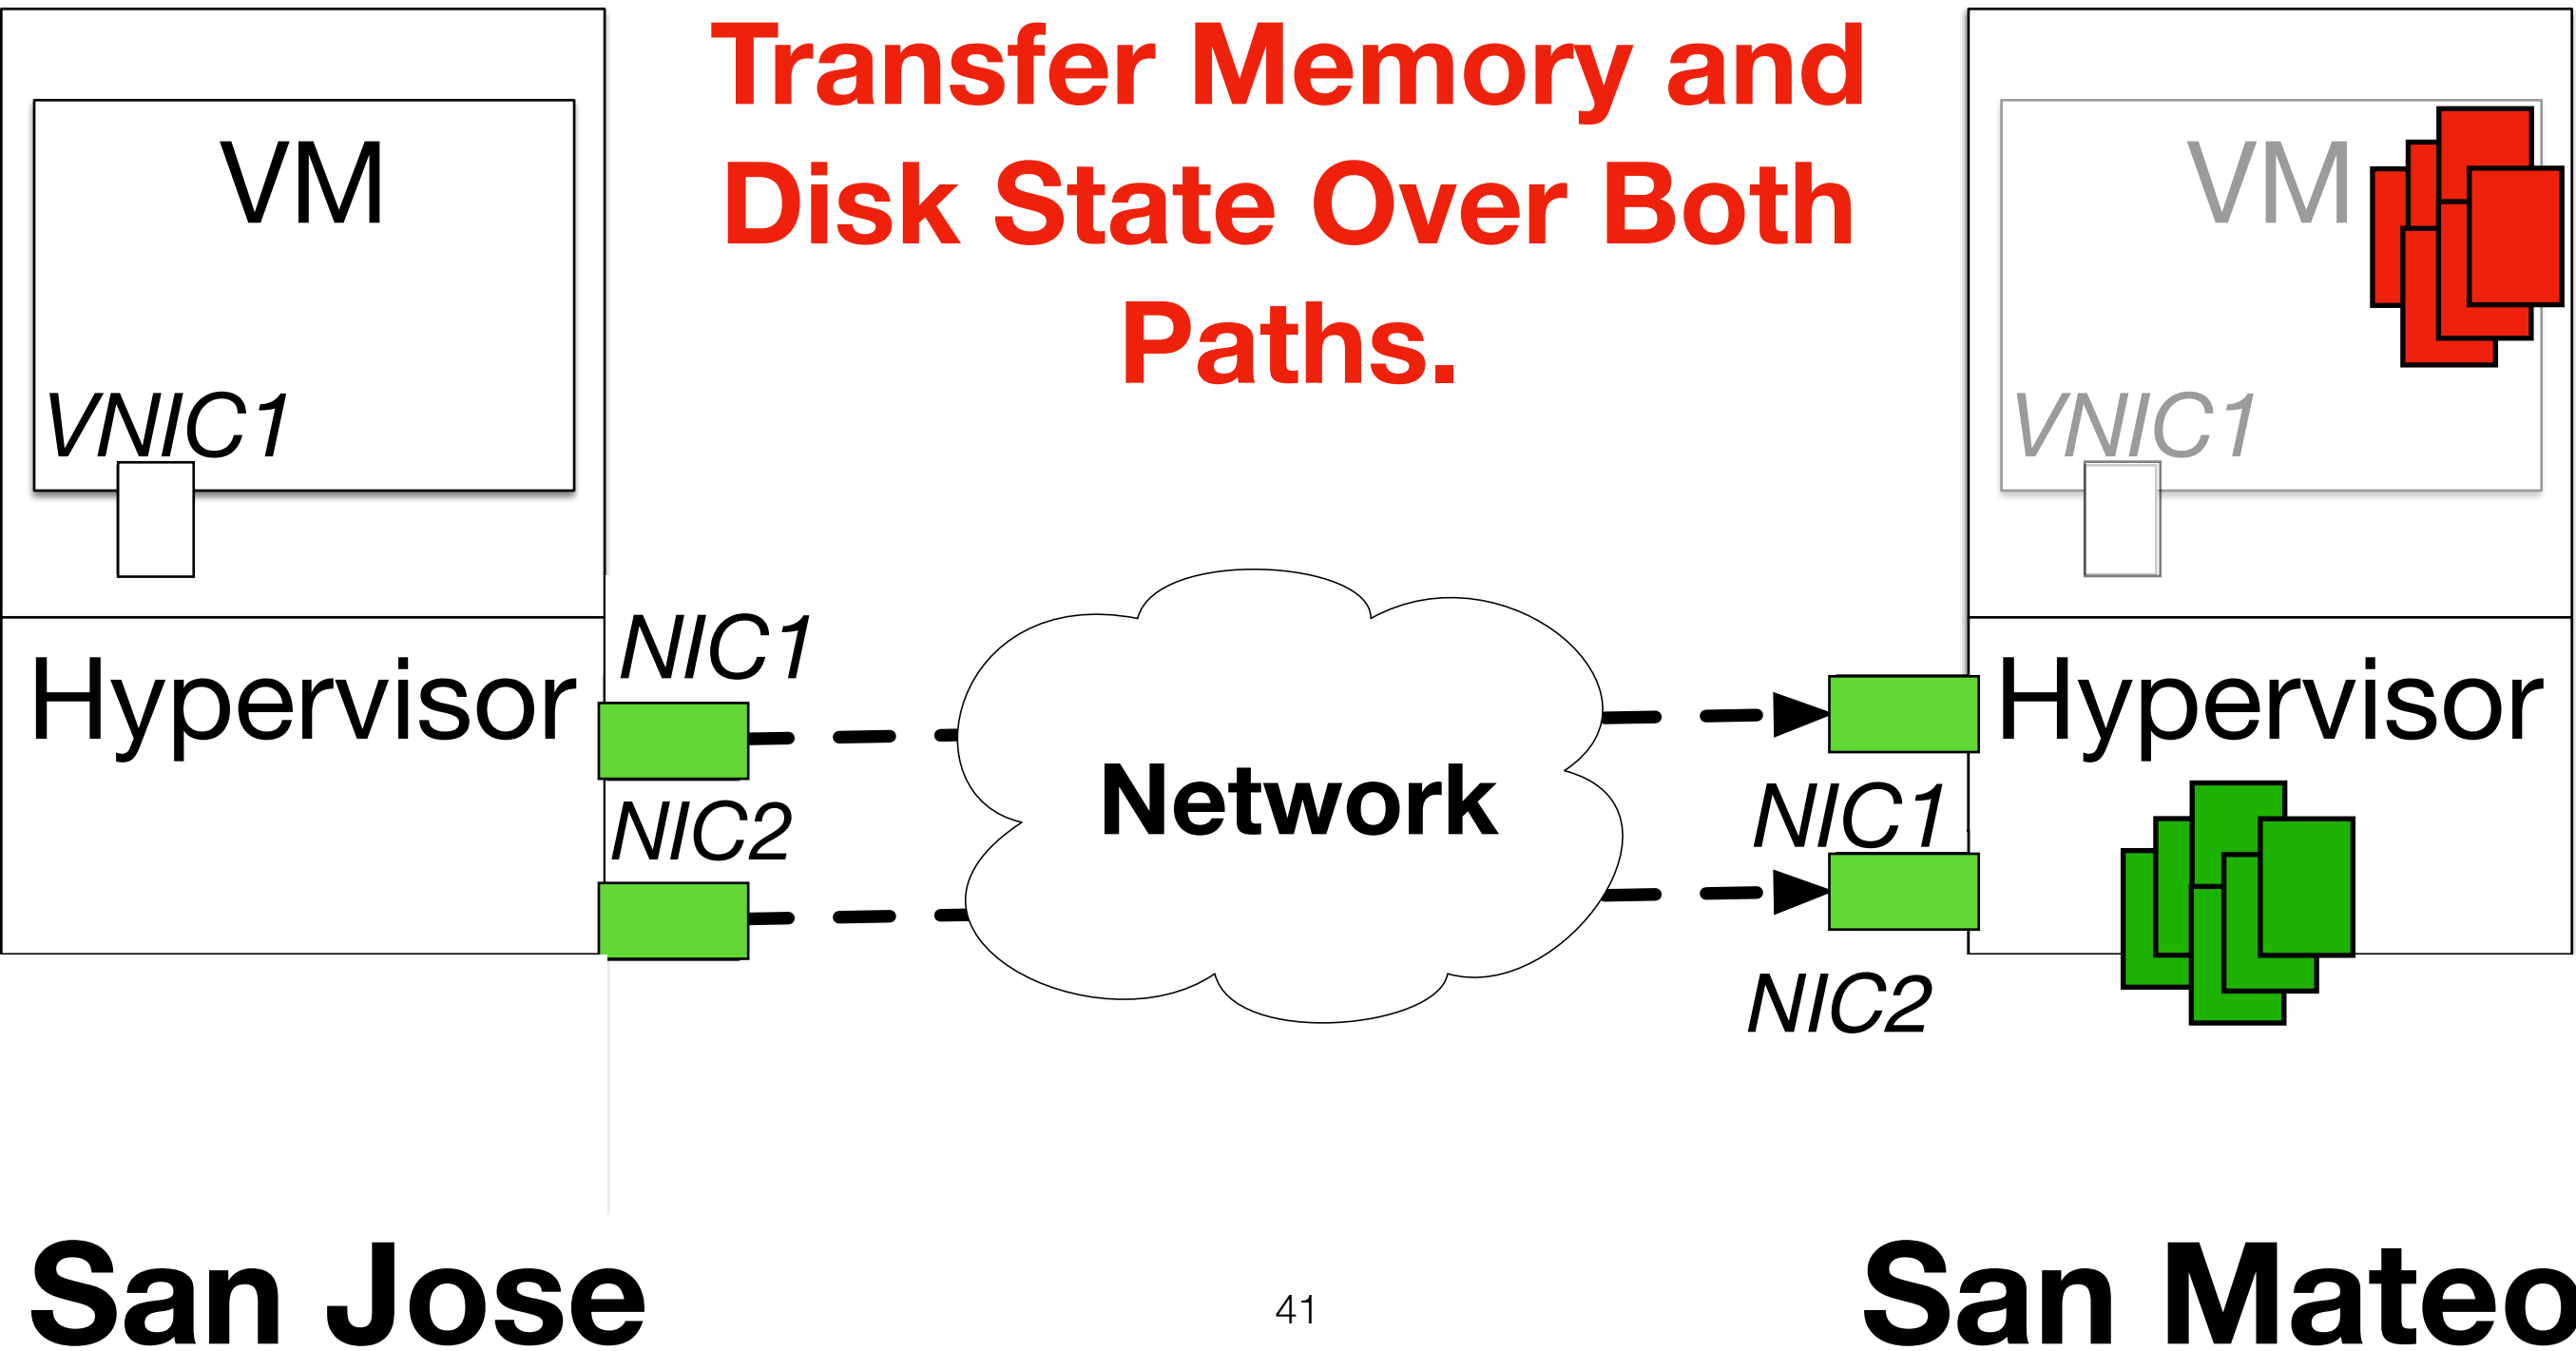
Migration

**Transfer Memory and
Disk State Over Both
Paths.**



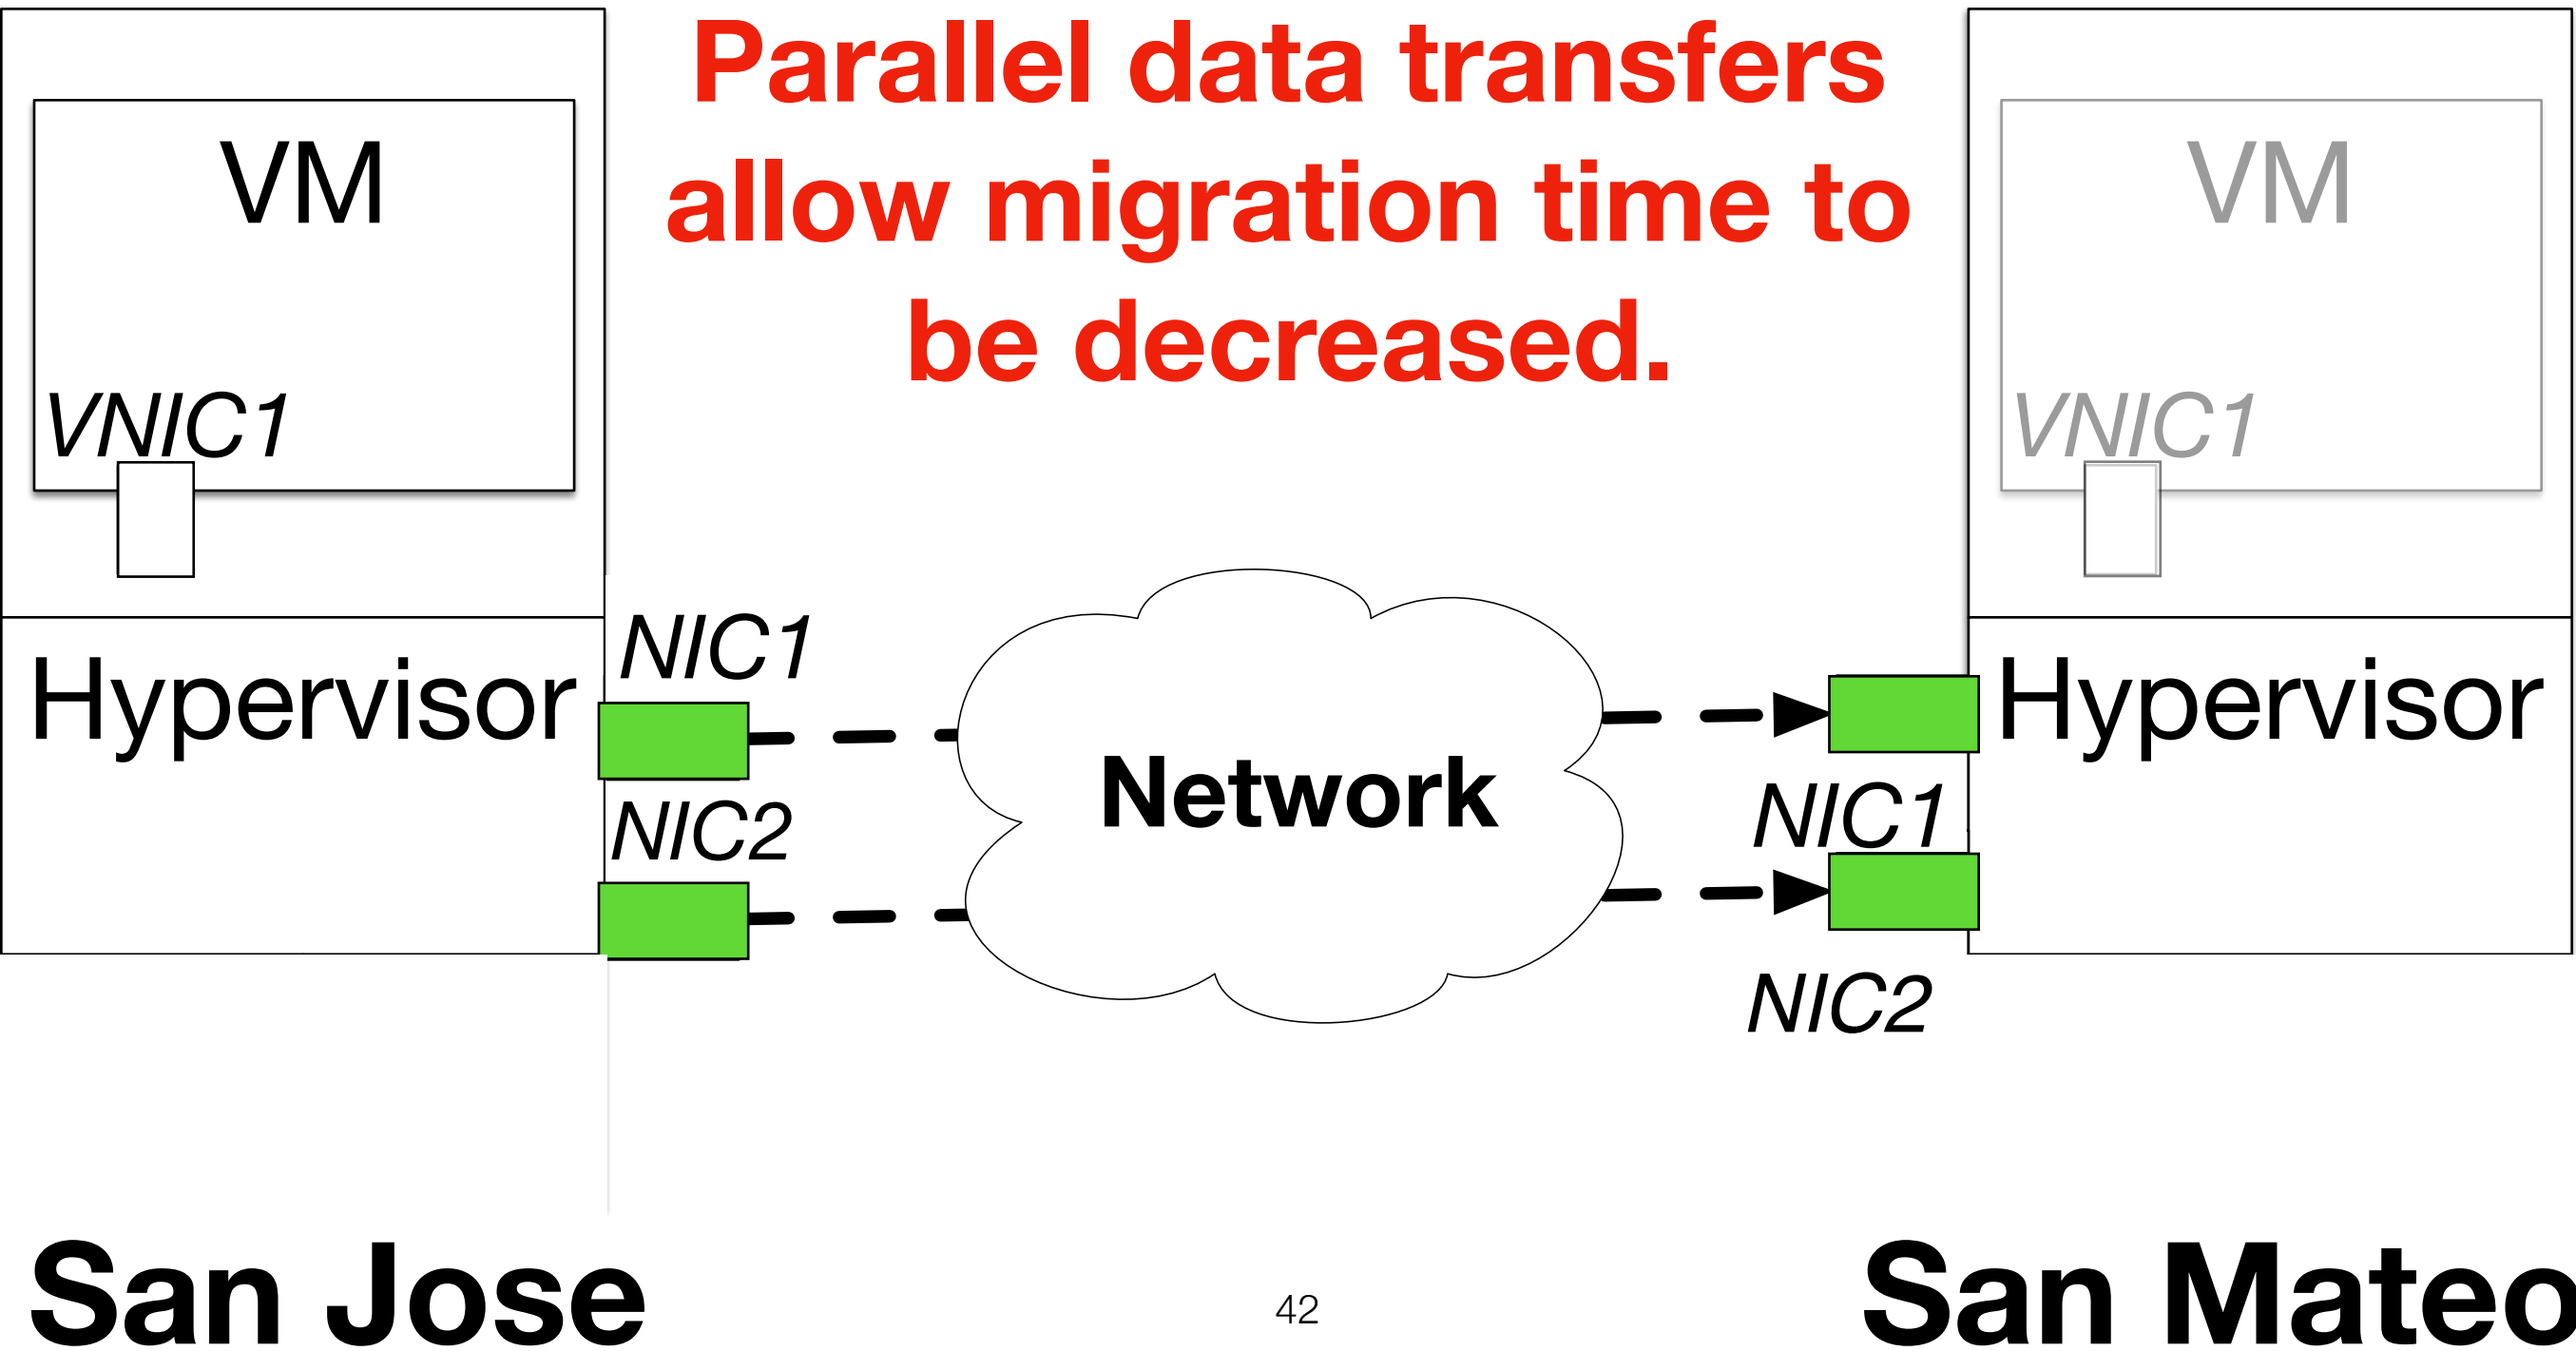
Migration

**Transfer Memory and
Disk State Over Both
Paths.**



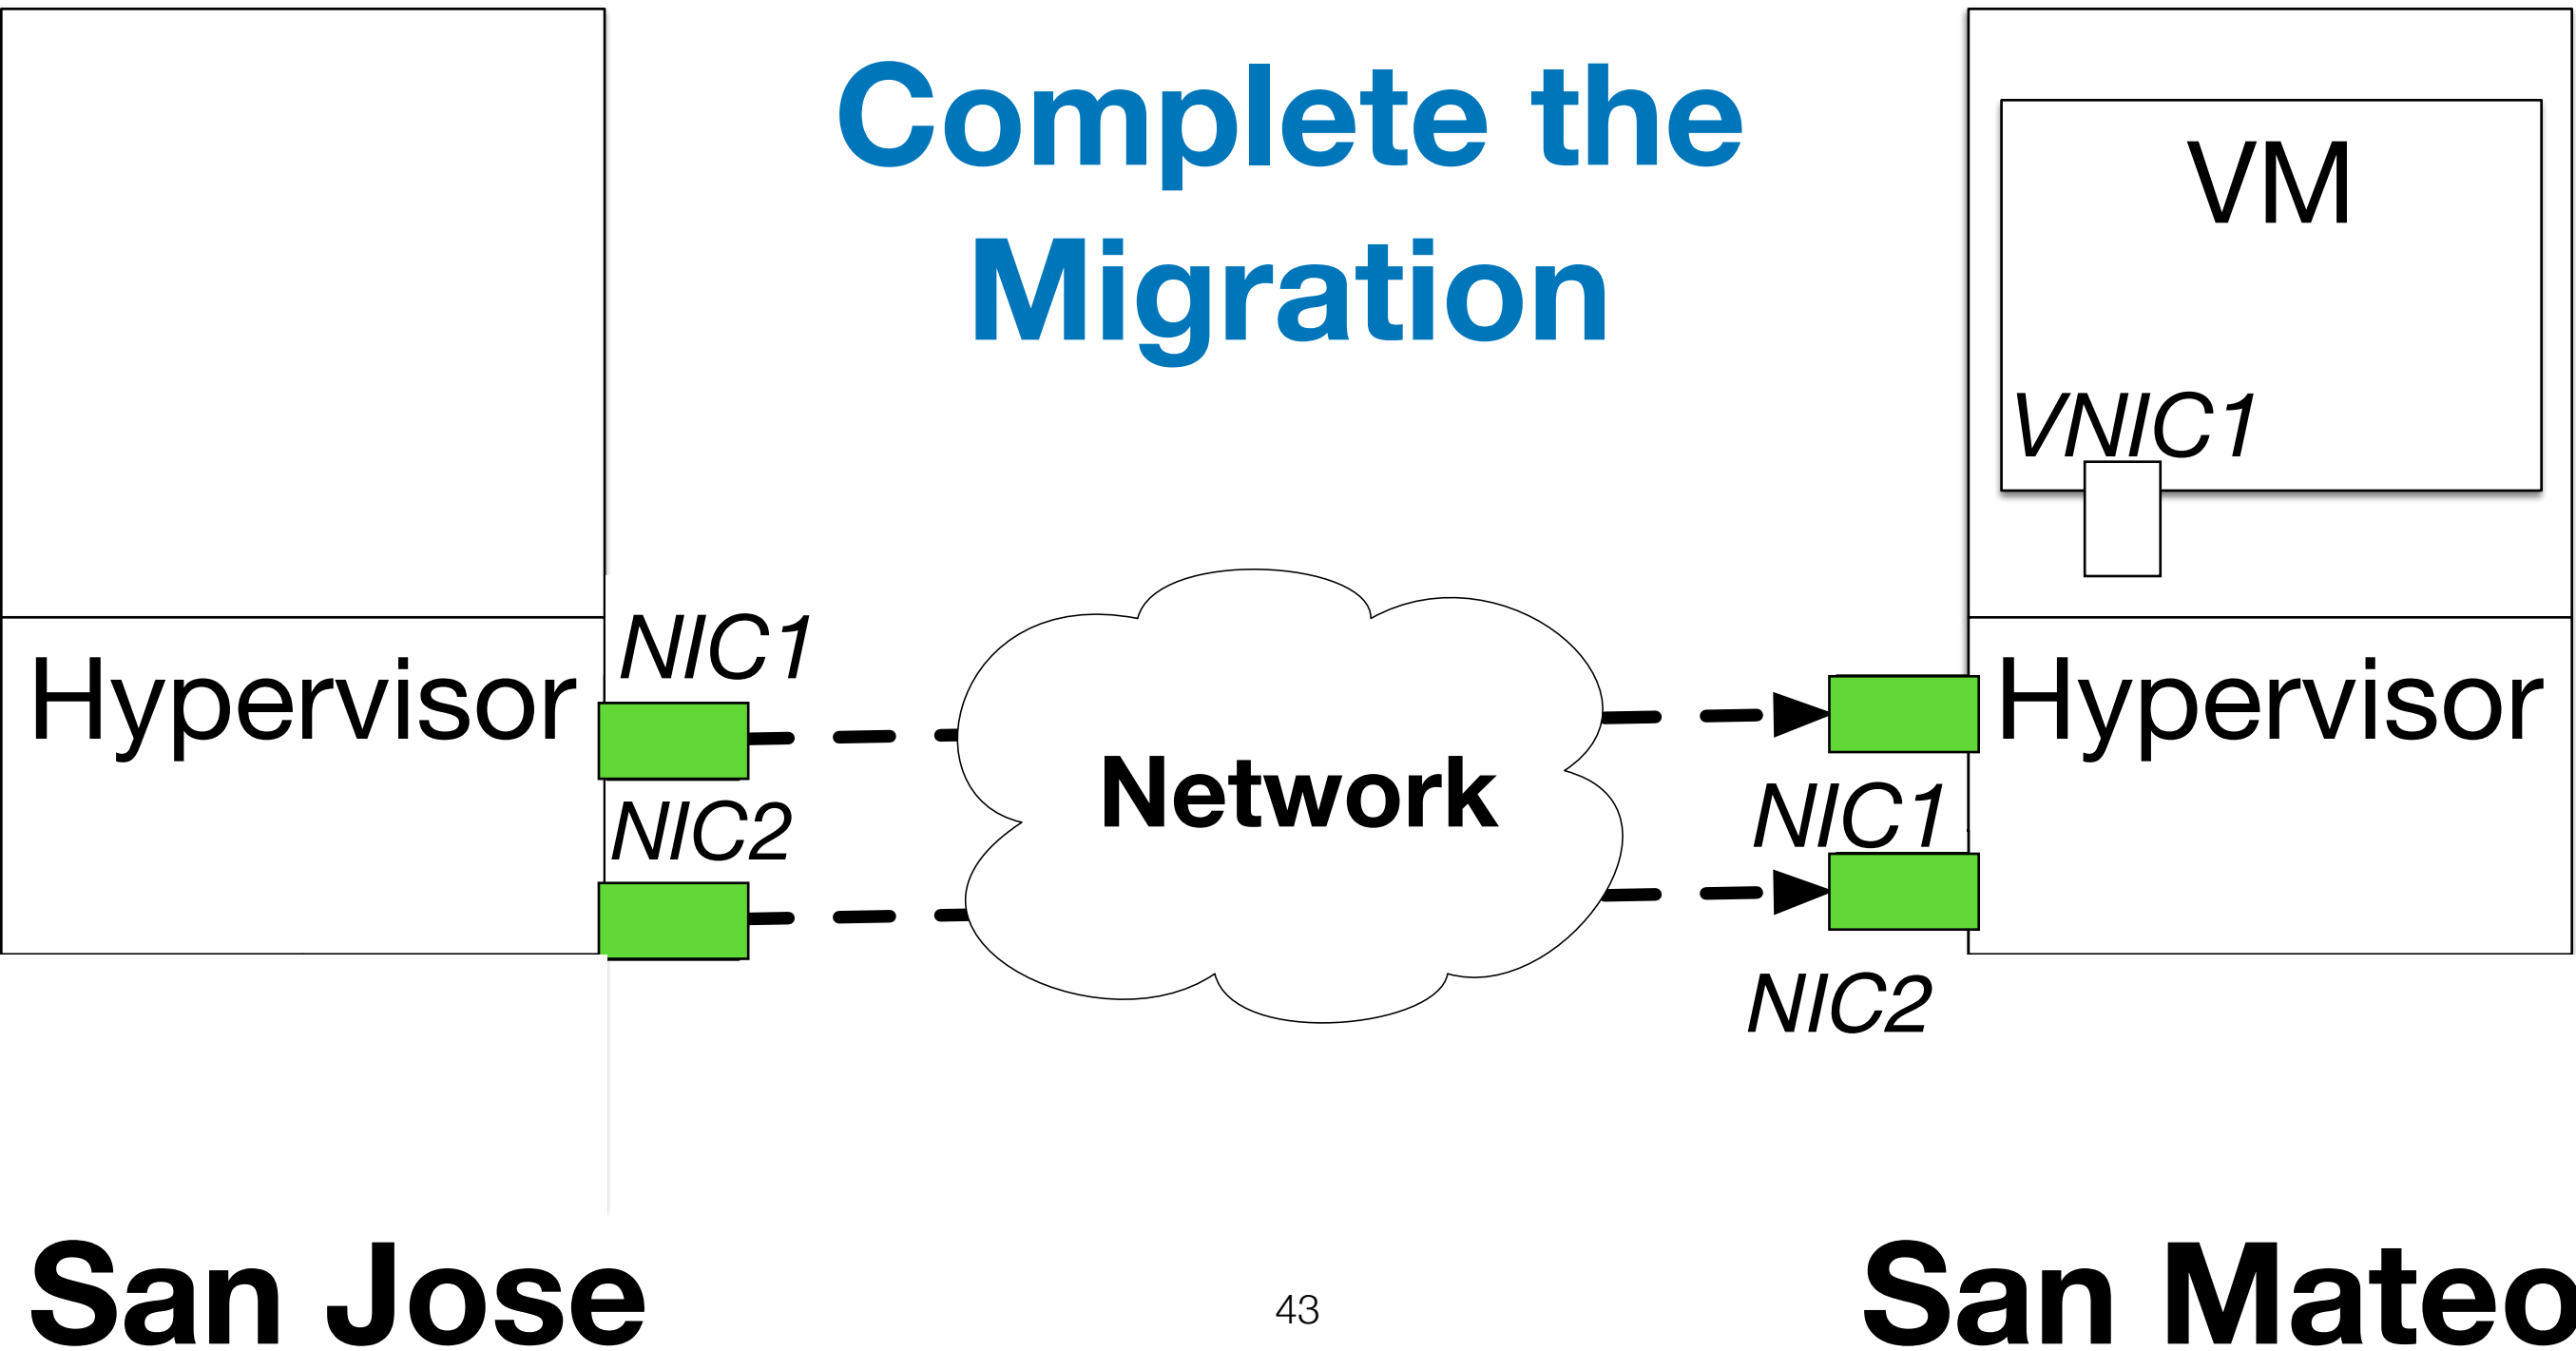
Migration

**Parallel data transfers
allow migration time to
be decreased.**



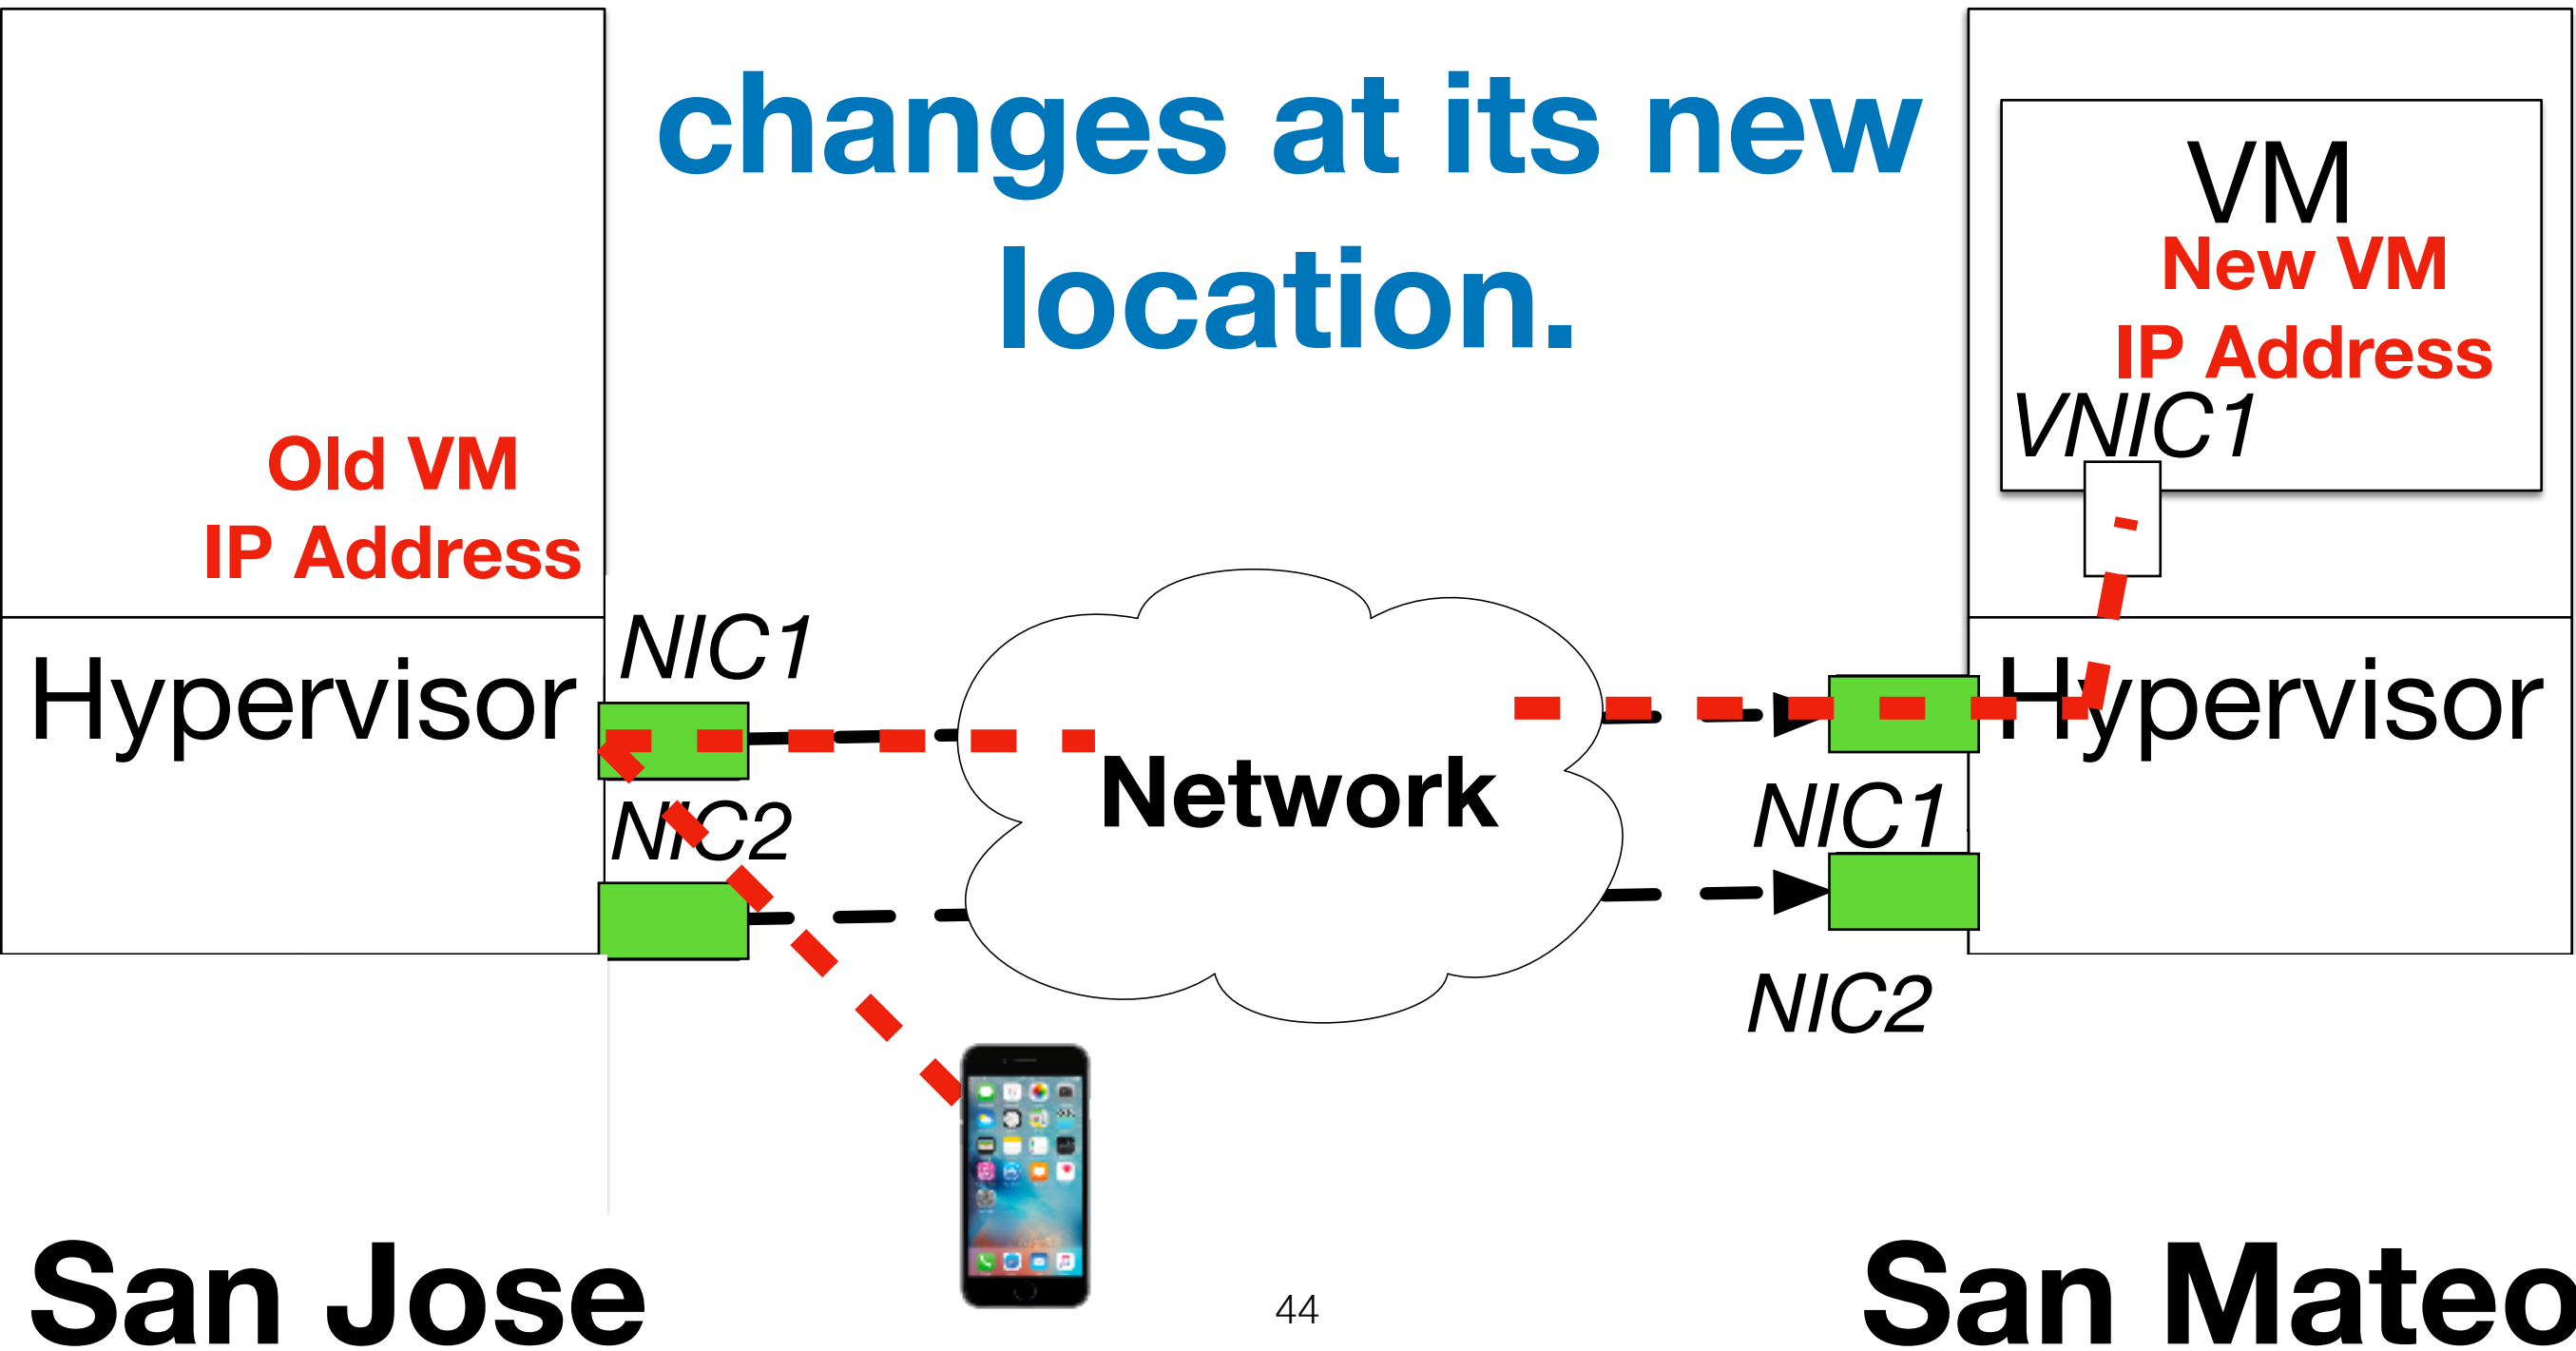
Migration

Complete the Migration



Migration

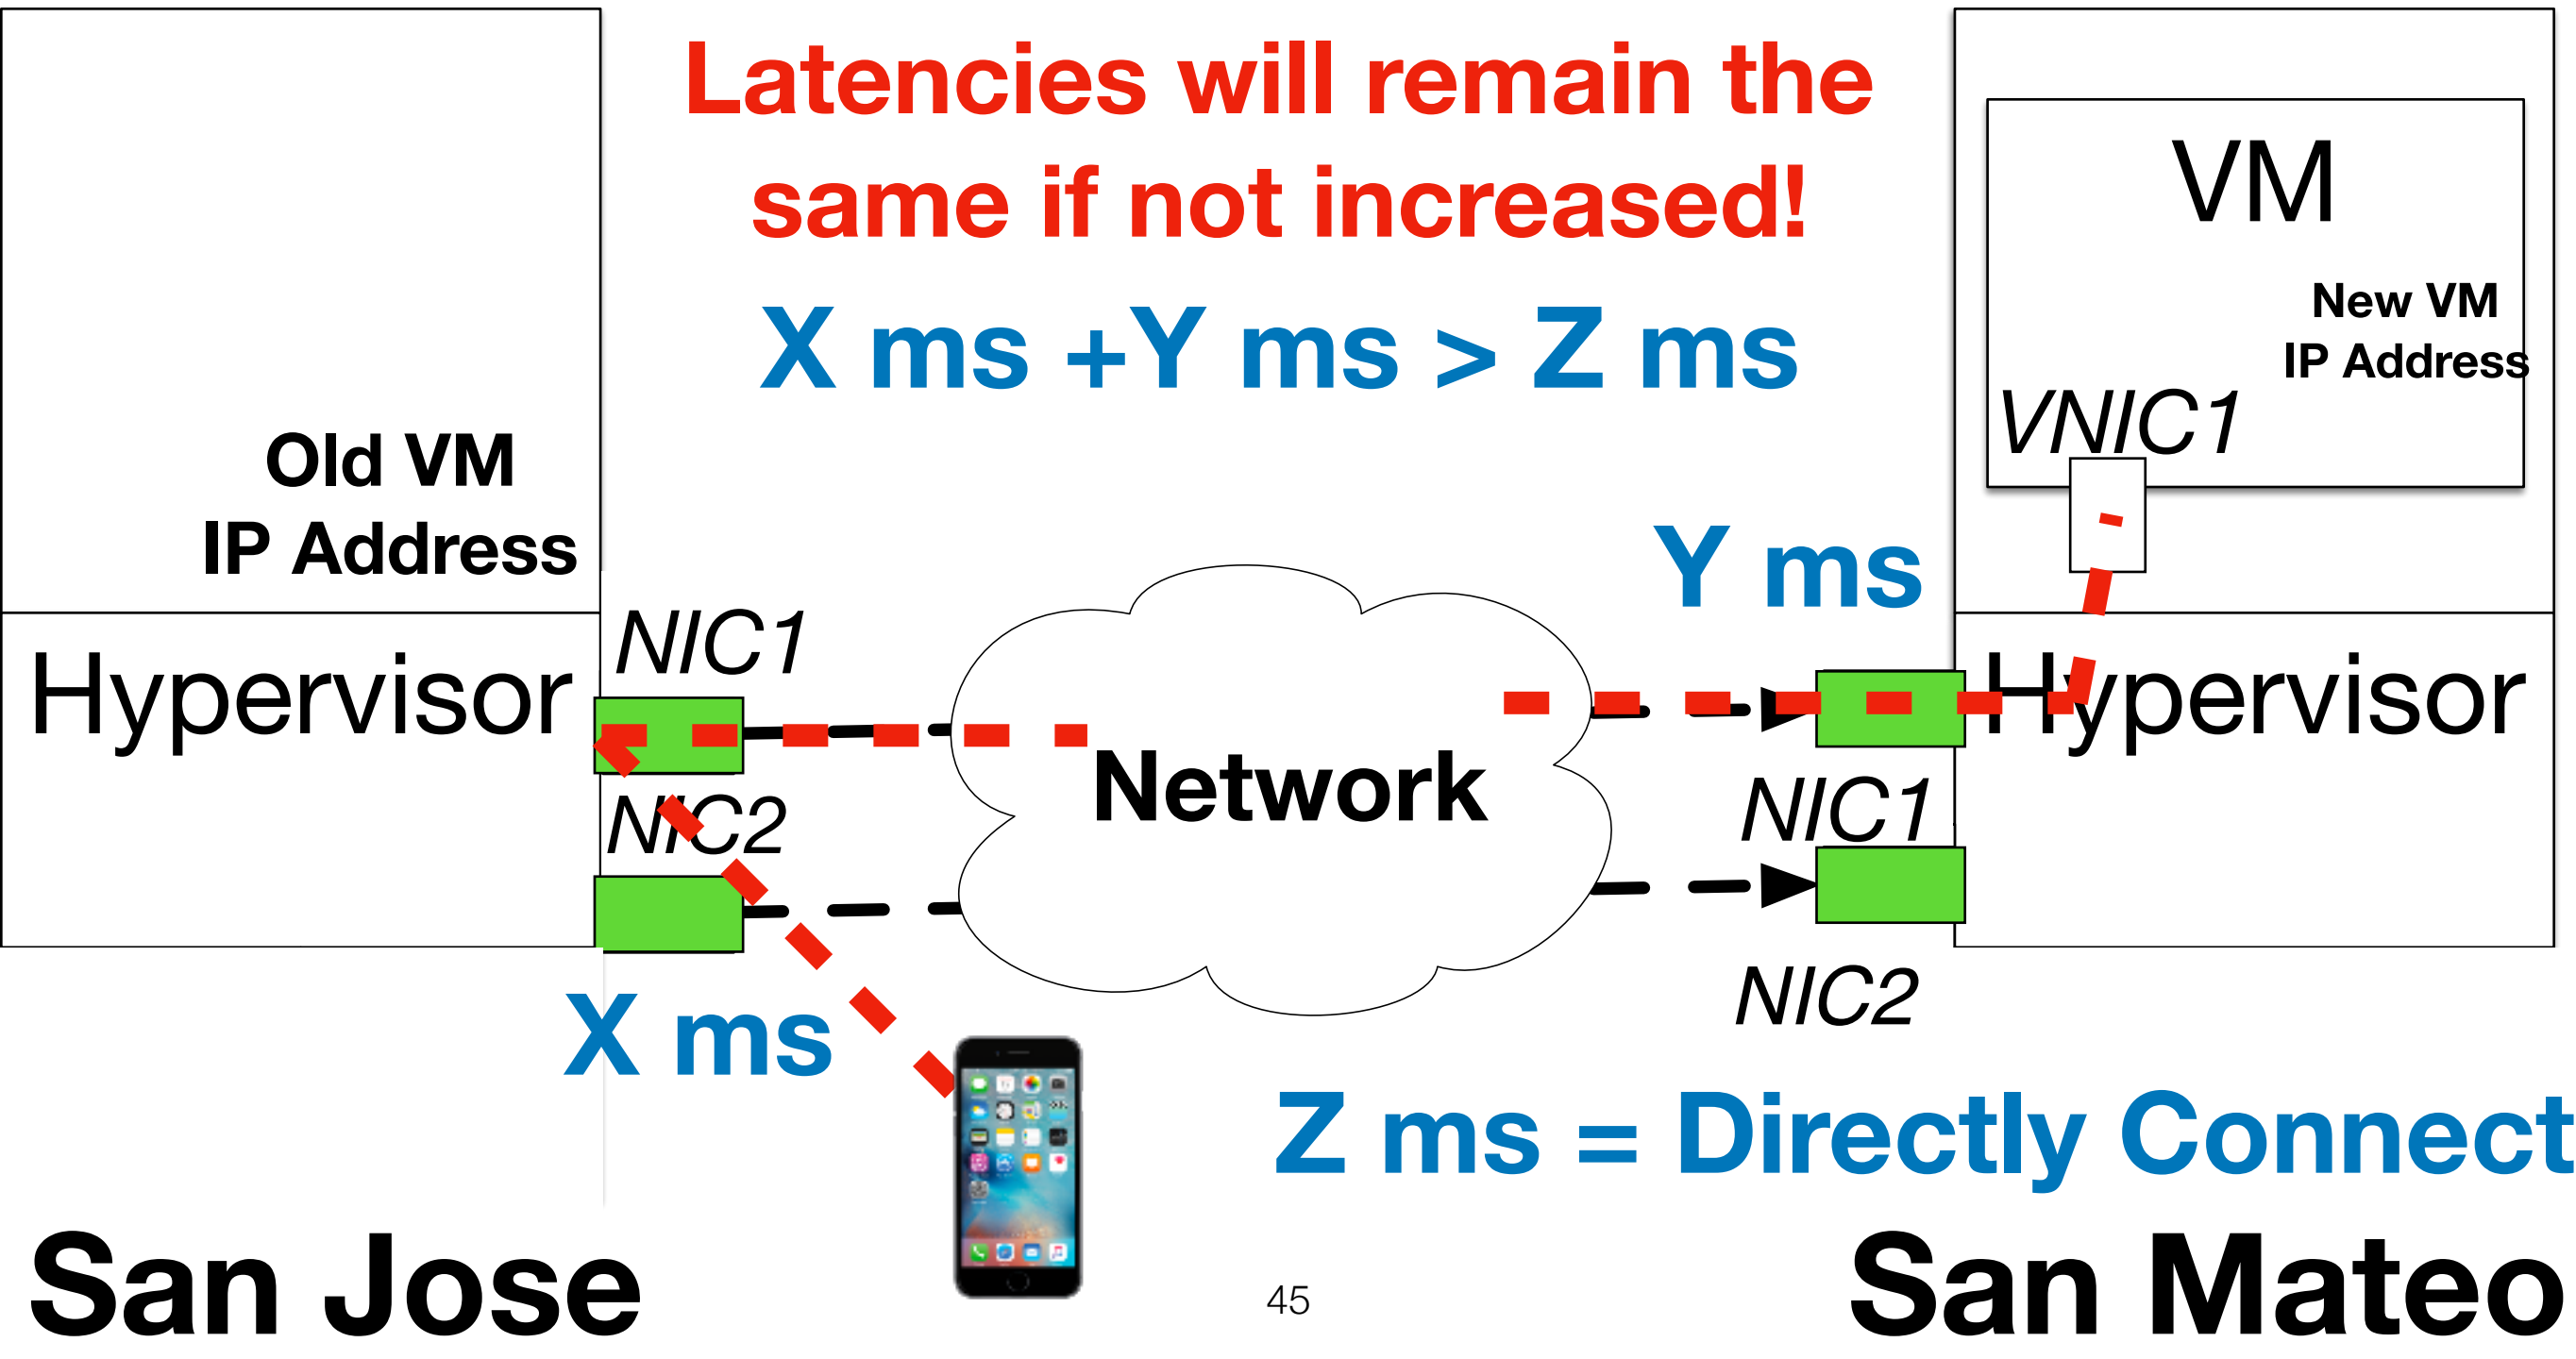
The VM IP address changes at its new location.



Migration

Latencies will remain the same if not increased!

$$X \text{ ms} + Y \text{ ms} > Z \text{ ms}$$




Maintaining Active Network Connections

- **Problem:** Current VM Migration approaches can cause latencies to increase for active network connections.


Maintaining Active Network Connections

- **Problem:** Current VM Migration approaches can cause latencies to increase for active network connections.
- **Solution:** MPTCP's seamless address handoff enables transparency in the network.


An MPTCP Socket




An MPTCP Socket



An MPTCP Socket




An MPTCP Socket




An MPTCP Socket

- MPTCP's ability to add and remove addresses enables transparency in networks.
- We can take advantage of MPTCP's adaptability to notify clients of new addresses.



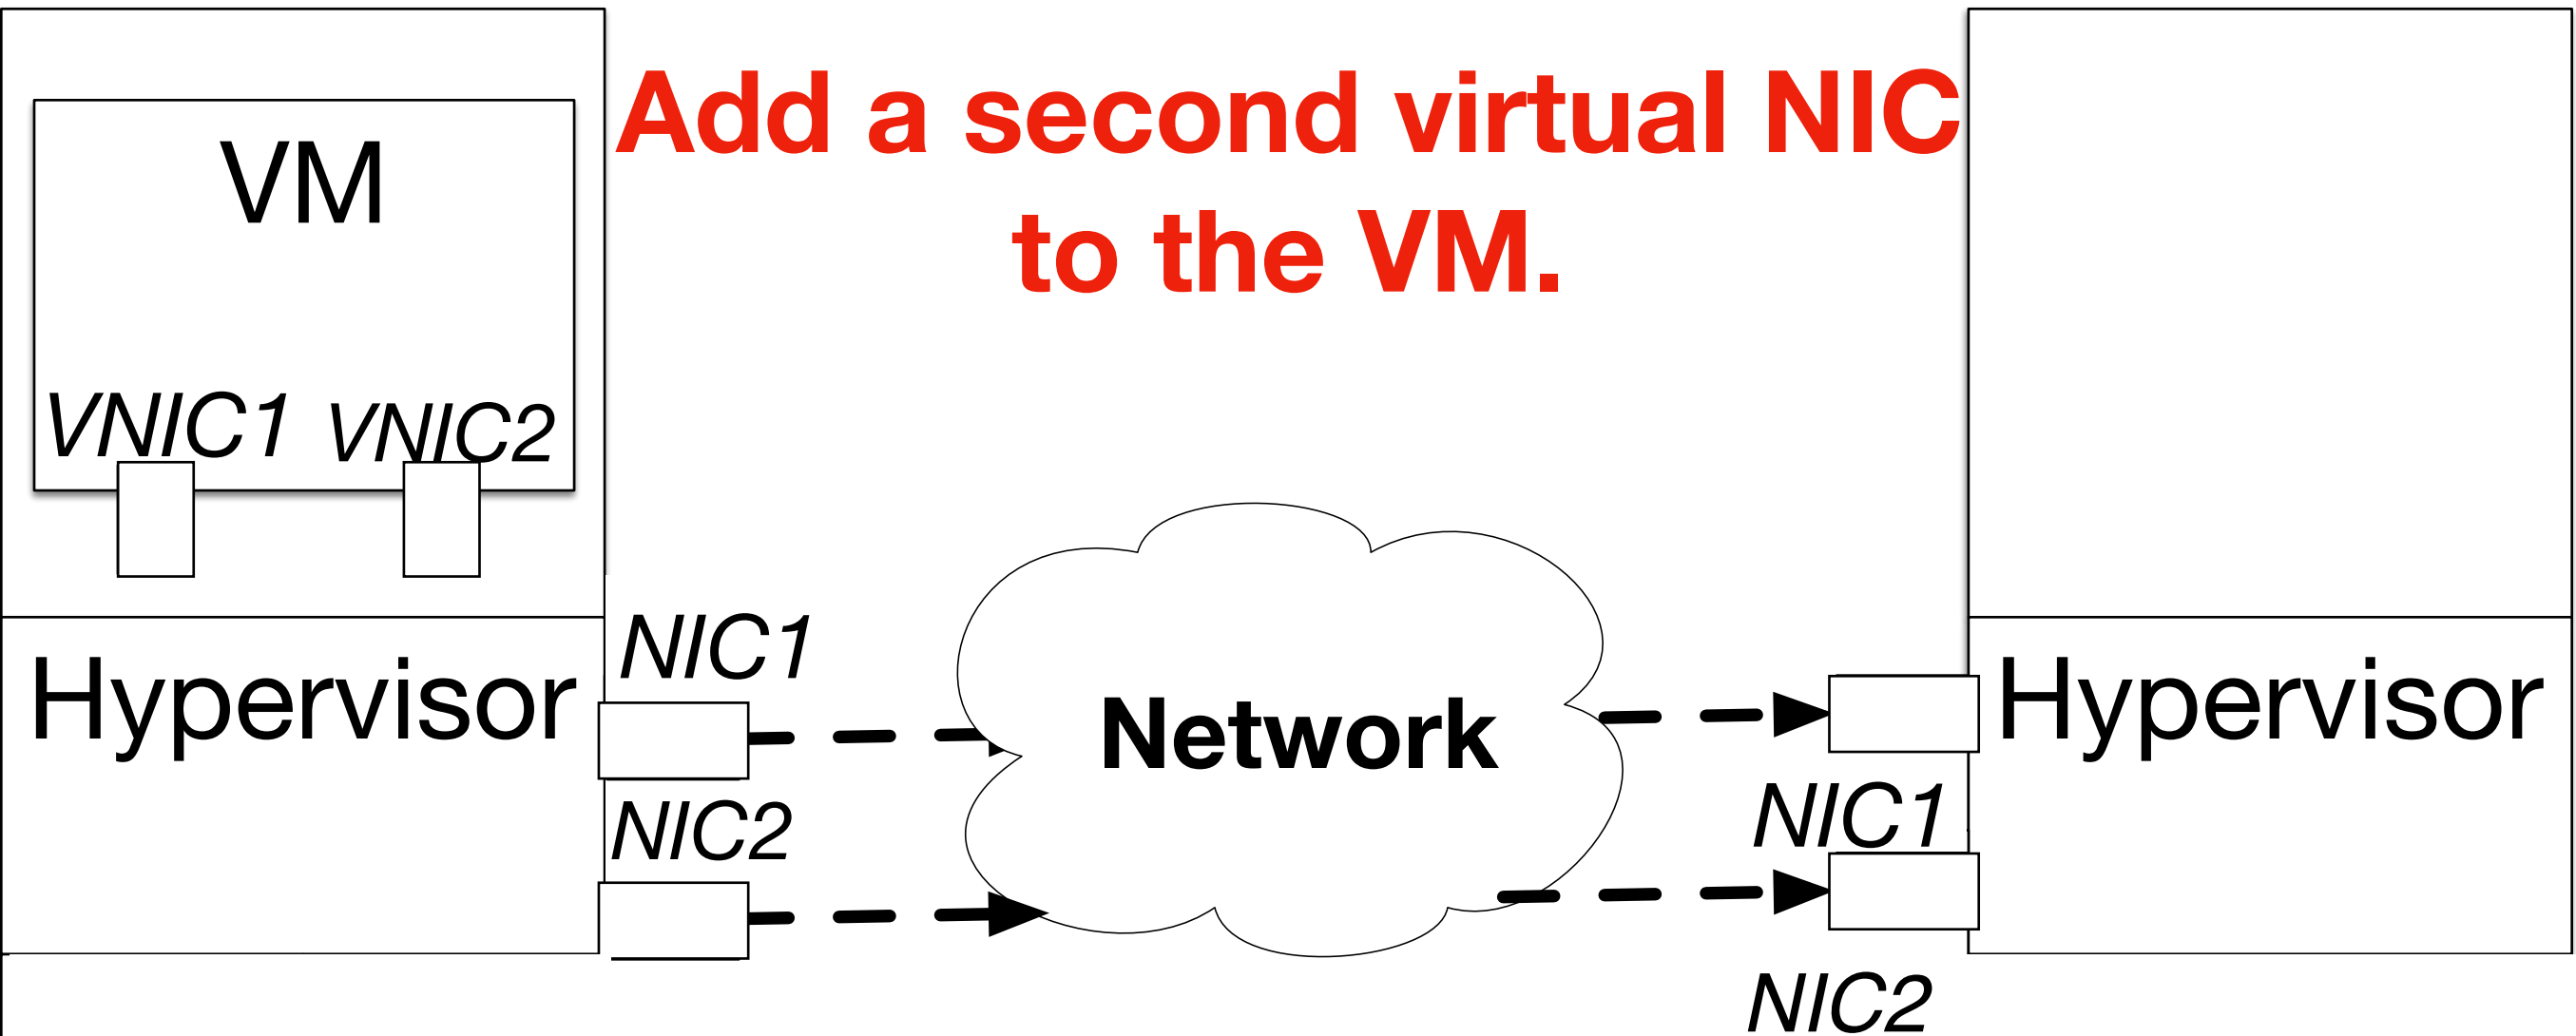
Migration



San Jose

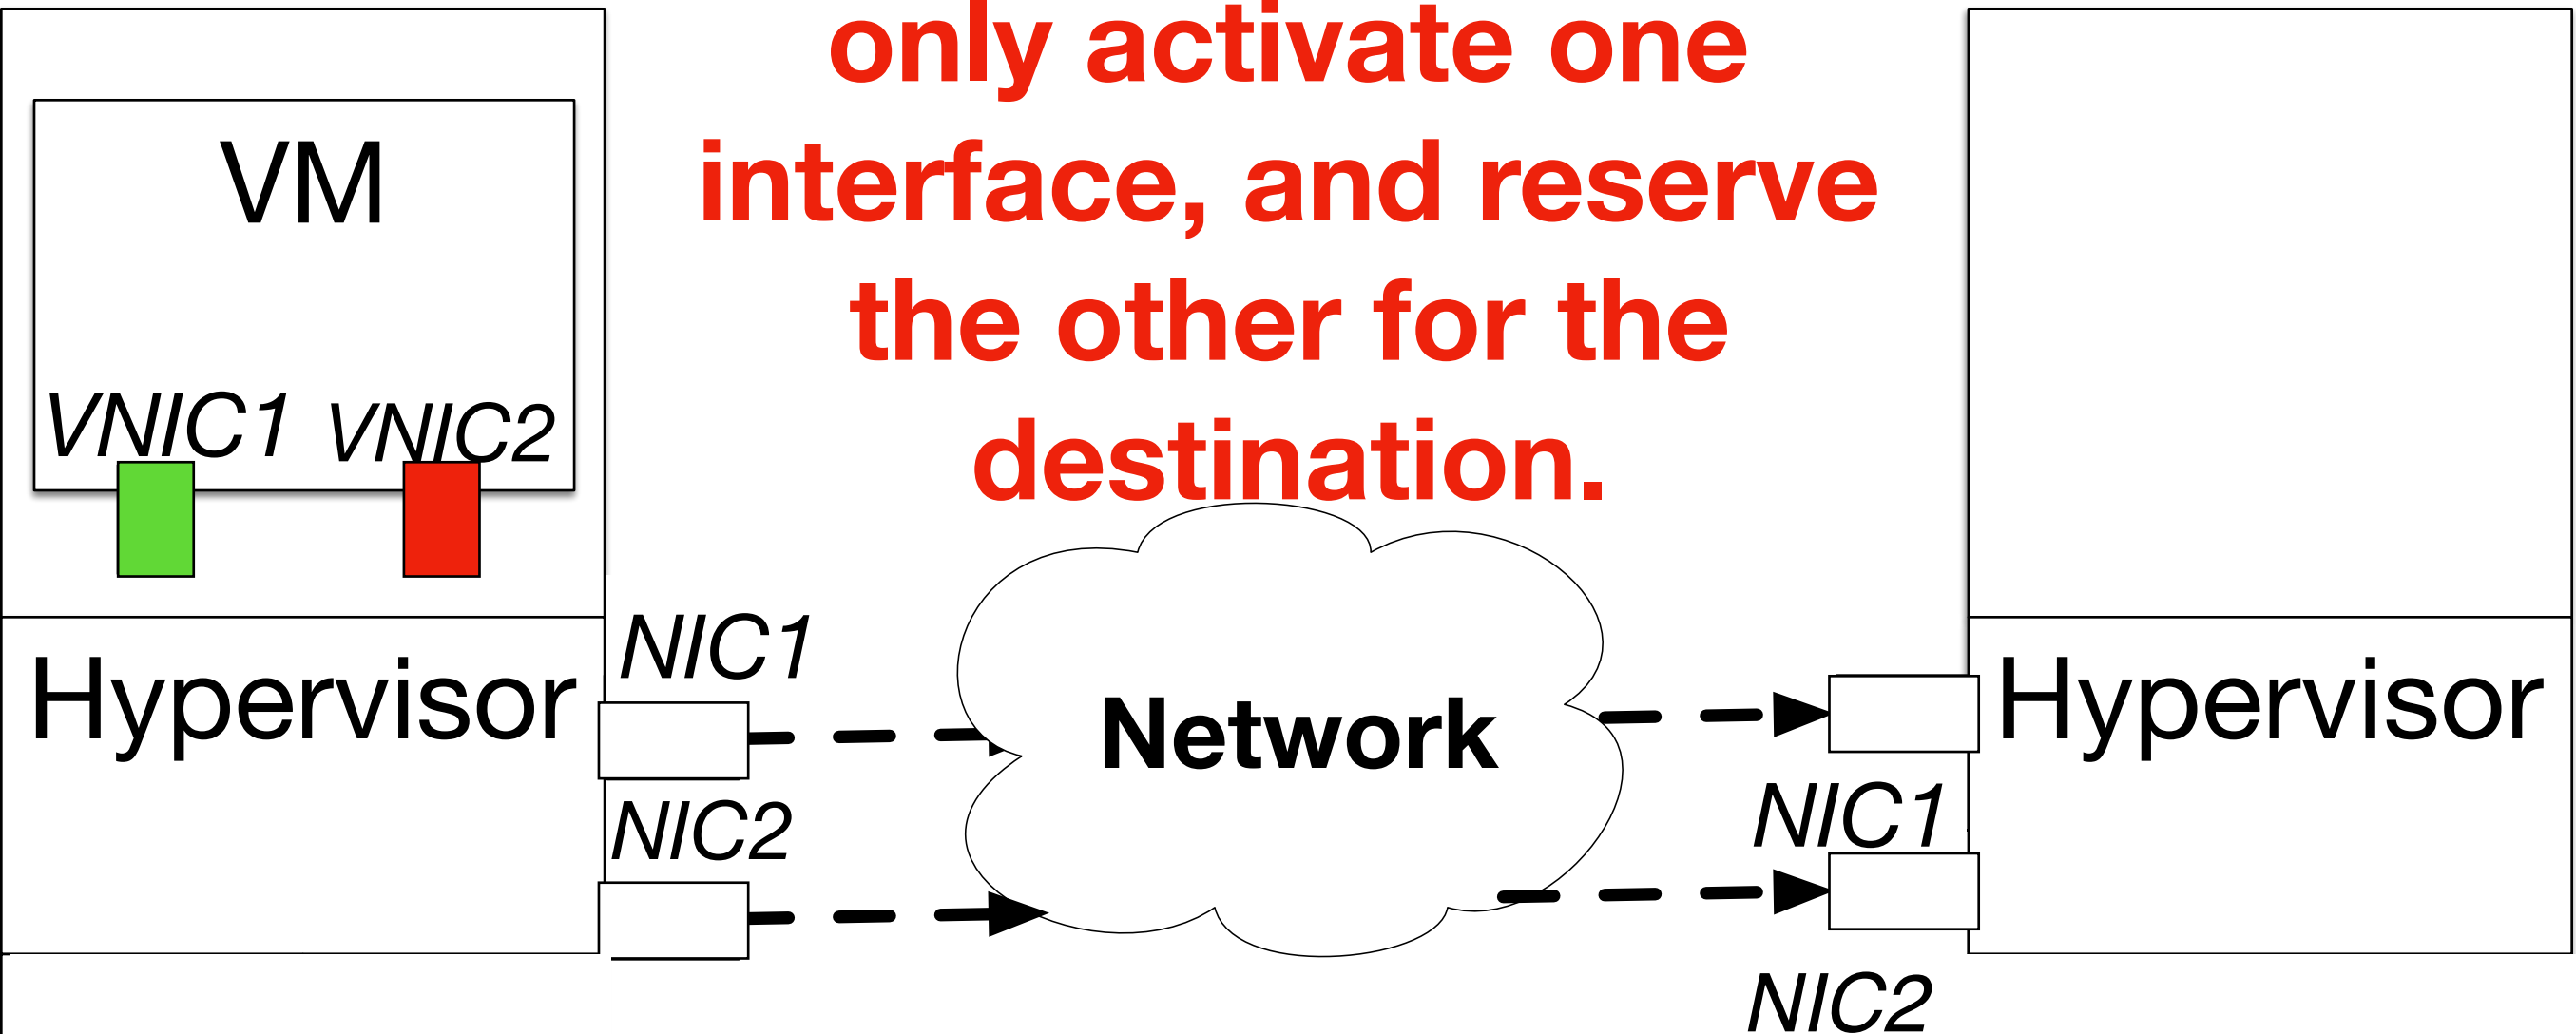
San Mateo

Migration



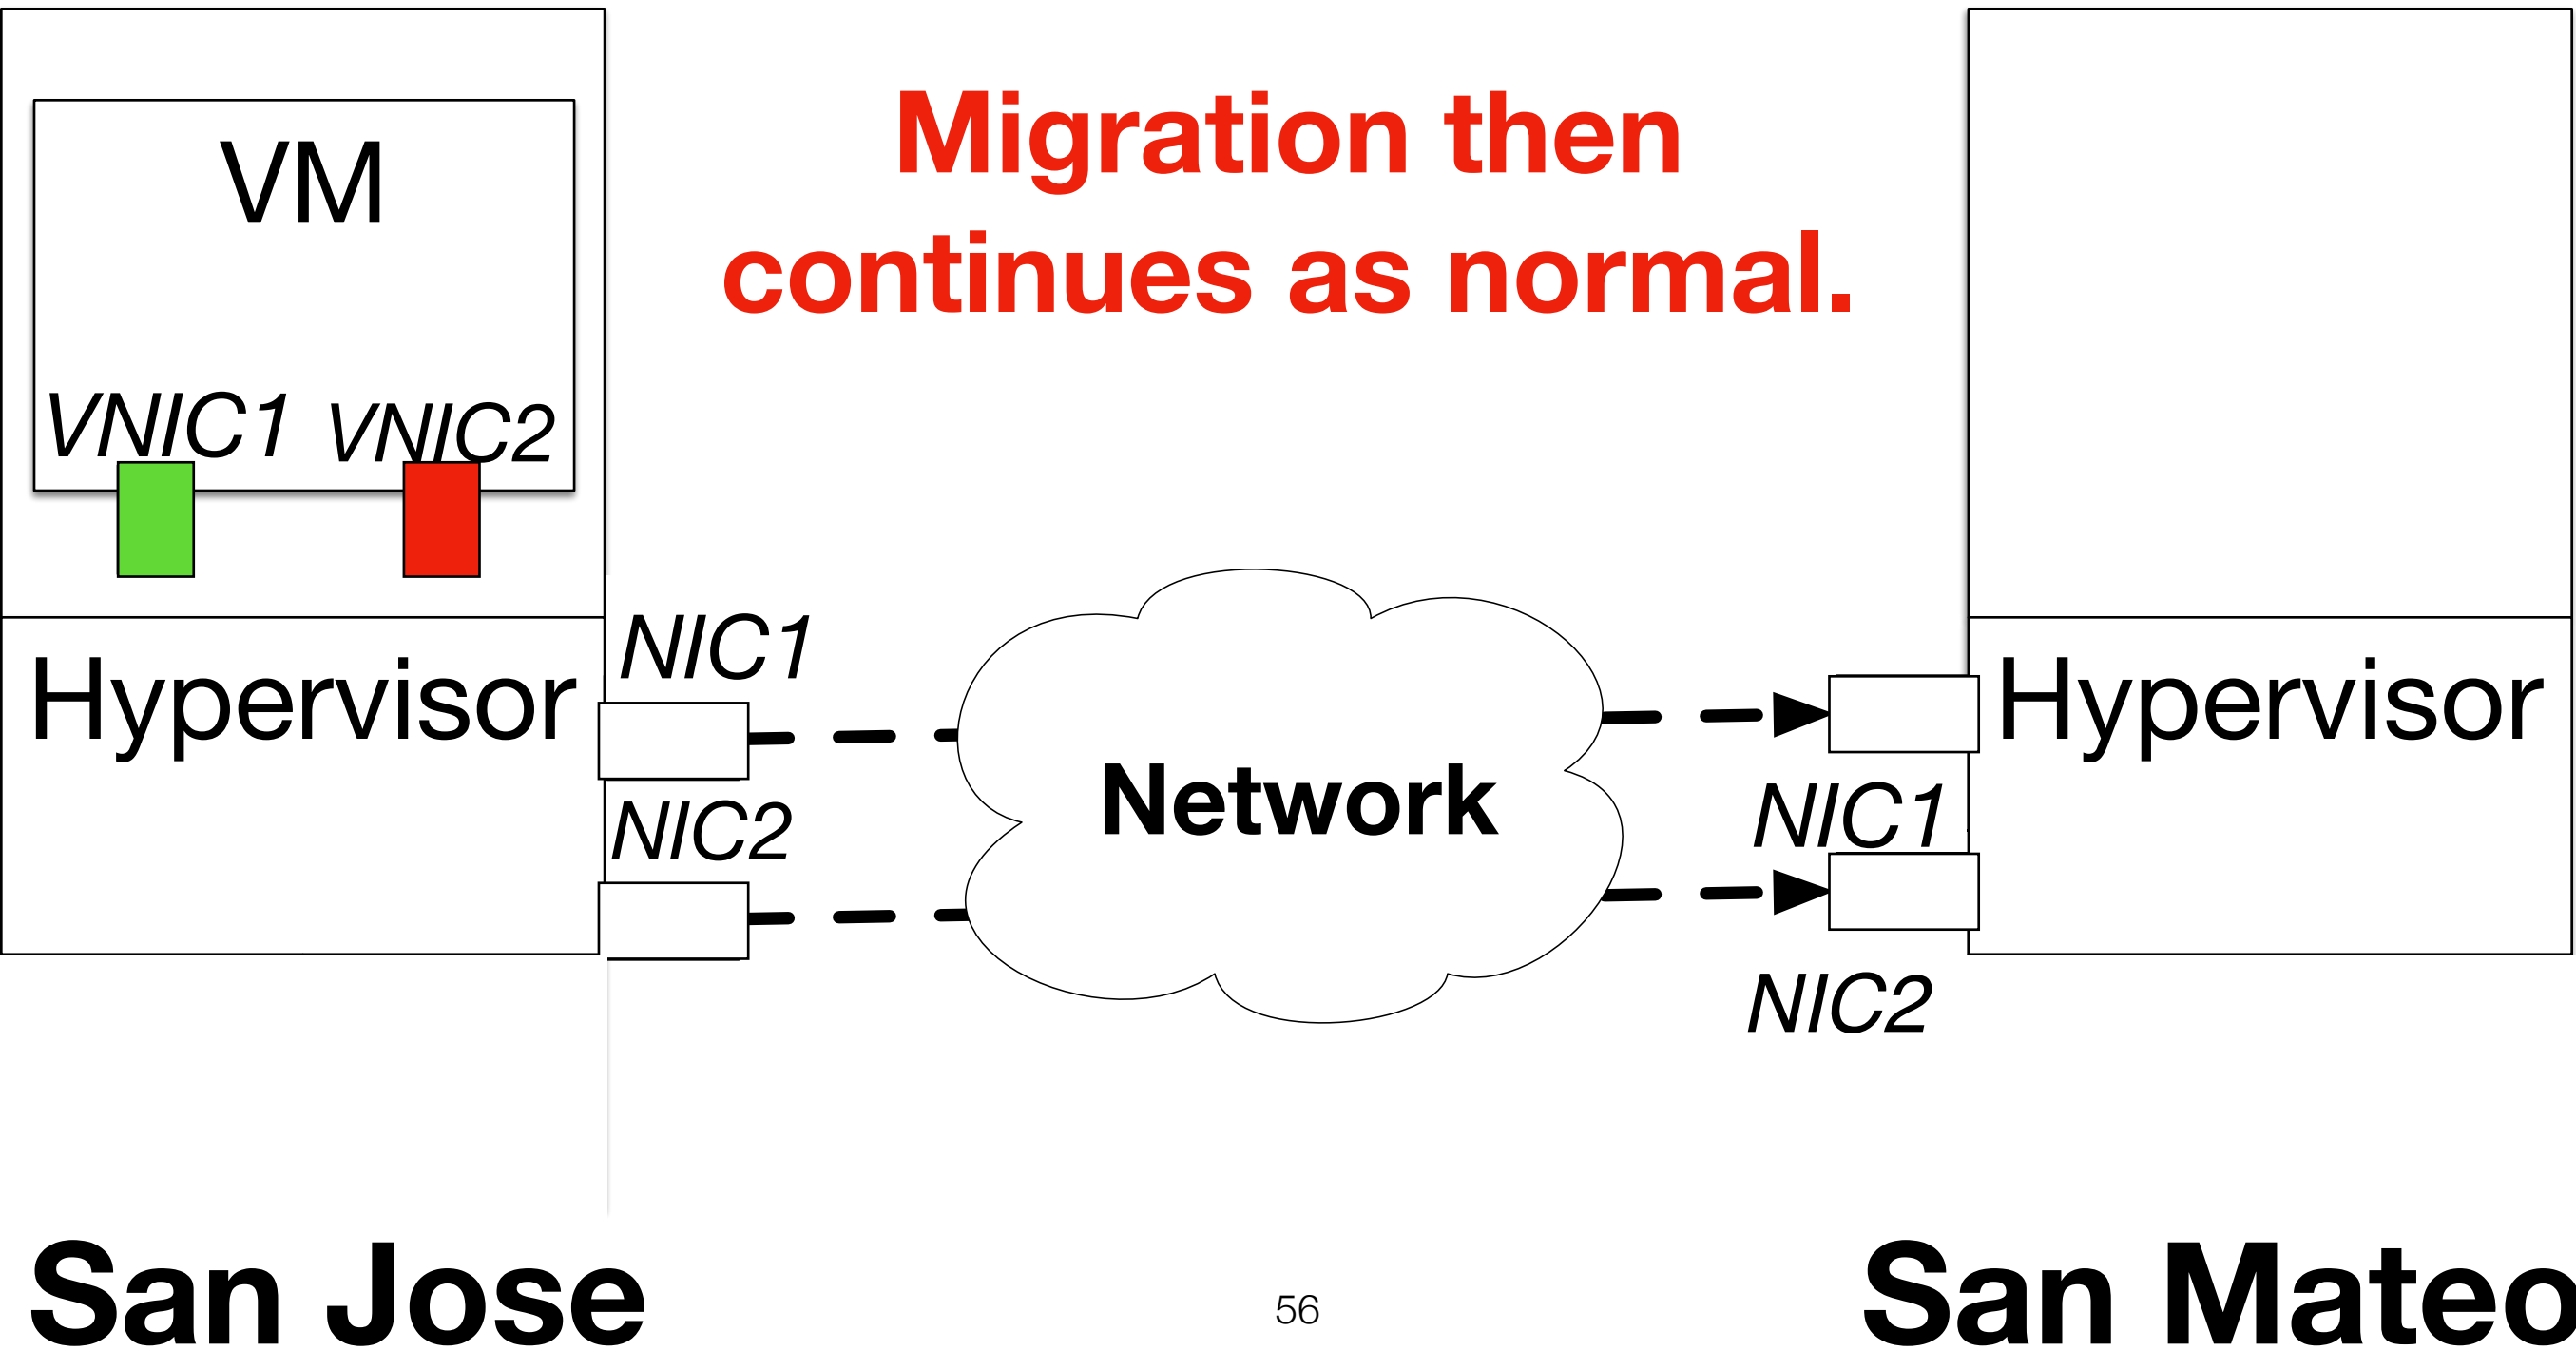
Migration

At the guest level we only activate one interface, and reserve the other for the destination.

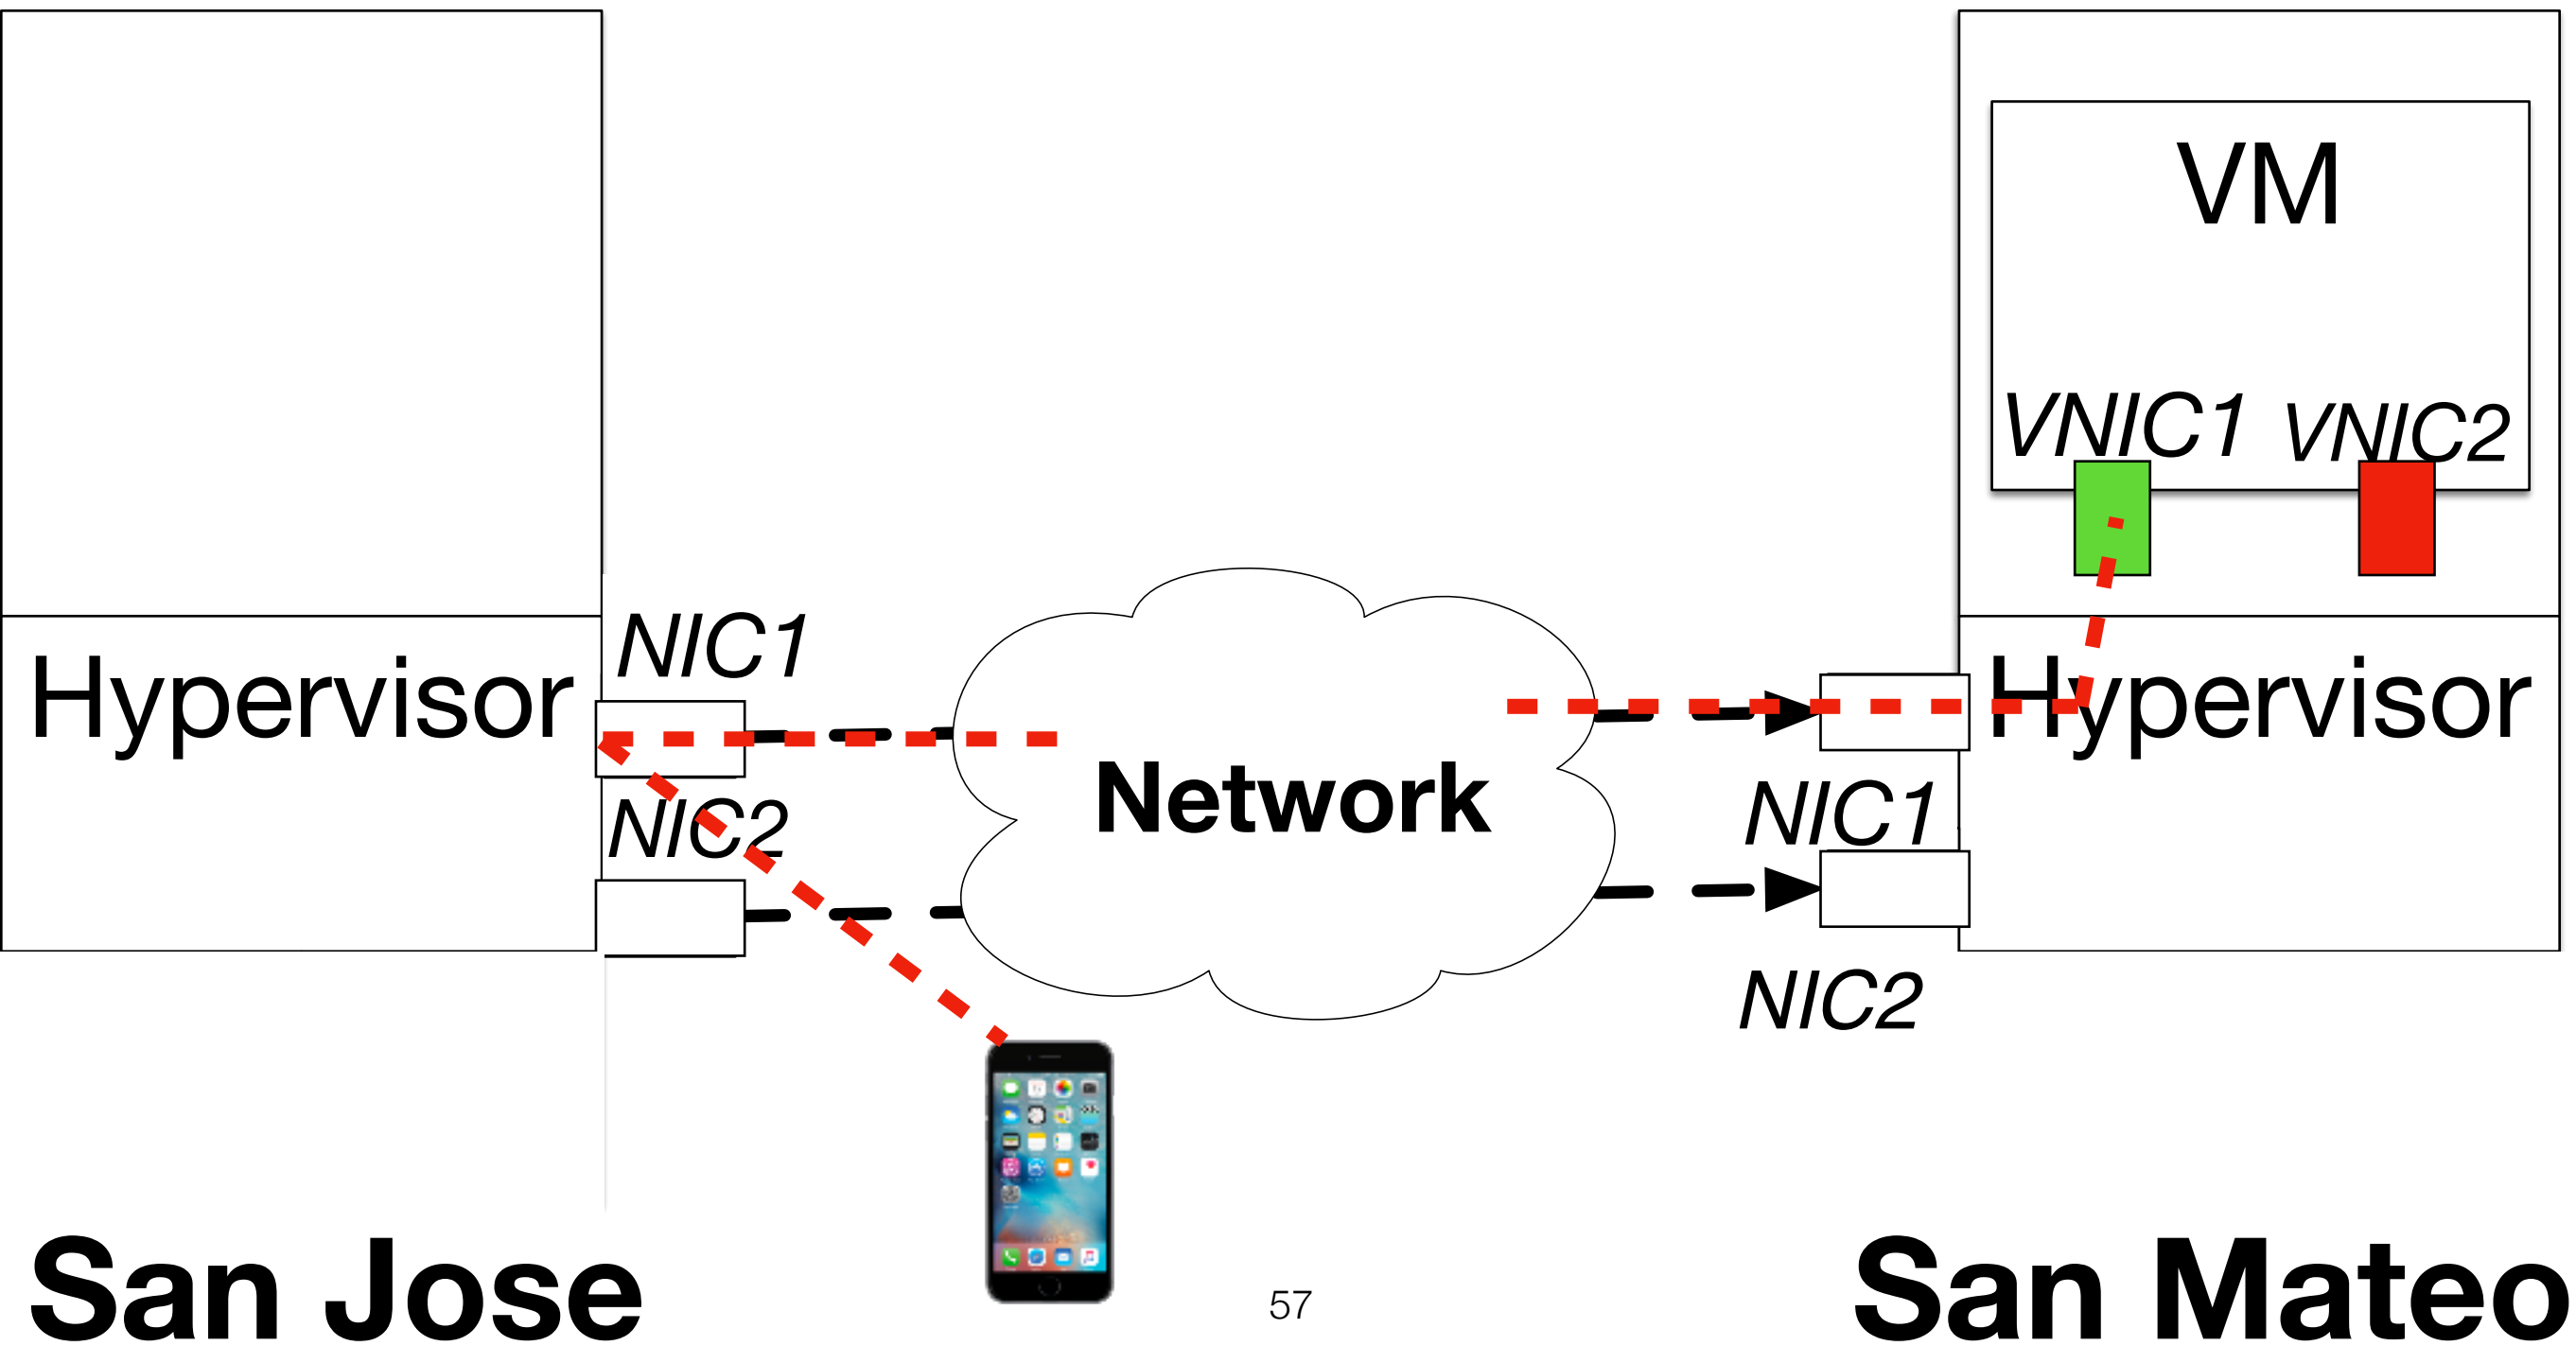


Migration

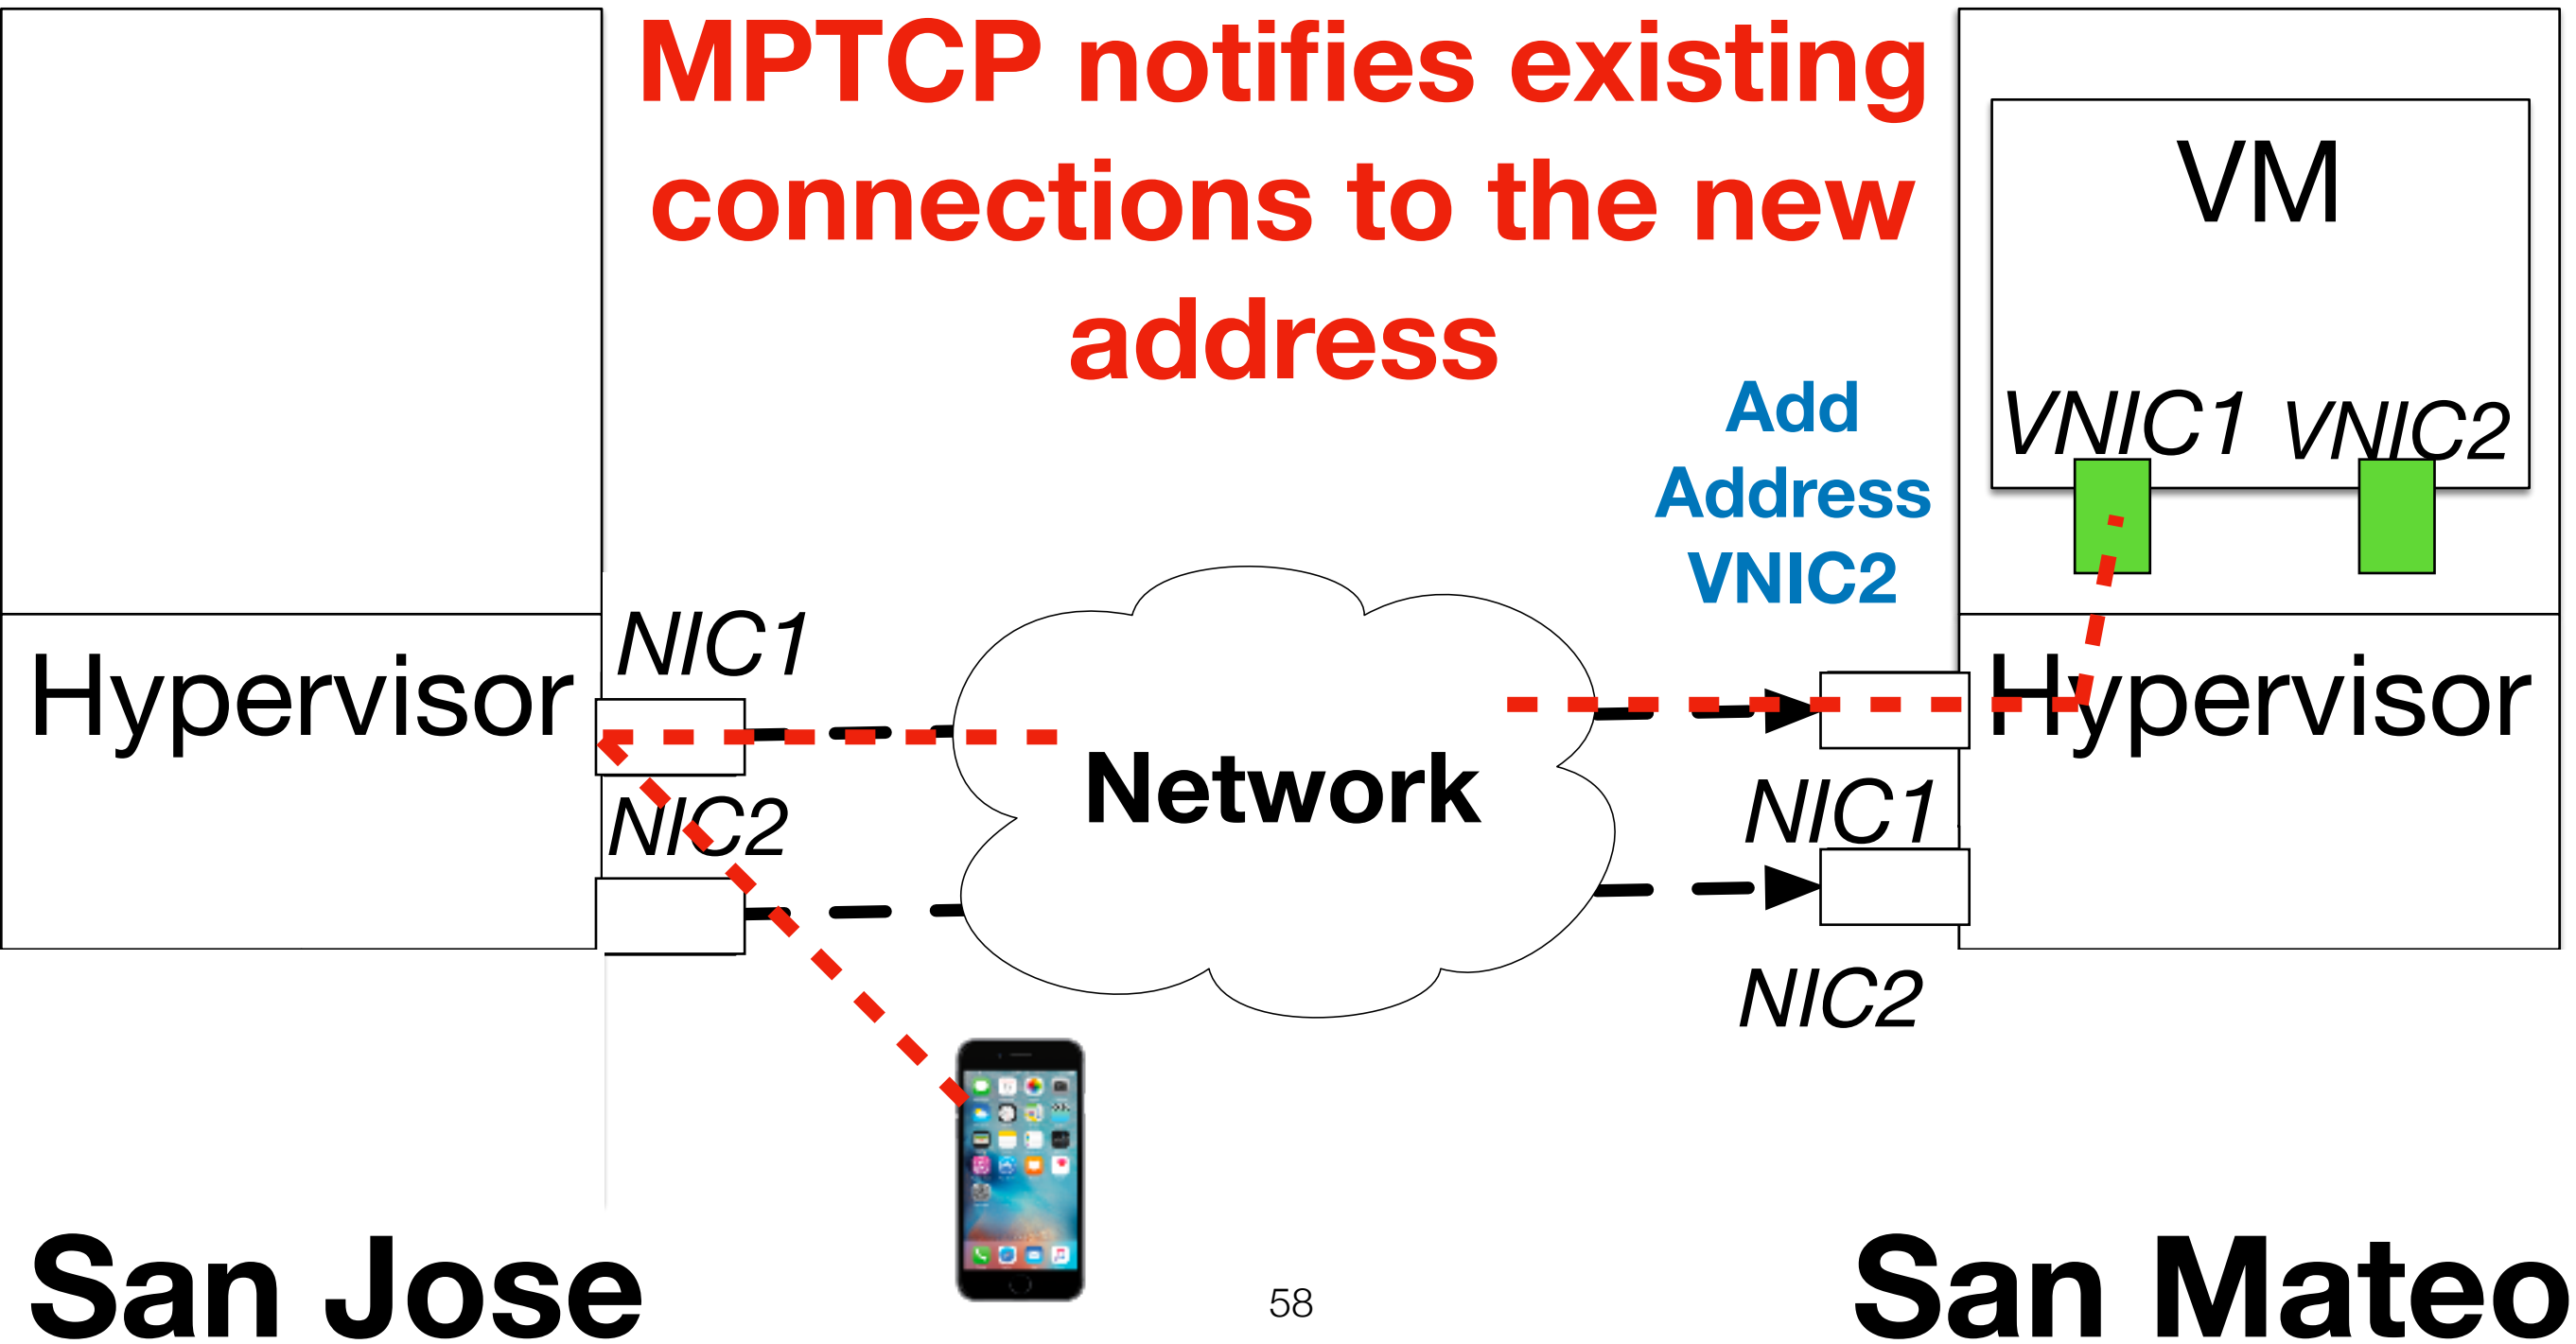
**Migration then
continues as normal.**



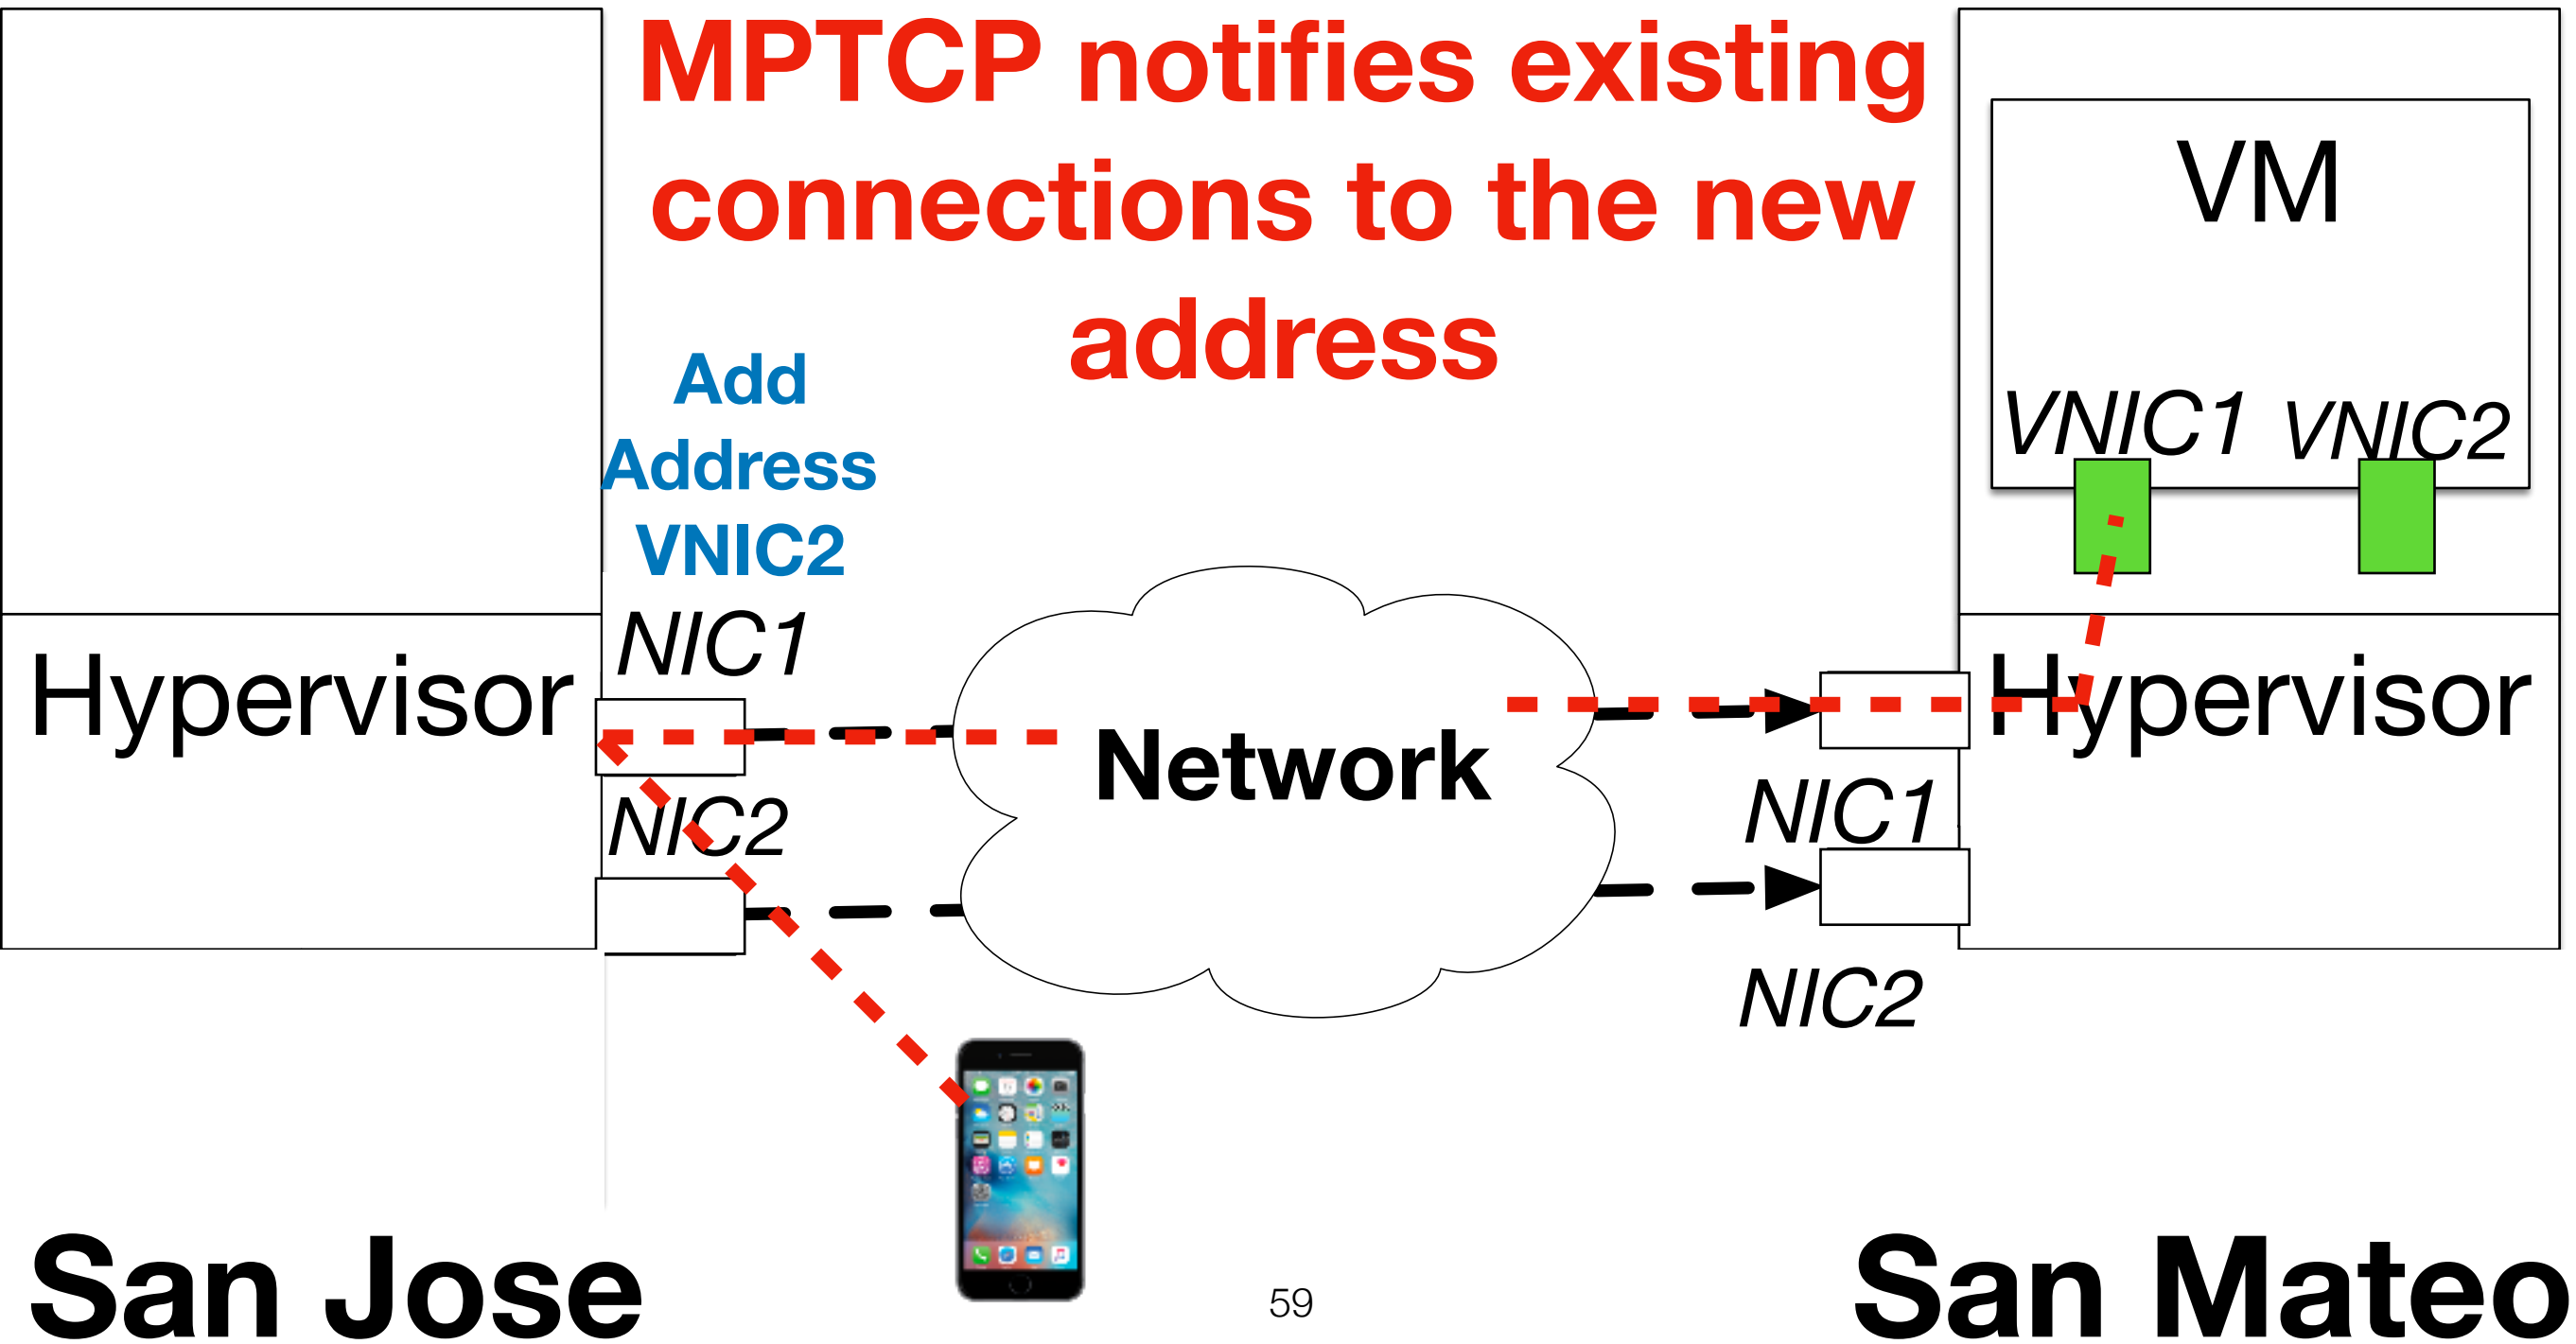
Migration



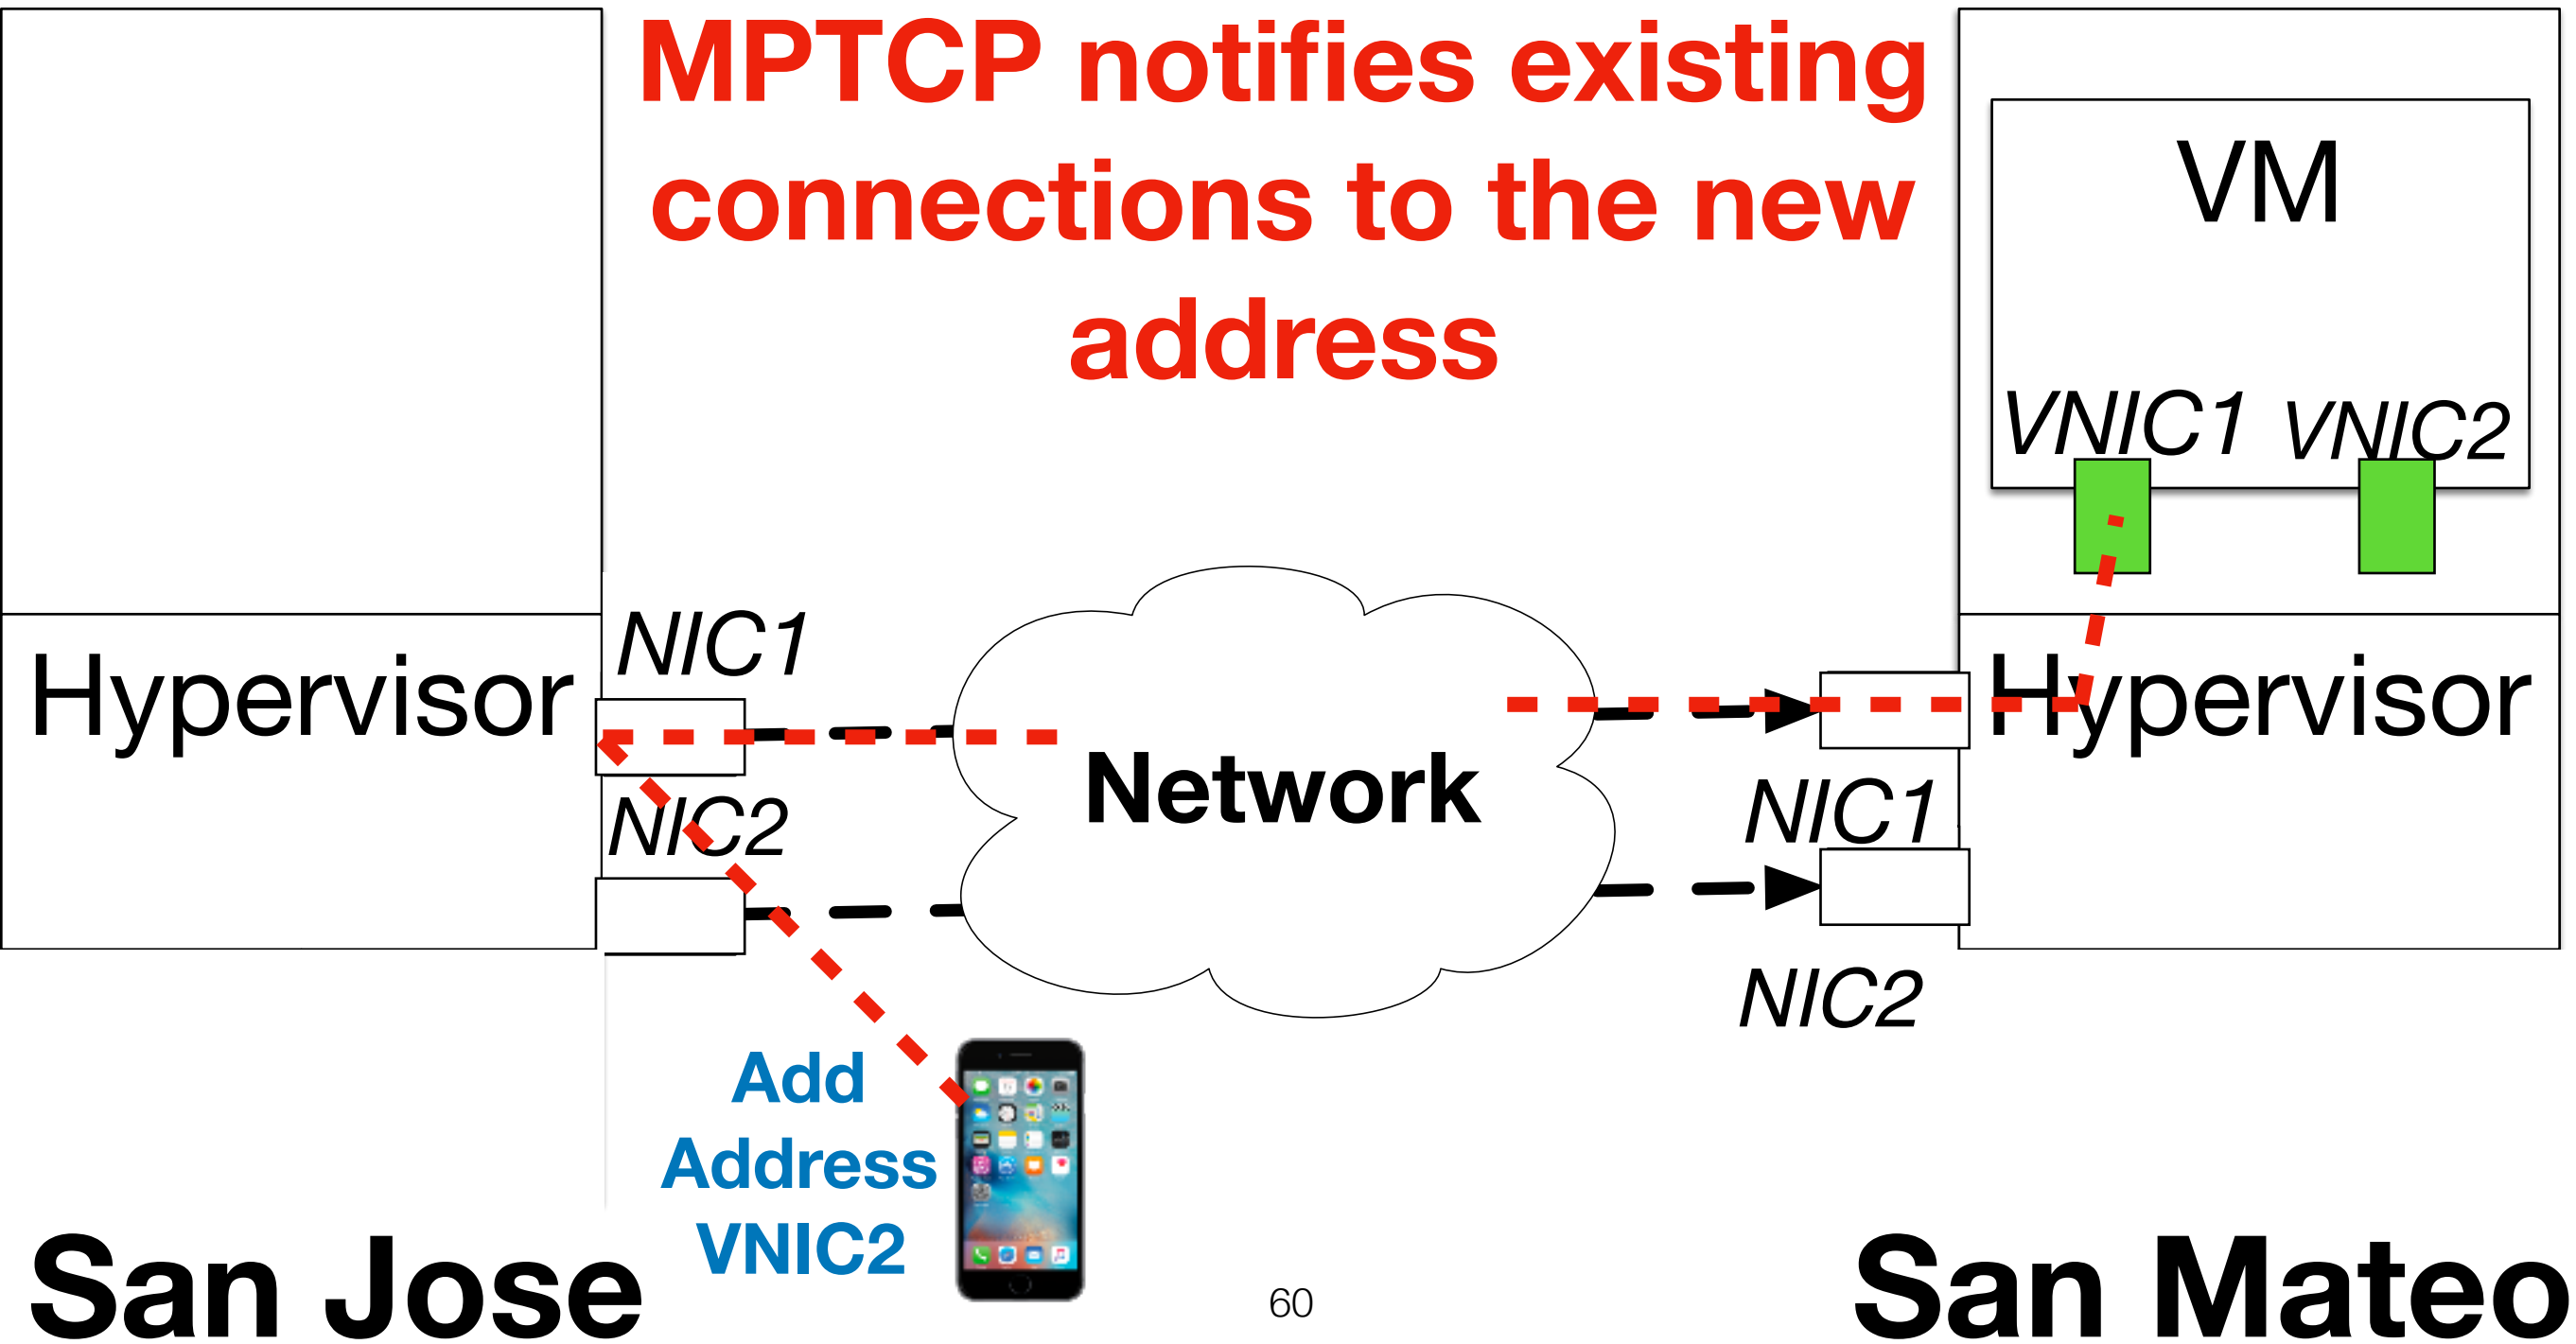
Migration



Migration

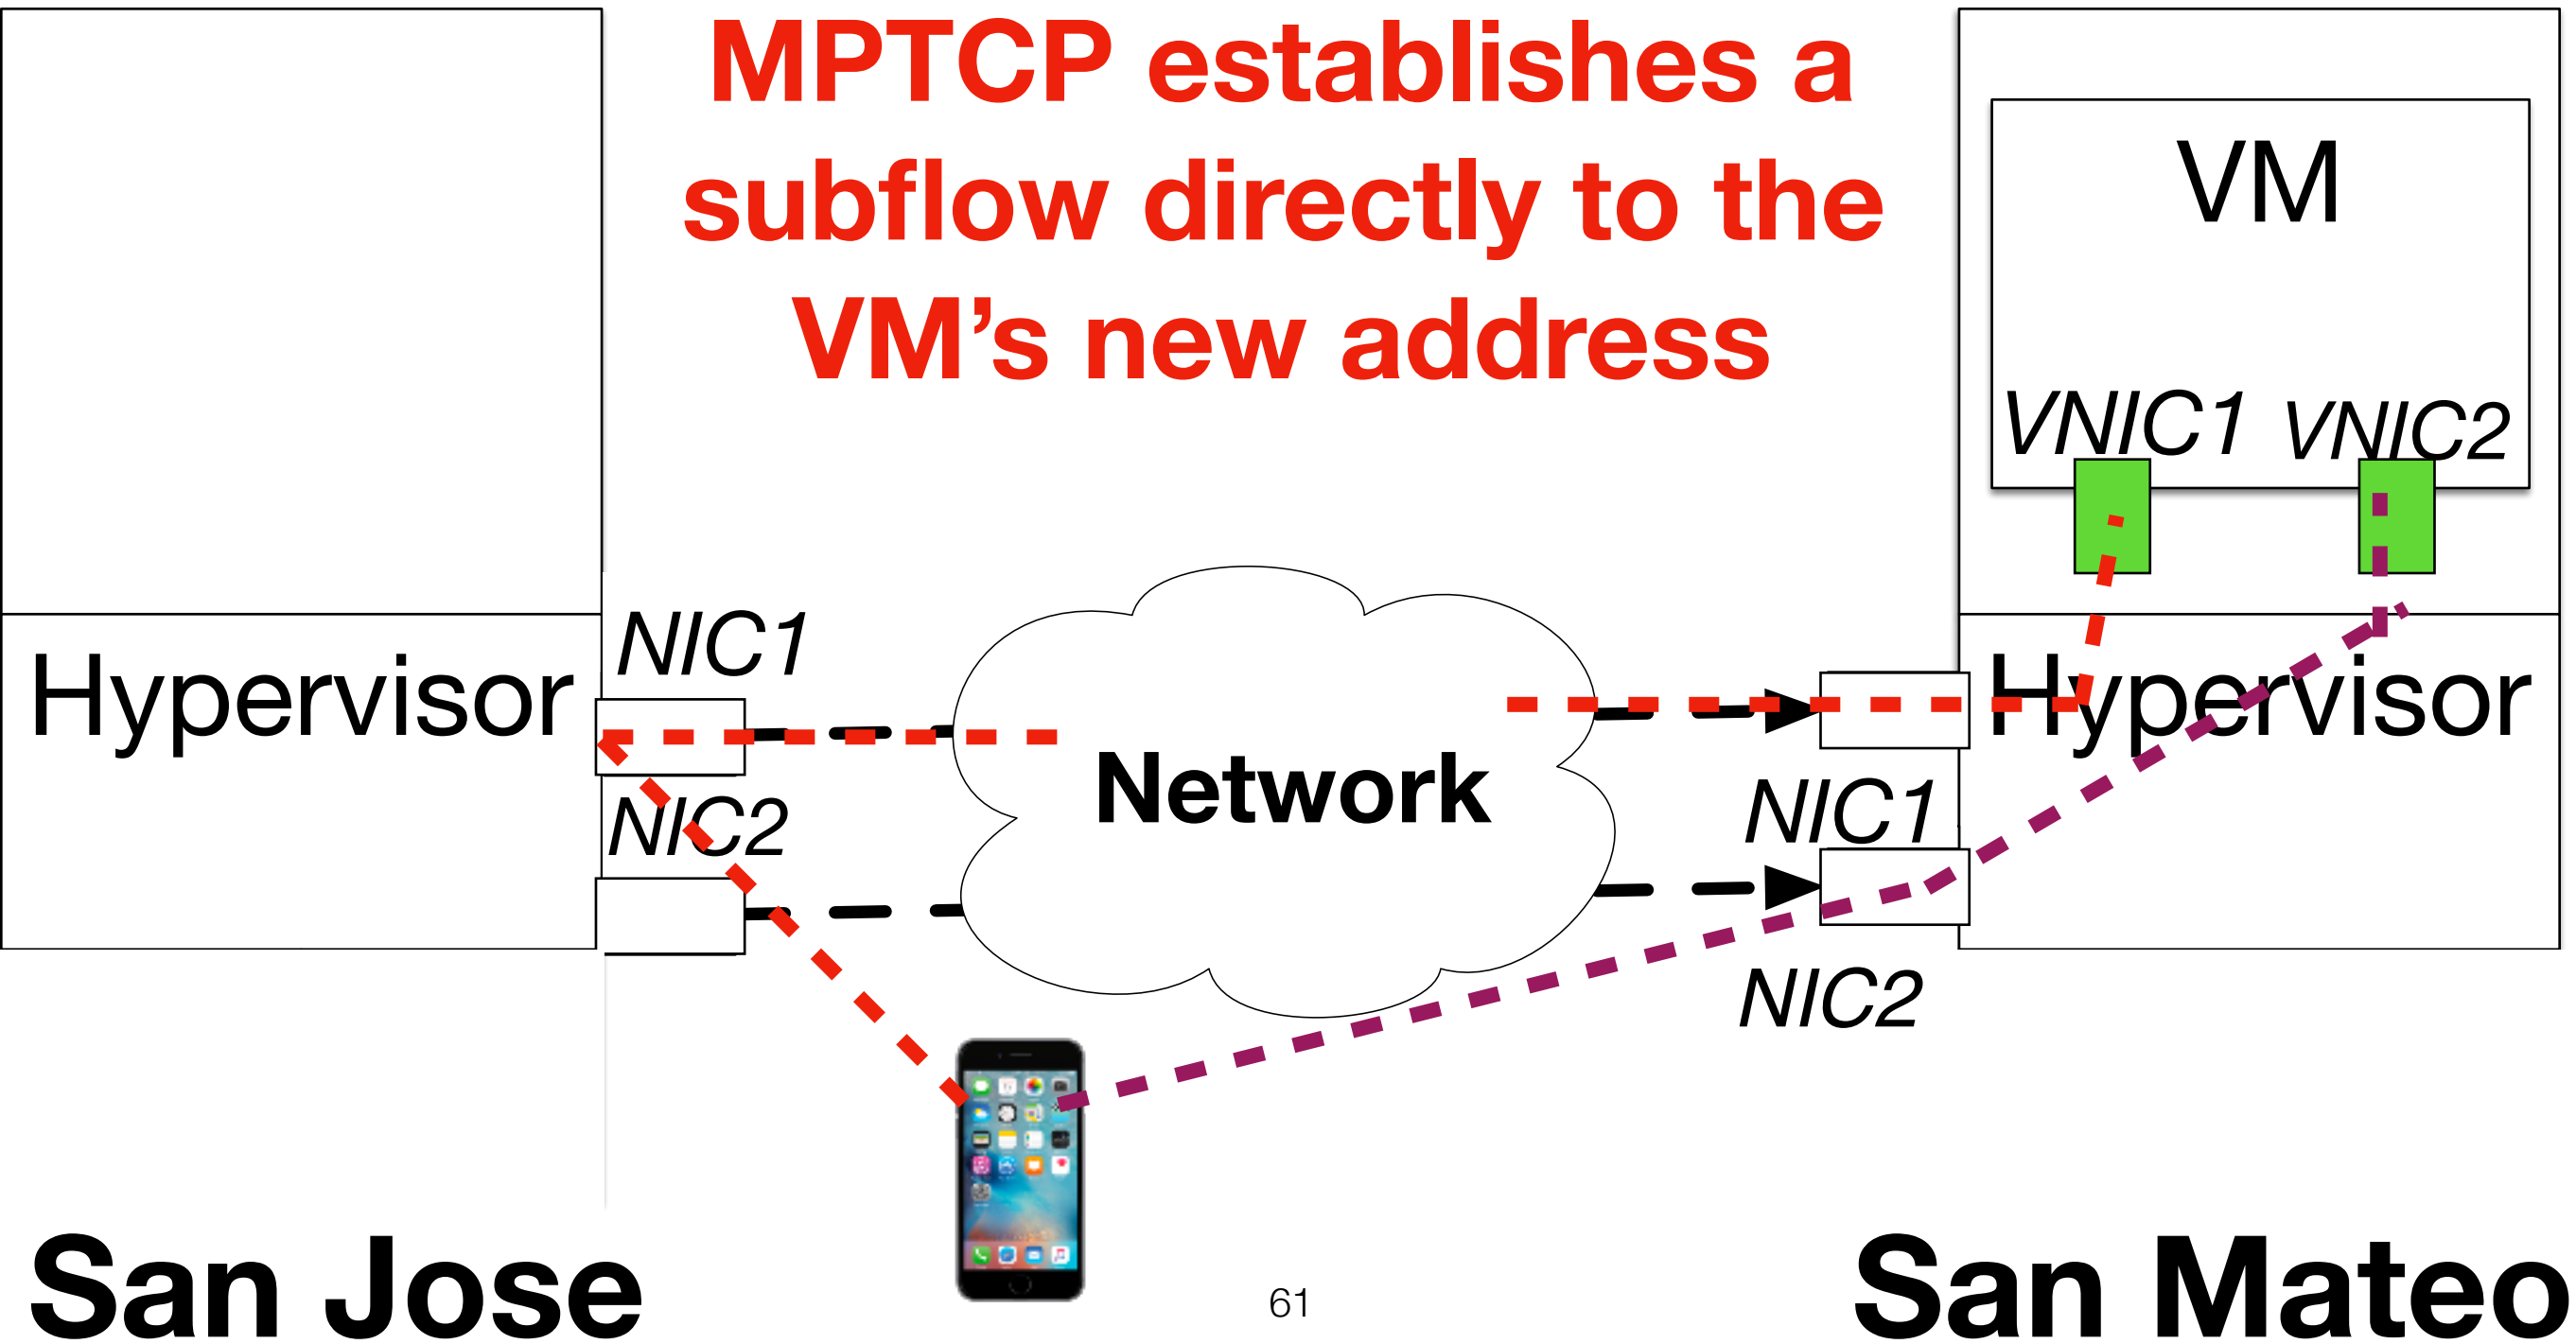


Migration



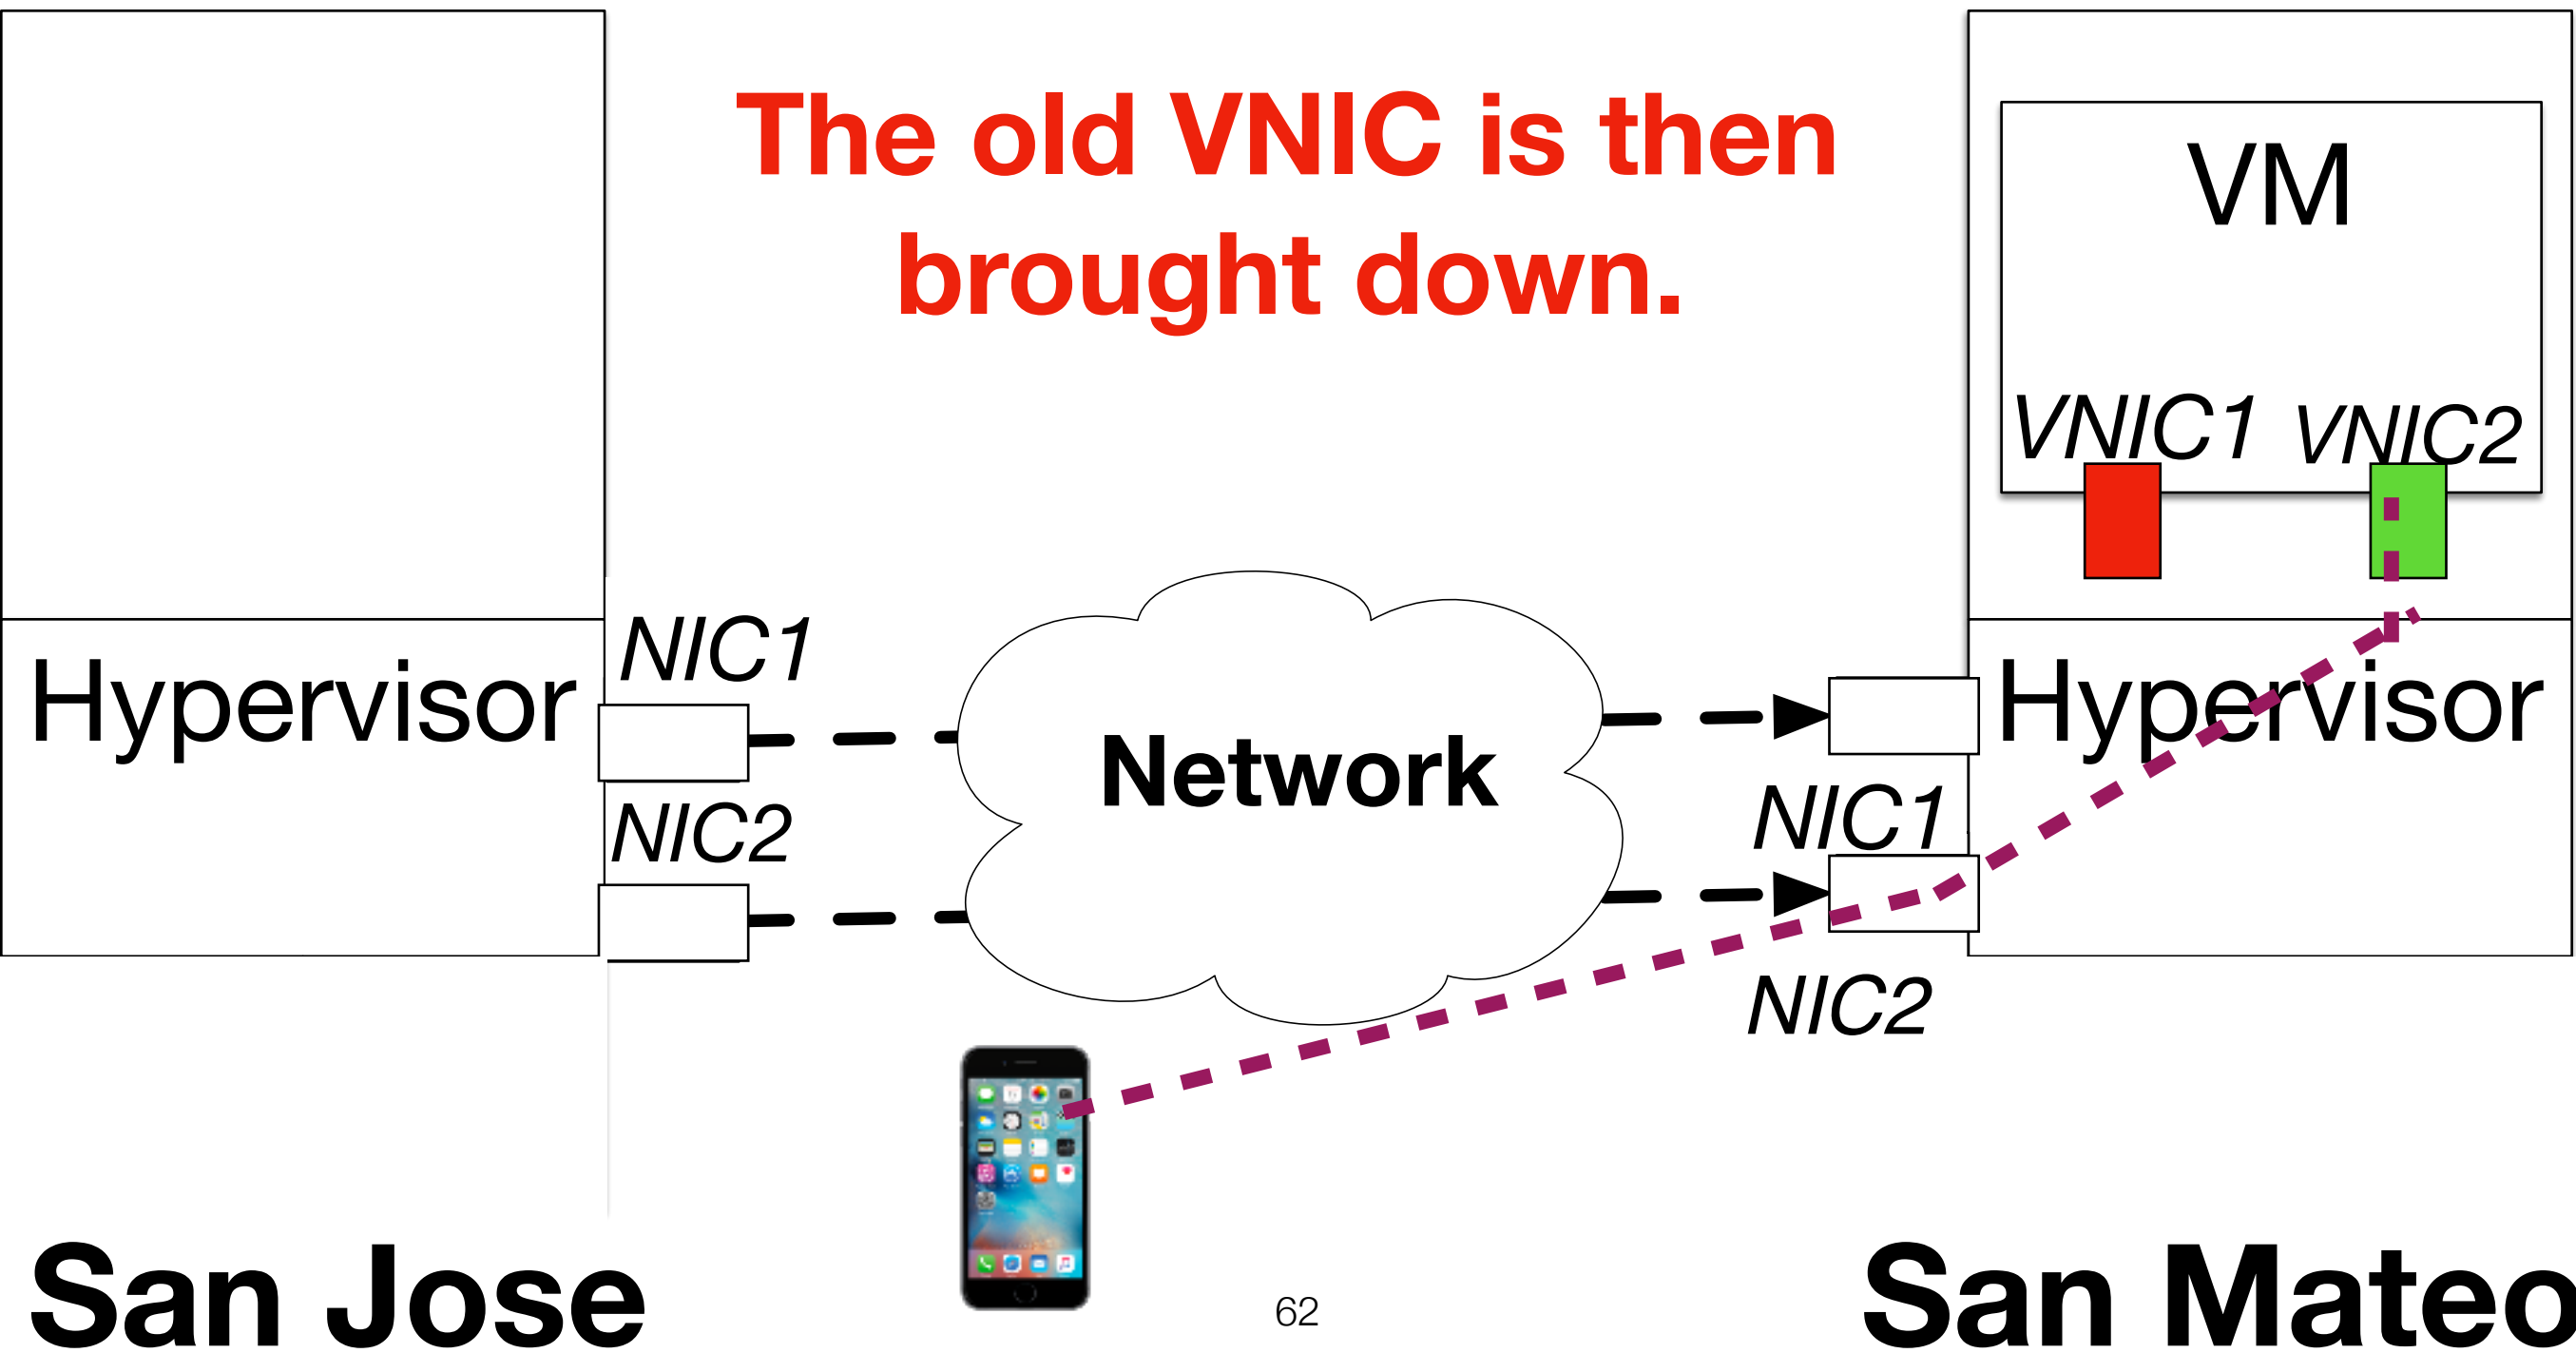
Migration

MPTCP establishes a subflow directly to the VM's new address



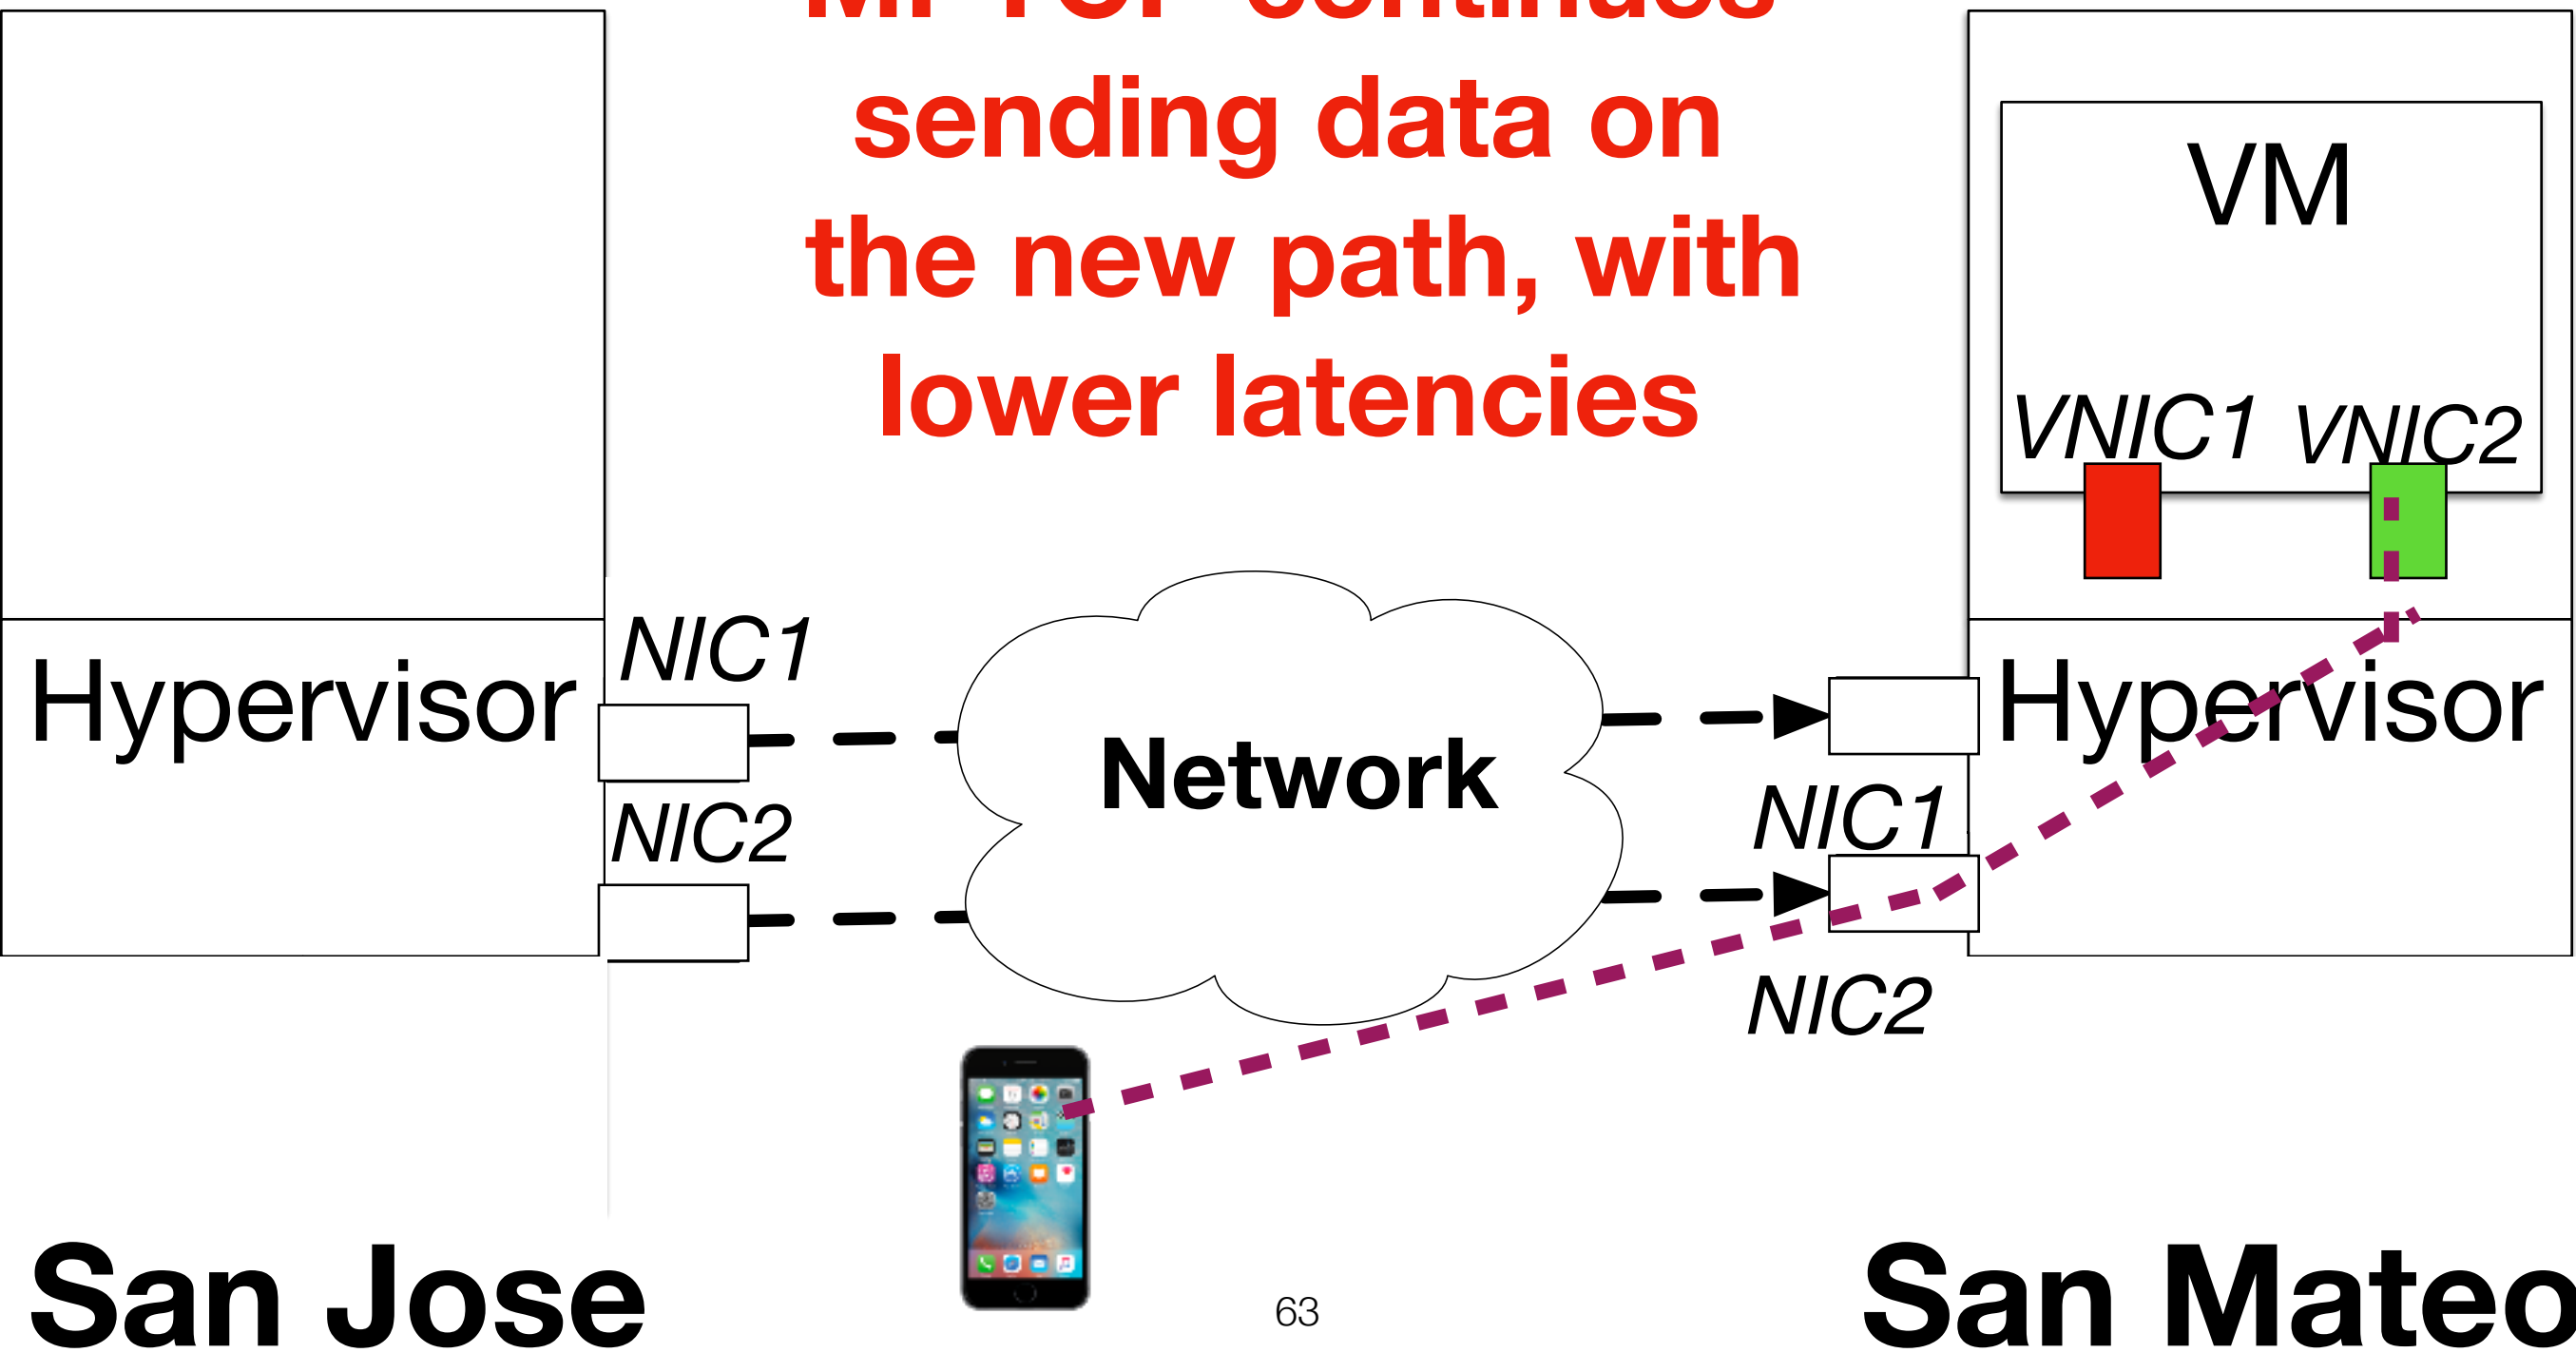
Migration

The old VNIC is then brought down.



Migration

**MPTCP continues
sending data on
the new path, with
lower latencies**




Outline


- Motivation and problem statement
- Migration & MPTCP Background
- Implementation
- **Results**
- Conclusion

Implementation & Evaluation Setup

- Tested our system on IBM's Softlayer Network along with a lab based edge cloud test bed (Dummynet).
- Use latest versions of MPTCP and KVM software.
- VM's run a workloads include dirtying VM memory and IPerf network benchmark.




Impact of Network Bandwidth




As bandwidths get lower MPTCP's value increases by almost almost 2x.

Impact of VM Application State



As VM memory size grow, MPTCP's performance increases. Allowing 2x faster migrations

Network Transparency & Guest Performance



**13x increase in after migration
throughput by with MPTCP.**

Related Work

- Barham et. al. — Xen and the art of virtualization
- Clark et. al. — Introduced Live Migration
- Wood et. al. — VM Migrations over WAN
- Shen et. al. — WAN migrations between cloud locations
- Nasim et. al. — MPTCP migrations with SDNs in LAN's
- Ha et. al. — VM Migrations in Edge Clouds

Conclusion

- VM Migrations in edge clouds is challenging.
- Our approach: Exploit MPTCP for Edge Migrations
 - Parallelizes data transfer
 - Provides Network Transparency
- 2x reduction in VM migration time while being network transparent
- Future Work: Edge Cloud Migration Policies

Thank You!
Questions?

<https://cs.umass.edu/~lucasch>